



Accelerator Architectures for Machine Learning (AAML)

Lecture 2: Basics of DNN Models

Tsung Tai Yeh
Department of Computer Science
National Yang-Ming Chiao Tung University



Acknowledgements and Disclaimer

- Slides was developed in the reference with
Joel Emer, Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, ISCA 2019
tutorial
Efficient Processing of Deep Neural Network, Vivienne Sze, Yu-Hsin
Chen, Tien-Ju Yang, Joel Emer, Morgan and Claypool Publisher, 2020
Yakun Sophia Shao, EE290-2: Hardware for Machine Learning, UC
Berkeley, 2020
CS231n Convolutional Neural Networks for Visual Recognition,
Stanford University, 2020



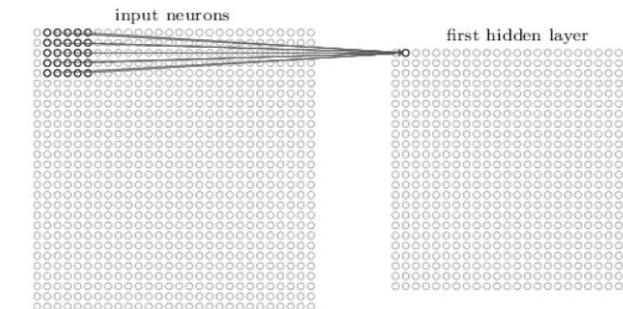
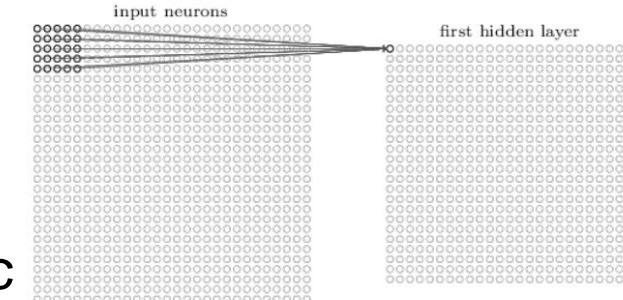
Outline

- Convolution Neural Network
 - Residual network
 - Depthwise convolution
- Transformer Model Architecture
 - Encoder-decoder based model
 - Large Language Models



Deep convolutional neural networks

- Each neuron only sees a “**local receptive field**”
 - 5 x 5 grid of neurons in this example
 - The first neuron is looking for feature in the top-left 5x5 corner of the image
 - Combines the 25 inputs with 25 synaptic weights to decide its output
 - The set of 5x5 weights as a “**filter**”

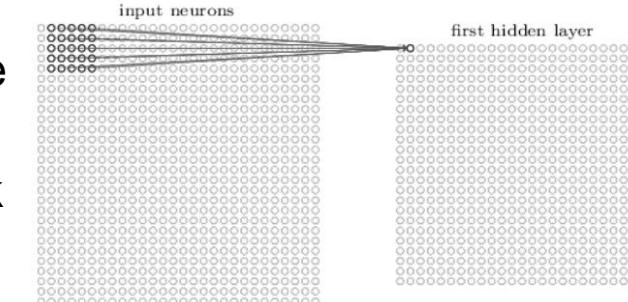
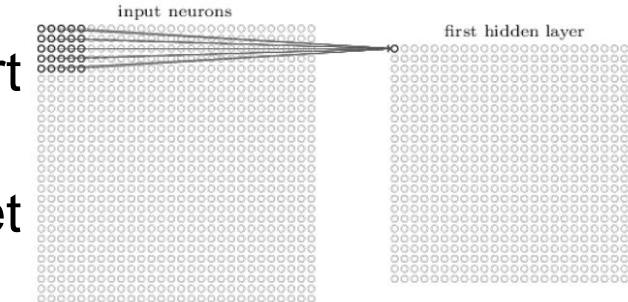




Deep convolutional neural networks

• Convolution

- Applying a 5x5 filter (kernel) to each part of the image
- All the neurons are sharing the same set of 25 weights (plus bias)
- Why do we create small size filter ?
 - The small local receptive field and the use of shared weights can help for slow learning rate in early layers of the network

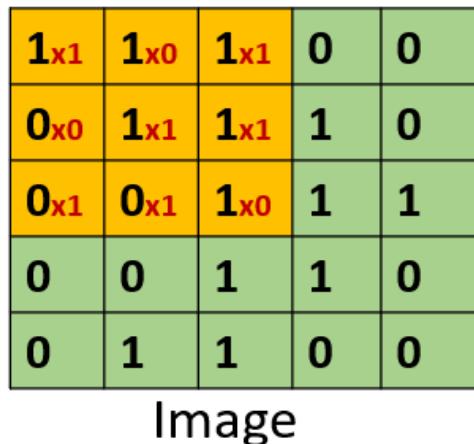




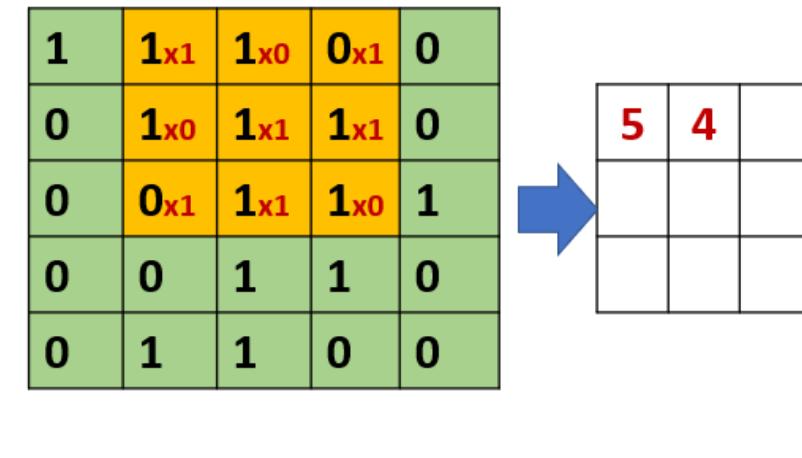
Convolutional Computation Details

- **Convolution**

- Sliding dot product or cross-correlation
- Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ filter kernel to get a $3 \times 3 \times 1$ convoluted feature



Convolved Feature





CNN Dimension Parameters

- N – Number of **input fmmaps/output fmmaps** (batch size)
- C – Number of 2D **input fmmaps/filters** (channels)
- H – Height of **input fmap** (activations)
- W – Width of **input fmap** (activations)
- R – Height of 2D **filter** (weights)
- S – Width of 2D **filter** (weights)
- M – Number of 2D **output fmmaps** (channels)
- F – Width of **output fmap** (activations)
- E – Height of **output fmap** (activations)



CONV Layer Tensor Computation

Output fmaps (Y)

Bias (B)

Input fmaps (X)

Filter weights (W)

$$Y[n][m][x][y] = \text{Activation}(B[m] + \sum_{i=0}^{R-1} \sum_{j=0}^{S-1} \sum_{k=0}^{C-1} X[n][k][Ux+i][Uy+j] \times W[m][k][i][j])$$

$$0 \leq n \leq N, 0 \leq m \leq M, 0 \leq y \leq E, 0 \leq x \leq F$$

$$E = (H - R + U)/U, F = (W - S + U)/U$$

Shape Parameter	Description
N	fmap batch size
M	# of filters or # of output fmap channels
C	# of input fmap or # of filter channels
U	Convolution stride



CONV Layer Implementation

```
for ( n = 0; n < N; n++) {  
    for (m = 0; m < M; m++) {  
        for (x = 0; x < F; x++) {  
            for (y = 0; y < E; y++) {  
                Y[n][m][x][y] = B[m];  
                for ( i = 0; i < R; i++) {  
                    for ( j = 0; j < S; j++) {  
                        for ( k = 0; k < C; k++) {  
                            Y[n][m][x][y] += X[n][k][Ux+i][Uy+j] x W[m][k][i][j]  
                        }  
                    }  
                }  
                Y[n][m][x][y] = Activation(Y[n][m][x][y]);  
            }  
        }  
    }  
}
```

For each output fmap value

CONV & Activation

How to run CONV in parallel ?



CONV Layer Parallel Implementation

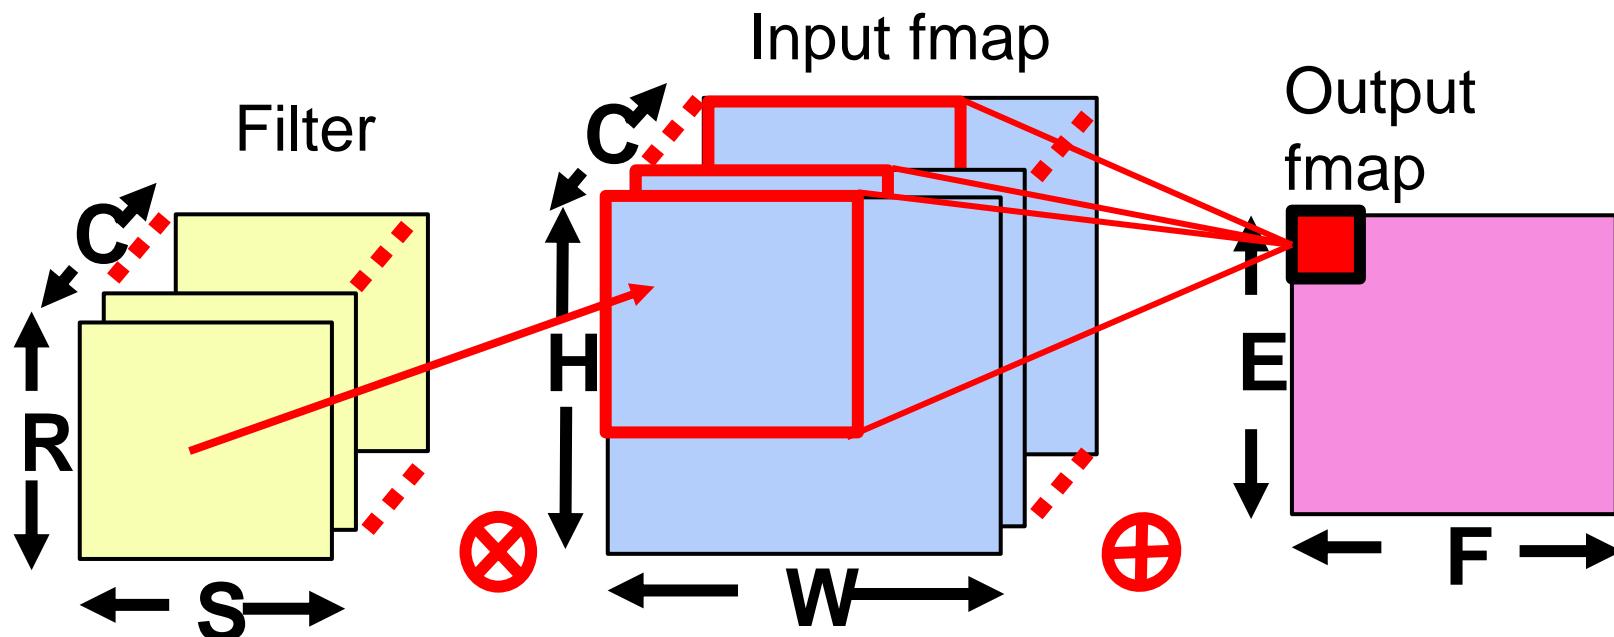
```
Parallel_for ( n = 0; n < N; n++) {  
    Parallel_for (m = 0; m < M; m++) {  
        Parallel_for (x = 0; x < F; x++) {  
            Parallel_for (y = 0; y < E; y++) {  
                Y[n][m][x][y] = B[m];  
                for ( i = 0; i < R; i++) {  
                    for ( j = 0; j < S; j++) {  
                        for ( k = 0; k < C; k++) {  
                            Y[n][m][x][y] += X[n][k][Ux+i][Uy+j] x W[m][k][i][j]  
                        }  
                    }  
                }  
                Y[n][m][x][y] = Activation(Y[n][m][x][y]);  
            }  
        }  
    }  
}
```

For each output fmap value

CONV & Activation



Convolution (CONV) Layer

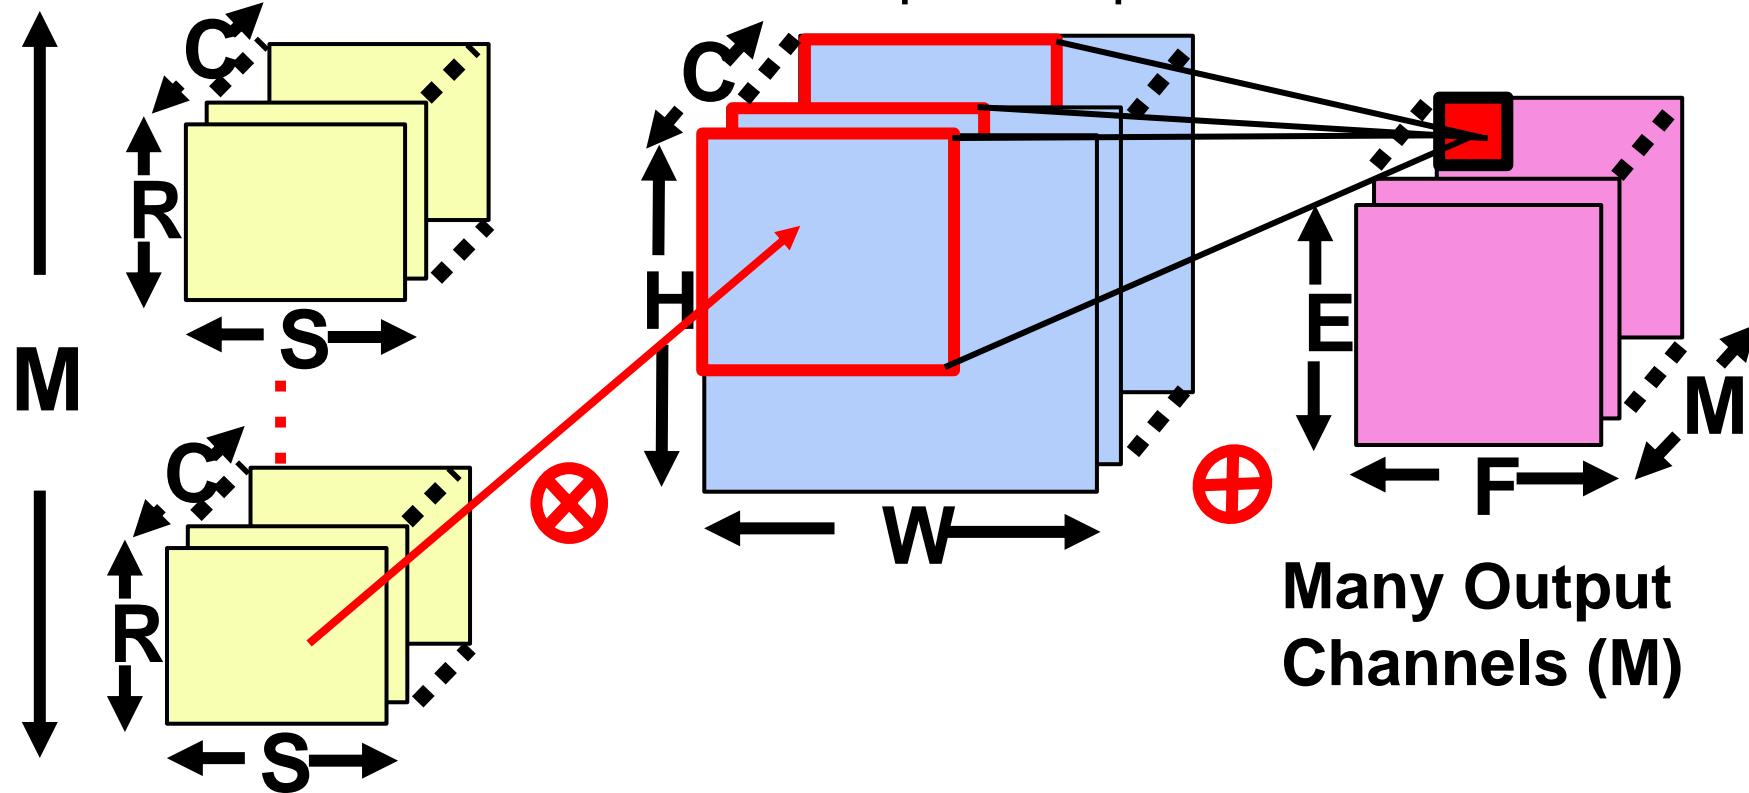


Many Input Channels (C), e.g. RGB in an image



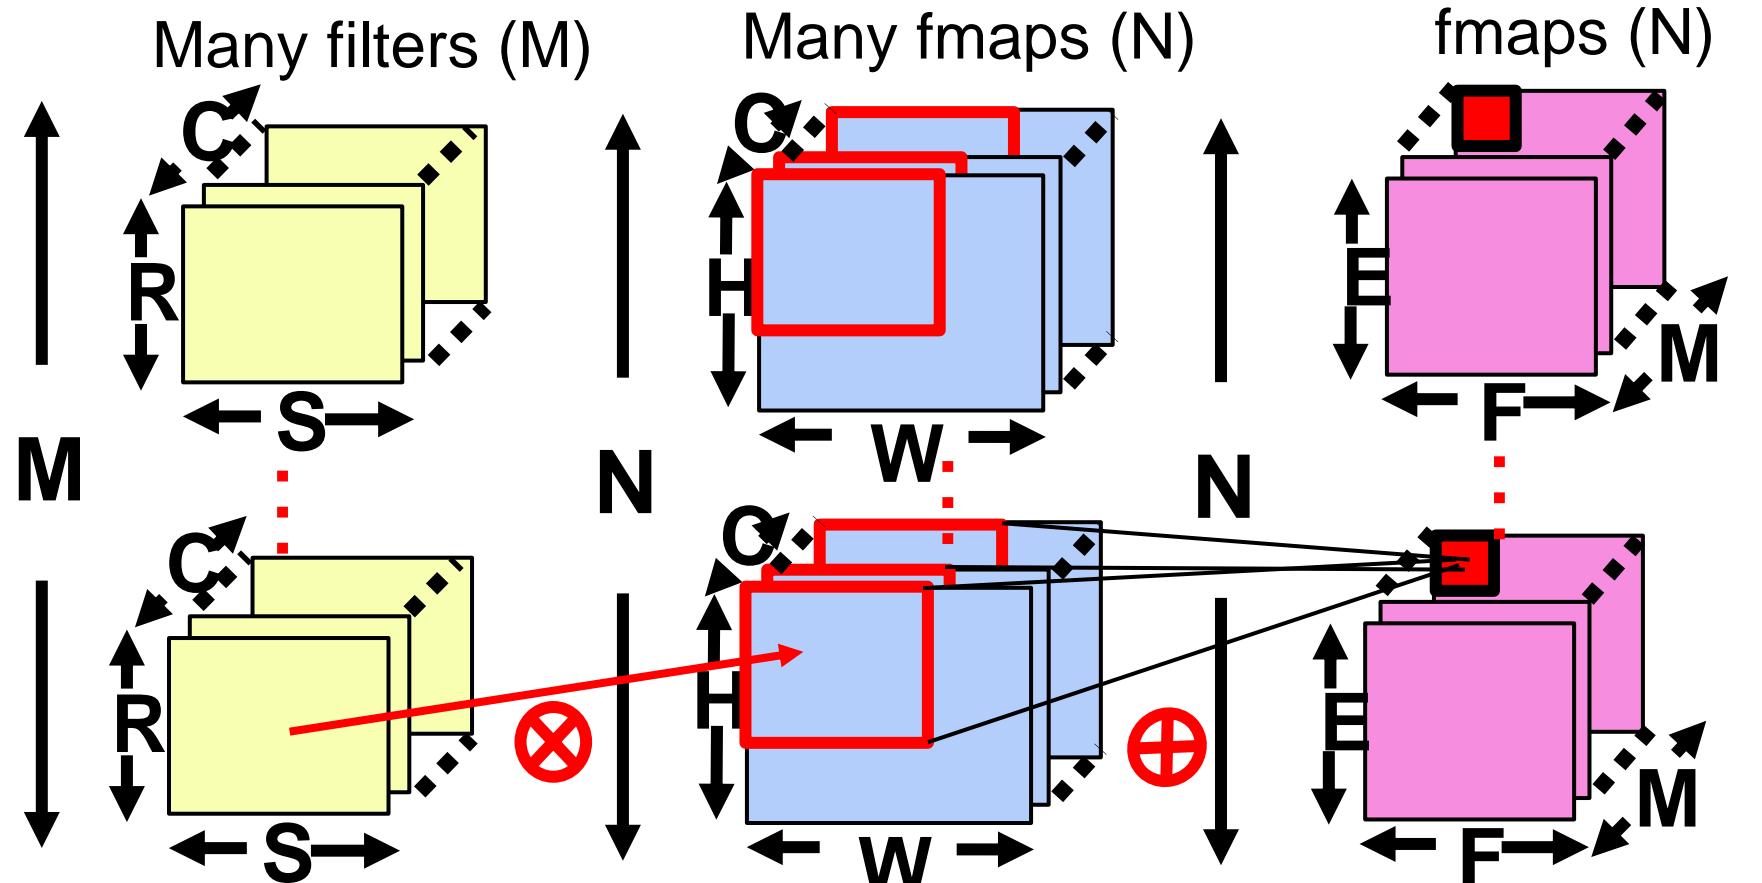
Convolution (CONV) Layer

Many filters (M)





Convolution (CONV) Layer





Pooling

- **Pooling**

- Once a feature has been found, its's exact location isn't as important as its relative location – help us reduce the parameters
- Further reduce the network, say reduce 4 neurons into a single one

2x2 pooling, stride = 2

22	8	4	9
11	3	10	0
5	8	11	3
2	1	5	9

Max pooling

22	10
8	11

11	6
4	7

Average pooling

Input fmap size: $W_1 \times H_1 \times C_1$

Spatial extent: F

Stride S

Output fmap after pooling: $W_2 \times H_2 \times C_2$

$$W_2 = (W_1 - F) / S + 1$$

$$H_2 = (H_1 - F) / S + 1$$

$$C_2 = C_1$$



POOL Layer Implementation

```
for ( n = 0; n < N; n++) {  
    for (m = 0; m < M; m++) {  
        for (x = 0; x < F; x++) {  
            for (y = 0; y < E; y++) {  
                max = -Inf  
                for ( i = 0; i < R; i++) {  
                    for ( j = 0; j < S; j++) {  
                        if ( X[n][m][Ux+i][Uy+j] > max) {  
                            max = X[n][m][Ux+i][Uy+j];  
                        }  
                    }  
                }  
                Y[n][m][x][y] = max;  
            }  
        }  
    }  
}
```

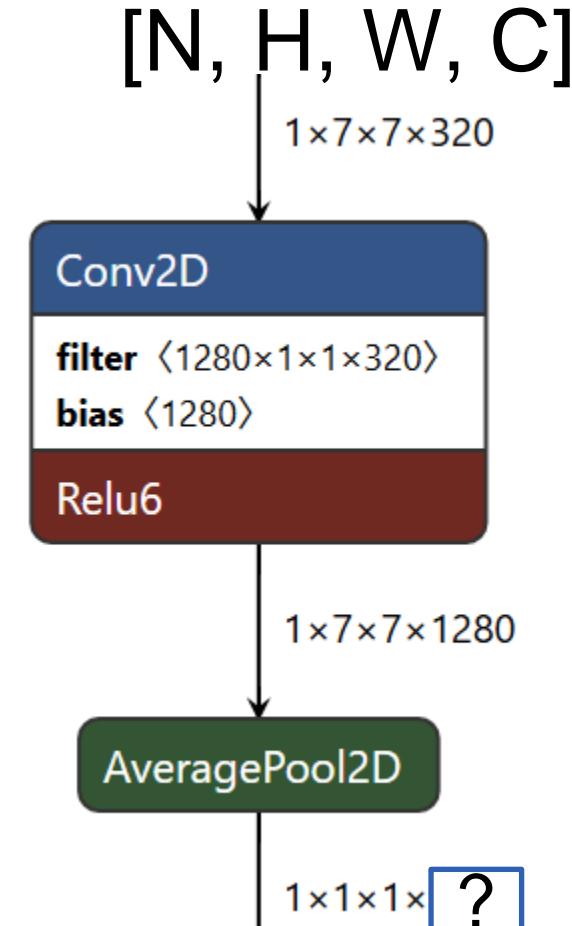
} for each pooled value

} Find the max in each window



Takeaway Questions

- What is the output channel of AveragePool2D?
 - (A) 1280
 - (B) $1280/7$
 - (C) 640





Online Survey

- How to reduce the MACs of the CNN model?

<https://reurl.cc/ZNqMZQ>

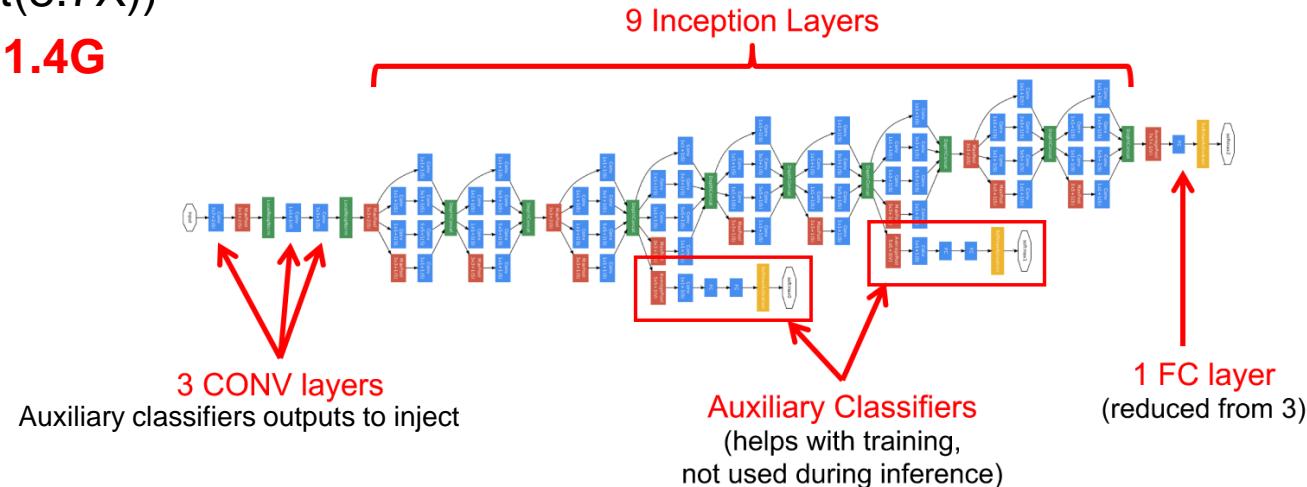




GoogleNet Inception Architecture

- 22 layers
- Fully-Connected Layers: 1
- **Weights: 7.0 M** (< VGG(19.7X)
AlexNet(8.7X))
- MACs: **1.4G**

ILSCVR14 Winner
GoogleNet is used to classify images
GoogleNet top-5 error rate is 6.67%
over VGG 7.3%

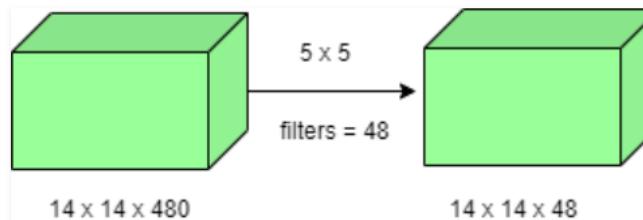




What's New in GoogleNet ?

- **1 x 1 CONV filter (why?)**

- Decrease the number of parameters (weights and biases)
- Increase the depth of the network

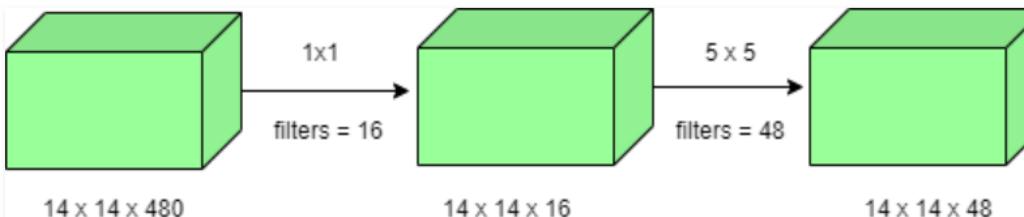


Case 1: 5x5 filter, # of filter = 48

Total MACs: $(14 \times 14 \times 48) \times (5 \times 5 \times 480) = 112.9M$

Case 2: 1x1 filter, # of filter = 16 as intermediate

Total MACs: $(14 \times 14 \times 16) \times (1 \times 1 \times 480) + (14 \times 14 \times 48) \times (5 \times 5 \times 16) = 5.3M$



GoogleNet can be trained in a single machine such as (GPU with limit memory space) !!

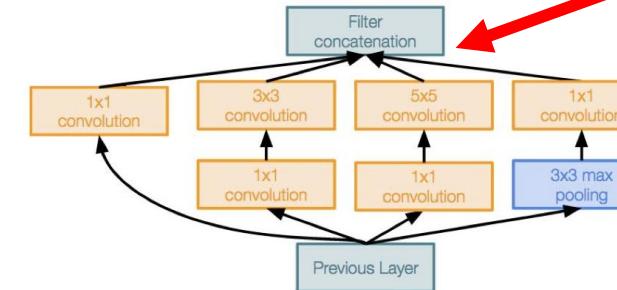


What's New in GoogleNet ?

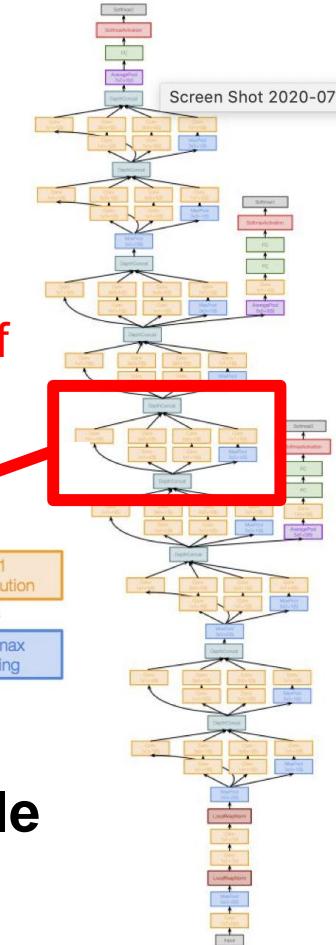
- **Inception module**

- A local network topology (network within a network)
- Stack modules on top of each other
- Multiple receptive field sizes for CONV (1×1 , 3×3 , 5×5)
- Pooling operation (3×3)
- Depth-wise filter concatenation

What's the problem of inception module?



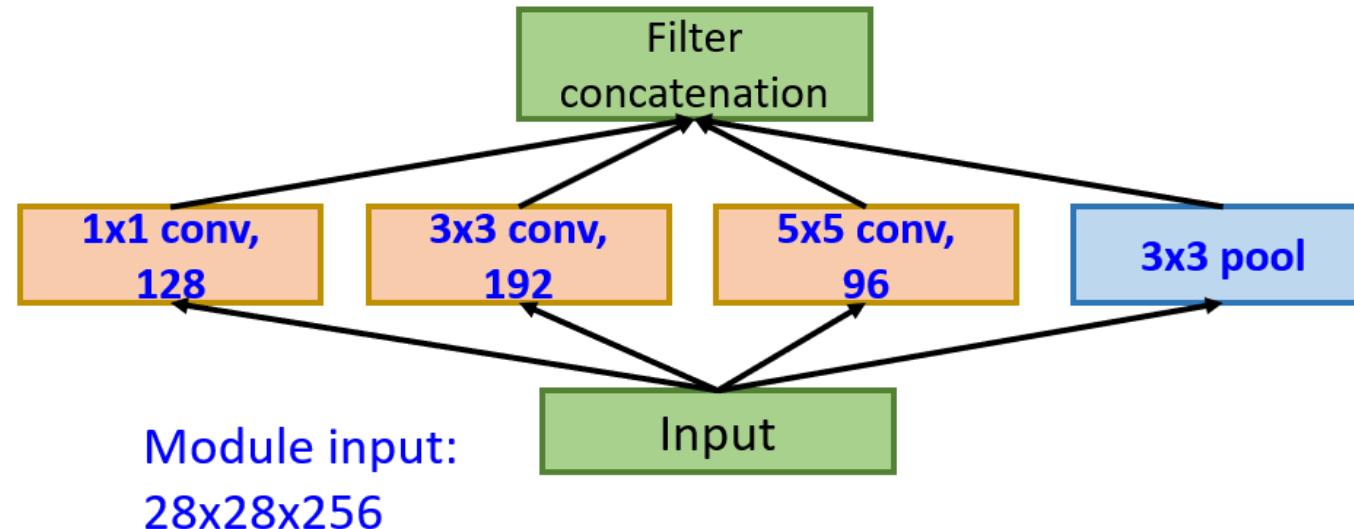
Inception module





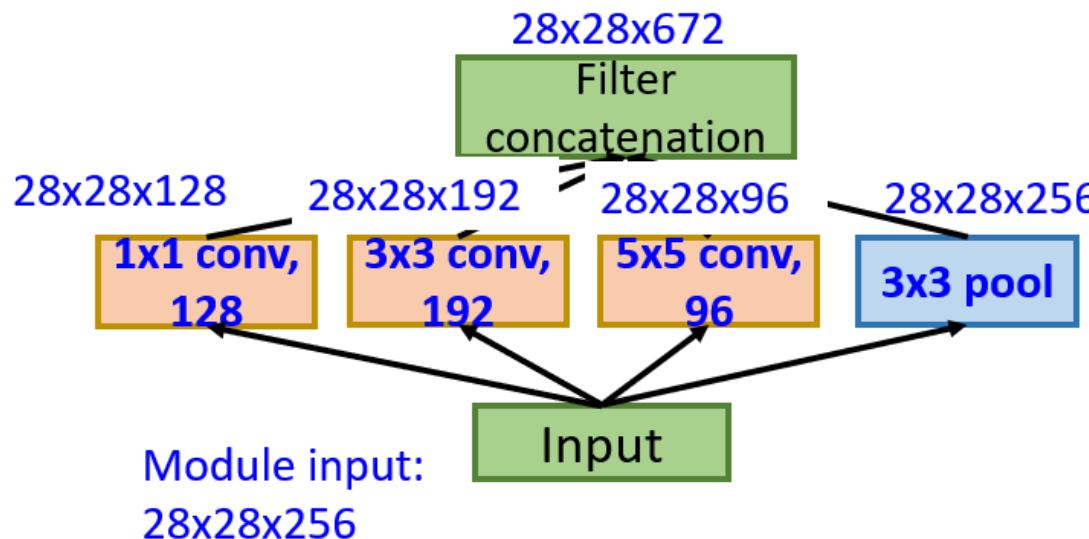
GoogleNet Inception Module Problems?

- What is the output size of 1x1 conv, with 128 filters ?





GoogleNet Inception Module Problems?



CONV Ops:

[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$
[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$
[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

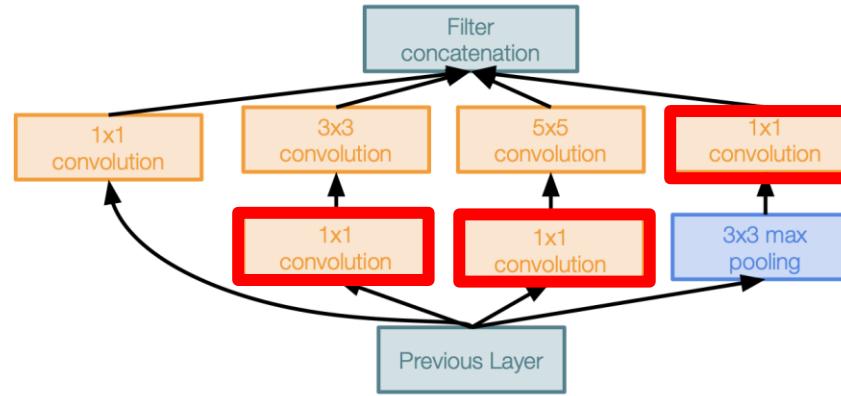
Total: 854 M ops

Very expensive compute

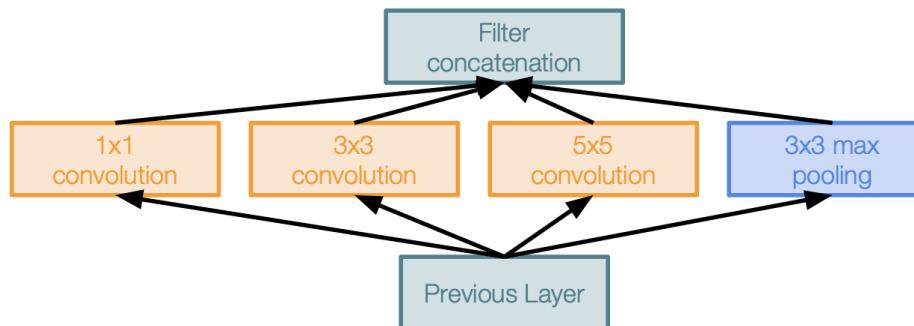
Solution: “bottleneck” layers that use 1x1 convolutions to reduce feature depth



Dimension Reduction on GoogleNet



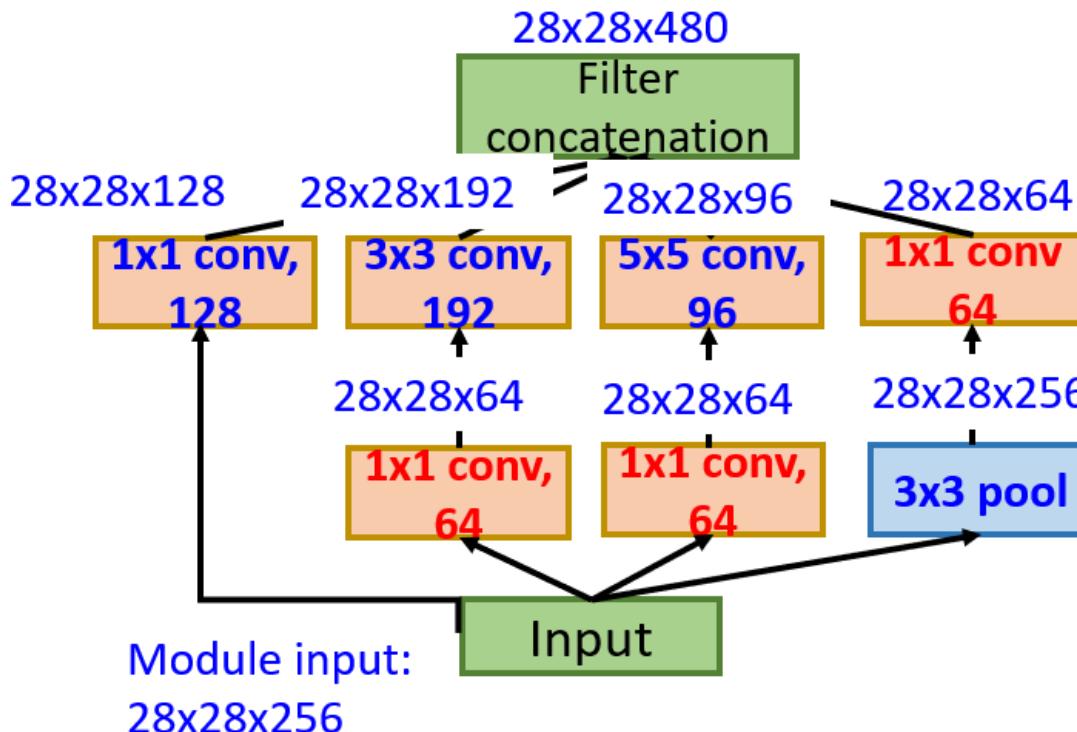
Inception module with
dimension reduction
using
**1x1 conv “bottleneck”
layers**



Naïve inception module



GoogleNet 1x1 Bottleneck Layer Ops



Conv Ops:

- [1×1 conv, 64] $28 \times 28 \times 64 \times 1 \times 1 \times 256$
- [1×1 conv, 64] $28 \times 28 \times 64 \times 1 \times 1 \times 256$
- [1×1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$
- [3×3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 64$
- [5×5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 64$
- [1×1 conv, 64] $28 \times 28 \times 64 \times 1 \times 1 \times 256$

Total: 358 M ops

Naïve version has **854M ops**

Bottleneck layer can reduce ops using dimension reduction



Online Survey

- How to train a deep neural network?

<https://reurl.cc/ZNqMZQ>

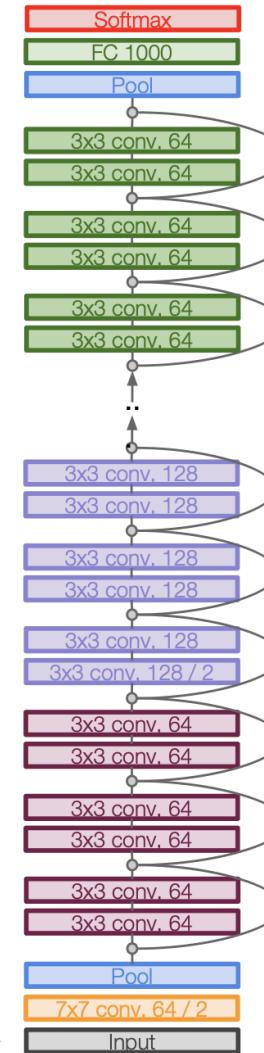
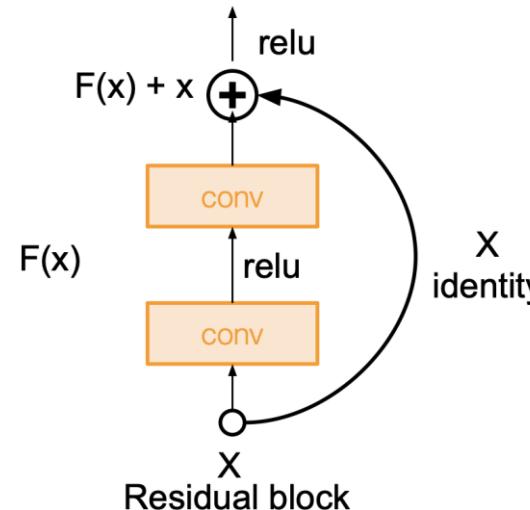




ResNet Model Overview

- 152-layer model for ImageNet Classification
- ILSVRC'15 winner (3.57% top-5 error)
- Using residual blocks and connections

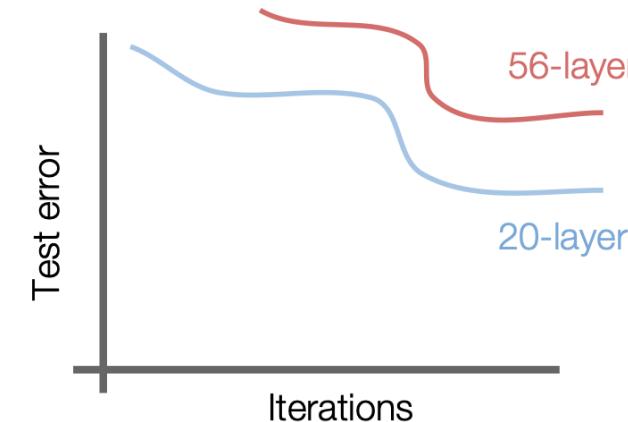
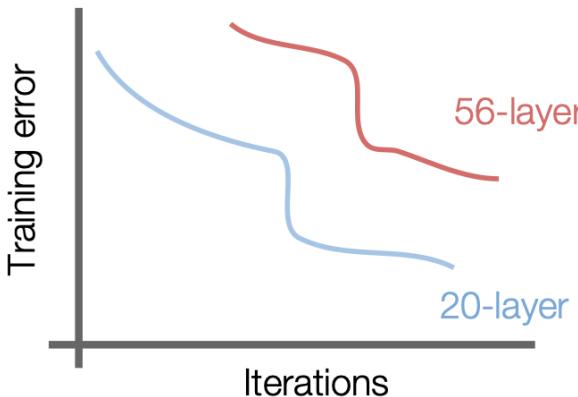
How to train the data in ULTRA-DEEP network (over 1000 layers)?





Deep Network Training Problems on ResNet

- Training and test error are increased with the length of networks



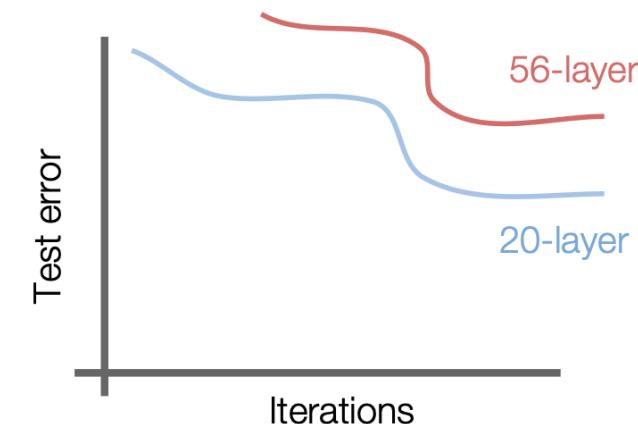
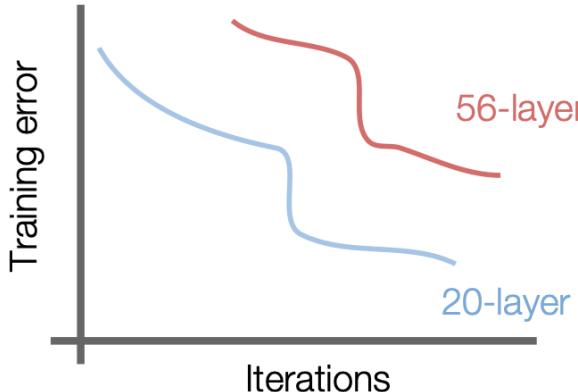
What is wrong when increasing the length of networks?

- Deeper model performs worse on both training and test error (overfitting?)



Deep Network Training Problems on ResNet

- Why overfitting isn't the main reason to increase error rate of 56-layer?
- Hypothesis: **vanishing gradient** raises error rate of ultra-deep networks?
 - Solution: Add layers to fit a residual mapping instead of fitting a desired underlying mapping directly (skipping connection)(What?)

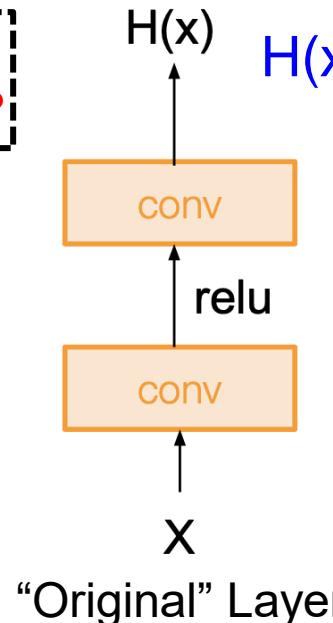




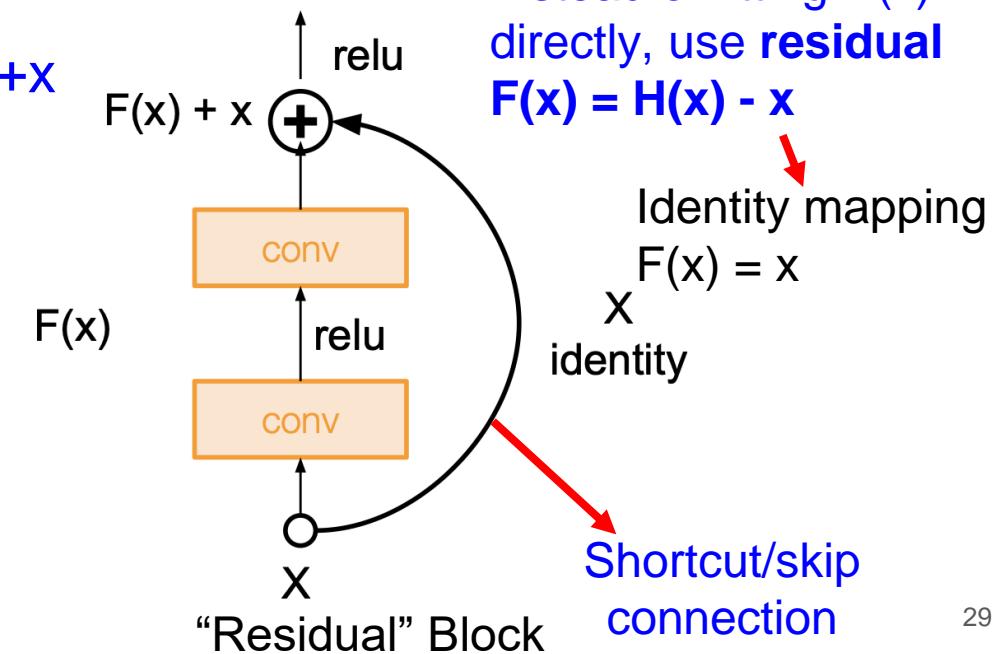
Deep Network Training Problems on ResNet

- Solution: Add layers to fit a residual mapping instead of fitting a desired underlying mapping directly

Why not use
 $H(x)$ directly?



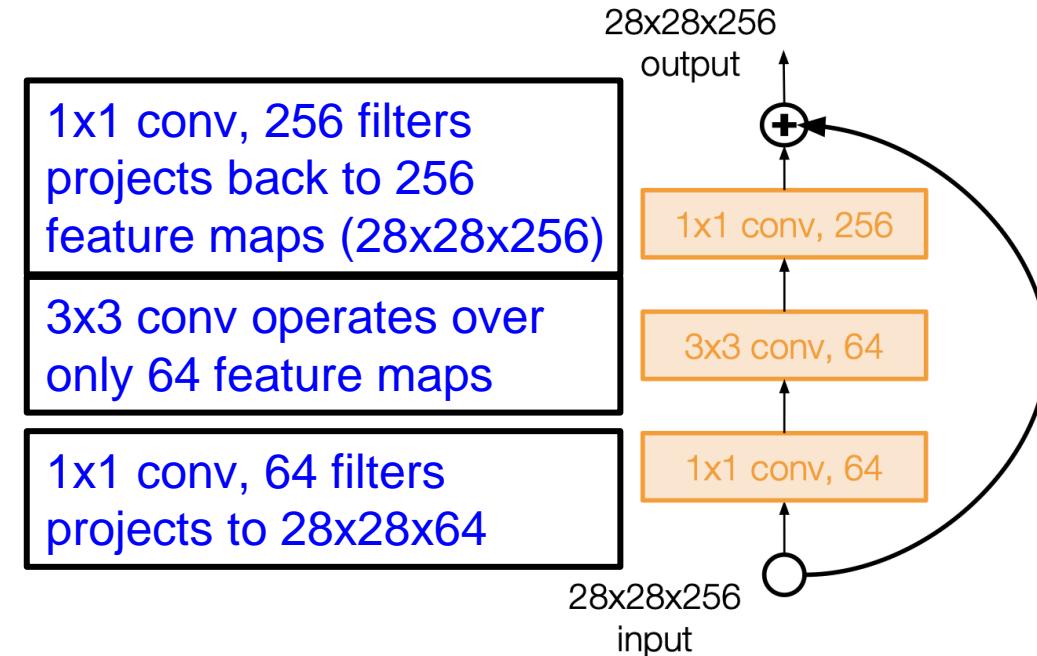
$$H(x) = F(x) + x$$





Bottleneck Layer on ResNet

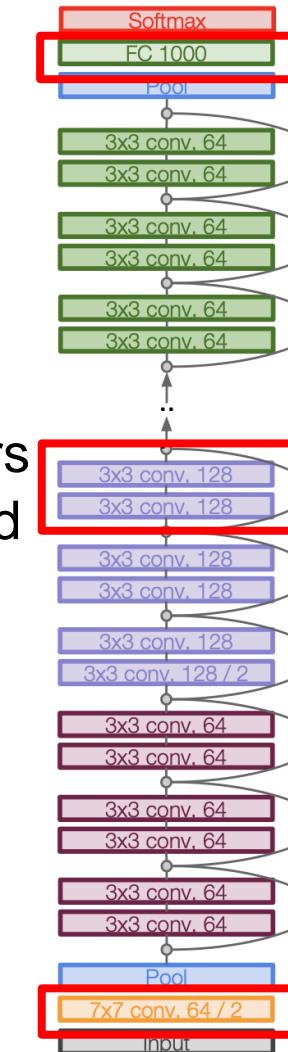
- ResNet50+ also uses “bottleneck” layer to improve efficiency for deep networks (similar to GoogleNet)





ResNet Model Details

- Full ResNet architecture
 - Stack residual blocks
 - Every residual block has two 3x3 conv layers
 - Periodically, double the number of filters and down-sample using stride 2 (/2 in each dimension)
 - Additional conv layer at the beginning
 - Only FC 1000 to output class
 - Total depths of 34, 50, 101 or 152 layers for ImageNet





Online Survey

- How to reduce model parameters?

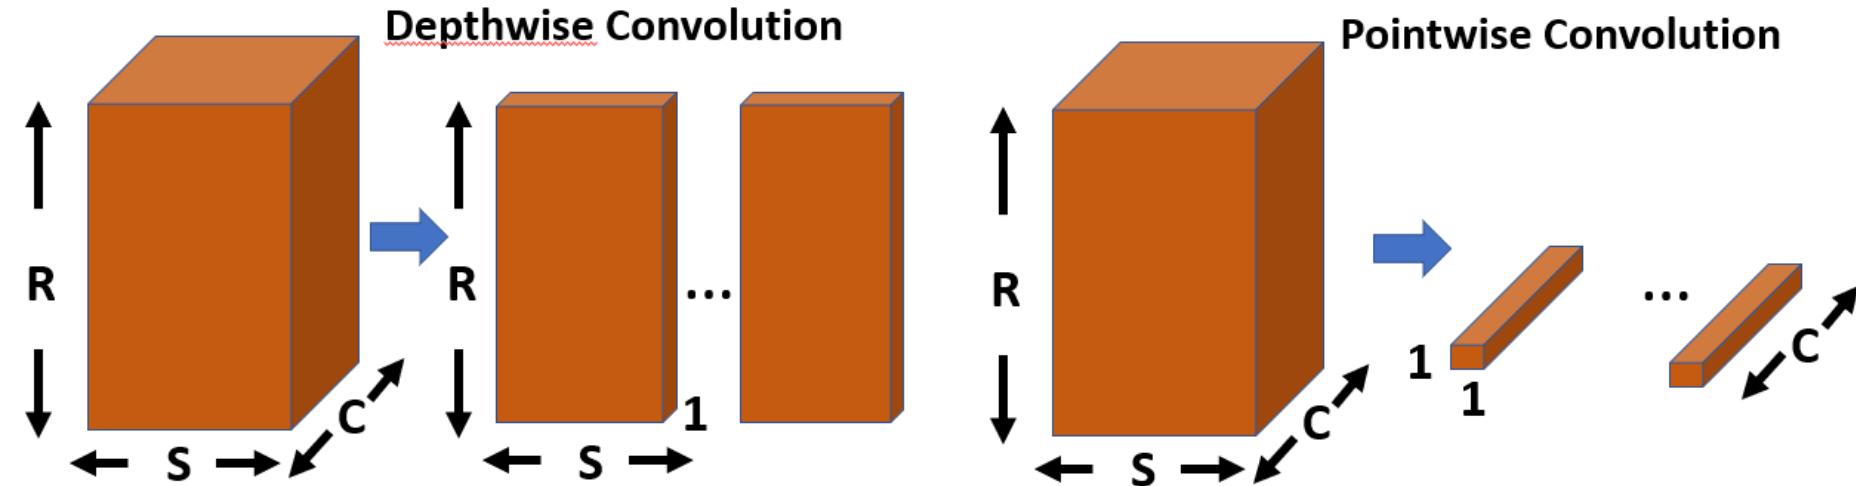
<https://reurl.cc/ZNqMZQ>





MobileNet v1: Depthwise Separable CONV

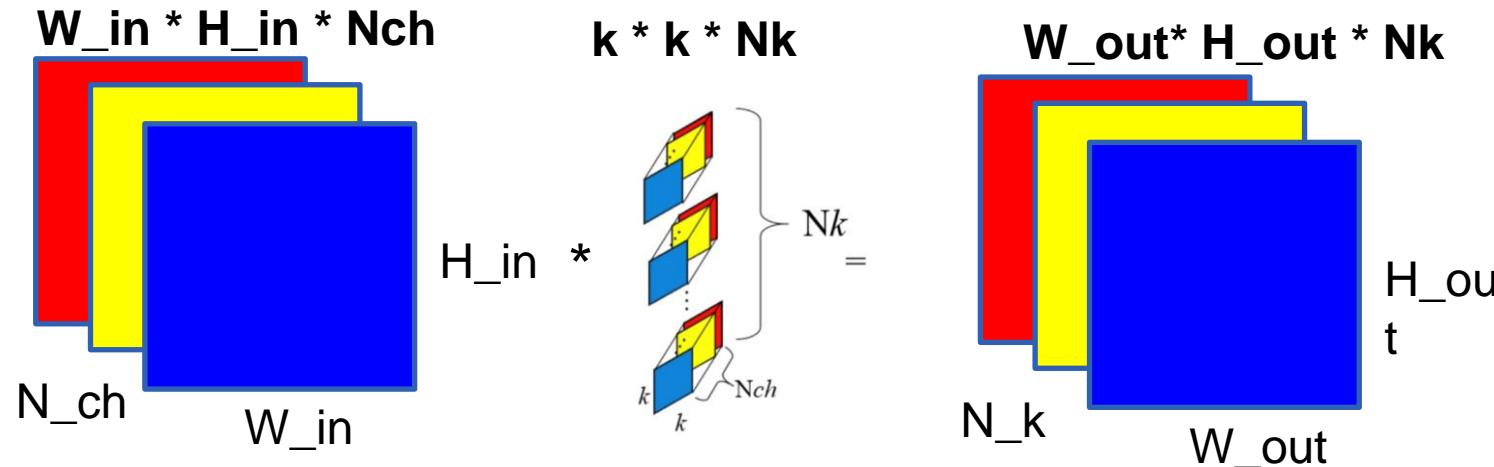
- Decouple **cross-channel correlations** and **spatial correlations** in the feature maps
- How to reduce # of parameters? [Andrew et. al. arxiv, 2017]





What is Depthwise Separable Convolution ?

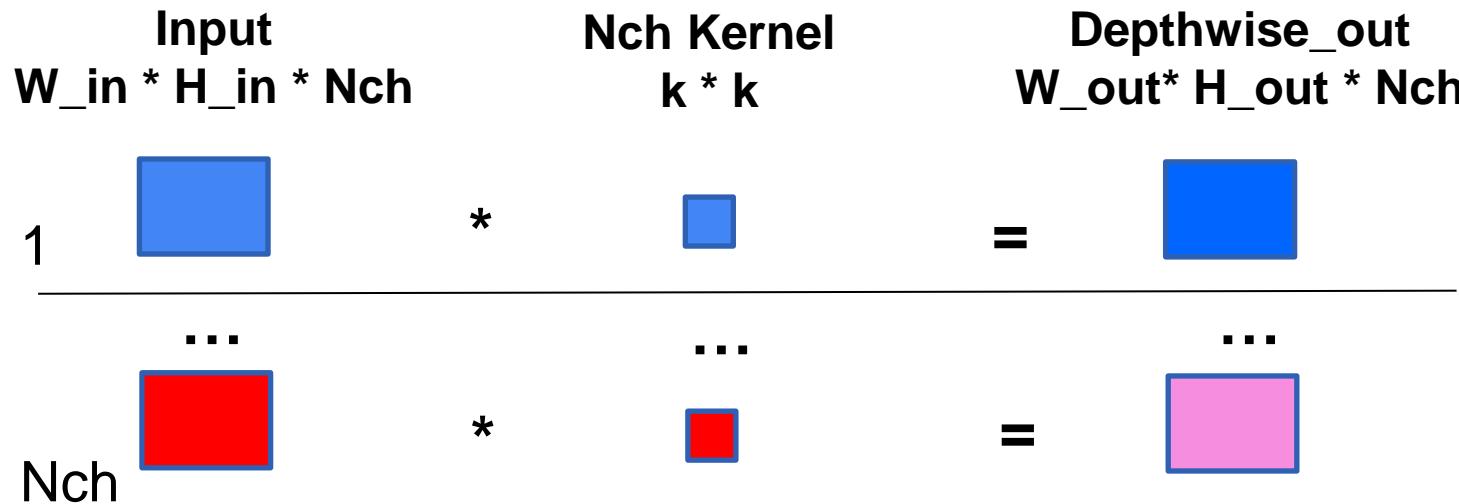
- **Purpose:** Reduce the amount of CONV computation
- **Input:** $W_{in} * H_{in} * Nch$ (# of channels)
- **Kernel:** $k * k * Nk$ (# of kernels)
- **Output:** $W_{out} * H_{out} * Nk$ (# of kernels)





Depthwise Convolution

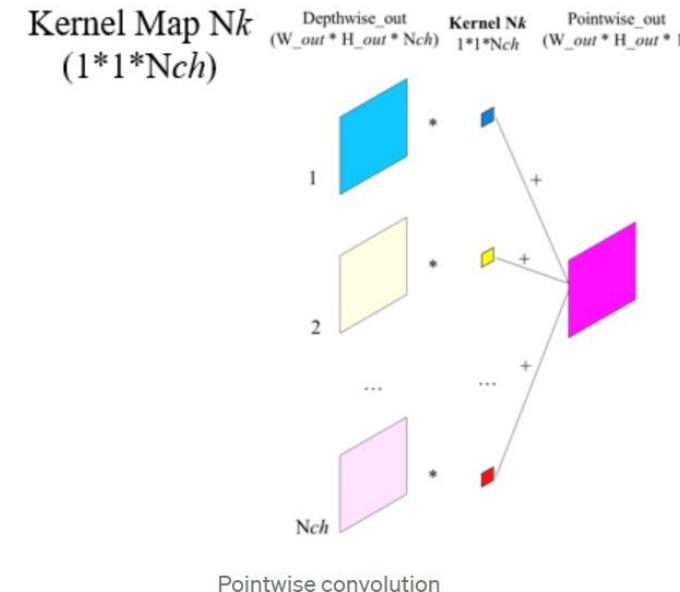
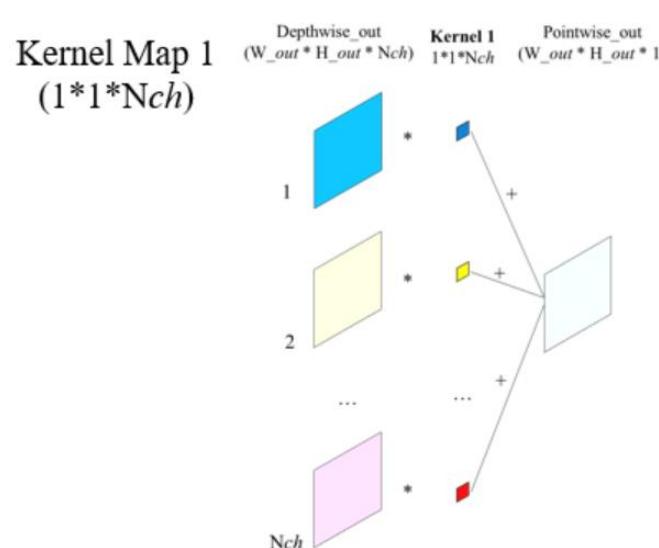
- Each channel of inputs has a $k * k$ kernel
- Separate the convolution of each channel
- **Difference:** Every kernel convolves with all channels in standard CONV





Pointwise Convolution

- The number of kernel: N_k with $(1 * 1 * N_{ch})$ size
- Do CONV on the outputs of depthwise convolution





Depthwise + Pointwise Convolution

Depthwise convolution

Input: $W_{in} * H_{in} * Nch$

Nch Kernel ($k * k$)

Output: $W_{out} * H_{out} * Nch$

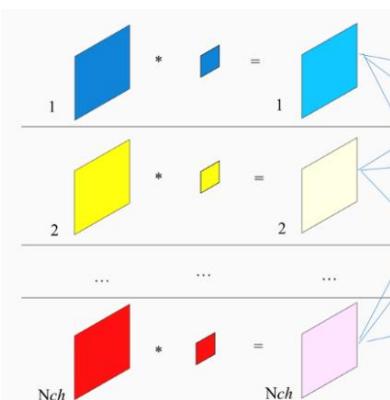
Pointwise convolution

Input: $W_{out} * H_{out} * Nch$

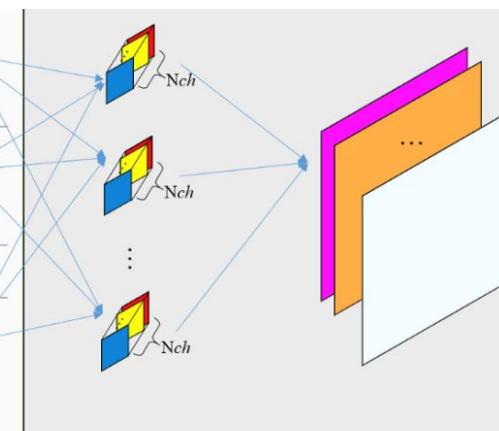
Nk kernel = $(1 * 1 * Nk)$

Output = $W_{out} * H_{out} * Nk$

Depthwise convolution



Pointwise convolution





Depthwise Separable Convolution

- **Standard CONV**

- Input: $W_{in} * H_{in} * Nch$
- Kernel: $k * k * Nk$
- Output: $W_{out} * H_{out} * Nk$
- Computation: $W_{in} * H_{in} * Nch * k * k * Nk$

- **Depthwise separable convolution**

- Depthwise CONV computation: $W_{in} * H_{in} * Nch * k * k$
- Pointwise CONV computation: $Nch * Nk * W_{in} * H_{in}$

$$\frac{\text{Depthwise separable CONV}}{\text{Standard CONV}} = \frac{W_{in} * H_{in} * Nch * k * k + Nch * Nk * W_{in} * H_{in}}{W_{in} * H_{in} * Nch * k * k * Nk}$$
$$= \frac{1}{Nk} + \frac{1}{K*k}$$



Depthwise Separable Convolution

- Depthwise separable convolution can save more computation when
 - kernel size is large
 - The number of kernel is increased
- Suppose input is $416 * 416 * 50$, # of filter is 10, its size is $3 * 3$.
- How much computation can be saved by depthwise separable convolution ?
 - $1/10 + 1/9 = 0.22$



Takeaway Questions

- What are problems in ultra-deep neural networks ?
 - (A) Over-fitting
 - (B) Gradient vanishing
 - (C) Low training accuracy
- Given a CNN model below, how many channels are in the second layer ?
 - (A) 4
 - (B) 8
 - (C) 16

	Input size	# of filter	Filter size	# of channel
Layer1	12 x 12	4	3x3	64
Layer2	12 x 12	16	3x3	



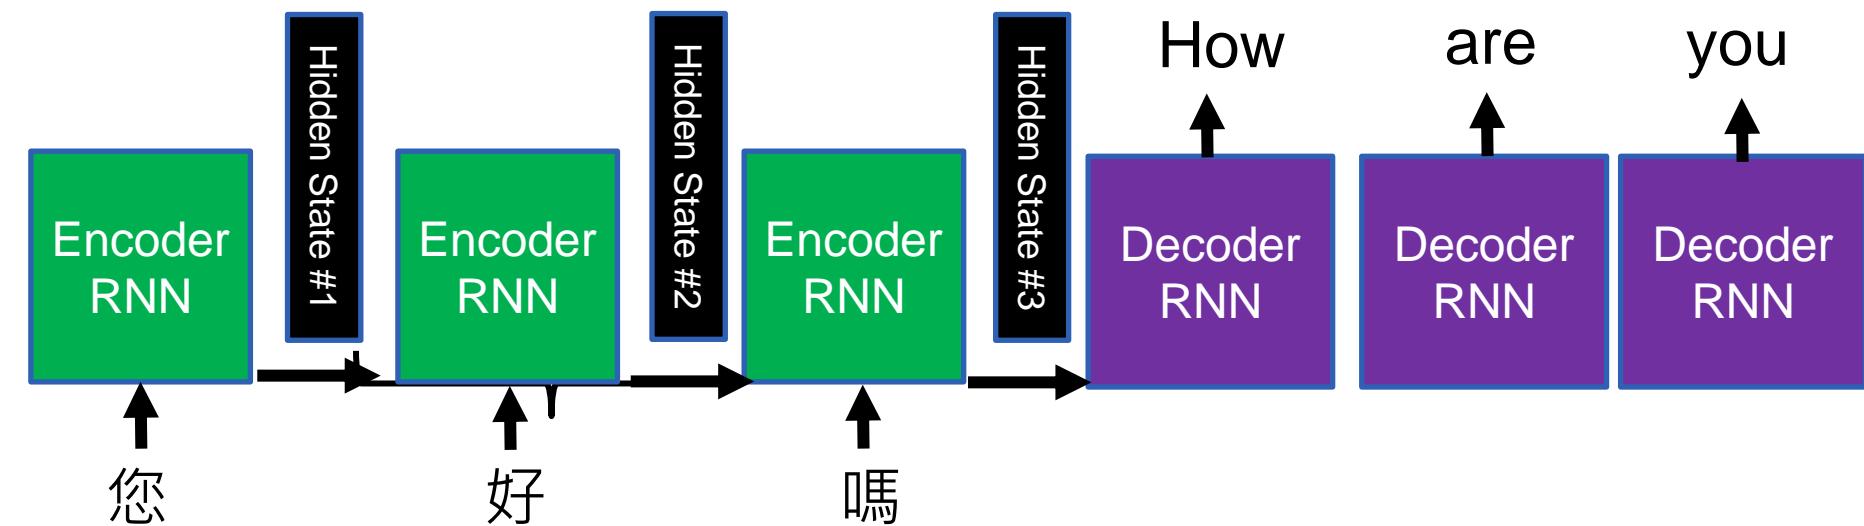
Takeaway Questions

- A standard CONV layer
 - Input: $W_{in} * H_{in} * Nch = (32 * 32 * 16)$
 - Kernel: $k * k * Nk = (3 * 3 * 8)$
 - Computation: $W_{in} * H_{in} * Nch * k * k * Nk = (32 * 32 * 16 * 3 * 3 * 8)$
- What is the amount of computation that is carried out by **depthwise separable convolution** ?
 - (A) $(32 * 32 * 16 * 3 * 3) + (3 * 8 * 8)$
 - (B) $(32 * 32 * 3 * 3 * 8) + (16 * 3 * 3)$
 - (C) $(32 * 32 * 16 * 3 * 3) + (16 * 8 * 32 * 32)$



Classical Sequence-to-Sequence Model

- Pass the last hidden state of the encoding stage
- Decoder uses this last hidden state to do the prediction





Word Representation

- **One-Hot Encoding**

- Representing each word as a vector that has as many values in it
- Each column in a vector is one possible word in a vocabulary
- Problem
 - In large vocabularies, these vectors can get very long
 - Contain all 0's except for one value
 - Sparse representation

*~100k columns, only one **1** in each vector*

the →	1	0	0	0	...	0	0	0
-------	---	---	---	---	-----	---	---	---

good →	0	0	1	0	...	0	0	0
--------	---	---	---	---	-----	---	---	---

movie →	0	0	0	0	...	1	0	0
---------	---	---	---	---	-----	---	---	---



Word Representation

- **Word Embedding**

- Map the word index to a continuous word embedding through a look-up table *~300 columns*
- Popular pre-trained word embeddings
 - Word2Vec, GloVe

the →	0.2	0.4	-0.1
good →	0.7	-0.5	0.3
movie →	0.1	0.2	0.6



Positional Encoding (PE)

- **Positional encoding (PE)**

- Information to each word about its position in the sentence
- **Unique** encoding for each word's position in a sentence
- Distance between any two positions is **consistent** across sentences with different lengths
- Encode words by using **sin()**, **cos()** with different frequencies
- **Deterministic** and **generalize** to longer sentences

$$\vec{p}_t^{(i)} = f(t)^{(i)} := \begin{cases} \sin(\omega_k \cdot t), & \text{if } i = 2k \\ \cos(\omega_k \cdot t), & \text{if } i = 2k + 1 \end{cases} \quad \omega_k = \frac{1}{10000^{2k/d}}$$



Positional Encoding (PE)

- Arguments
 - L: maximum # of possible positions
 - d_{model} : dimension of the embeddings
 - n: can be set to any value
 - k: position
 - i: dimension

$$\text{pos}(k) = \begin{bmatrix} \sin \omega_1 \cdot k \\ \cos \omega_1 \cdot k \\ \sin \omega_2 \cdot k \\ \cos \omega_2 \cdot k \\ \vdots \\ \sin \omega_{d_{model}/2} \cdot k \\ \cos \omega_{d_{model}/2} \cdot k \end{bmatrix}_{d_{model}}, \text{ where } \omega_i = 10000^{\frac{2i}{d_{model}}}$$

- For each $k = 0$ to $L - 1$
 - For each $i = 0$ to $\frac{d_{model}}{2}$
 - $PE_{(k,2i)} = \sin\left(\frac{k}{2i}\right) \frac{1}{n^{d_{model}}}$
 - $PE_{(k,2i+1)} = \cos\left(\frac{k}{2i}\right) \frac{1}{n^{d_{model}}}$

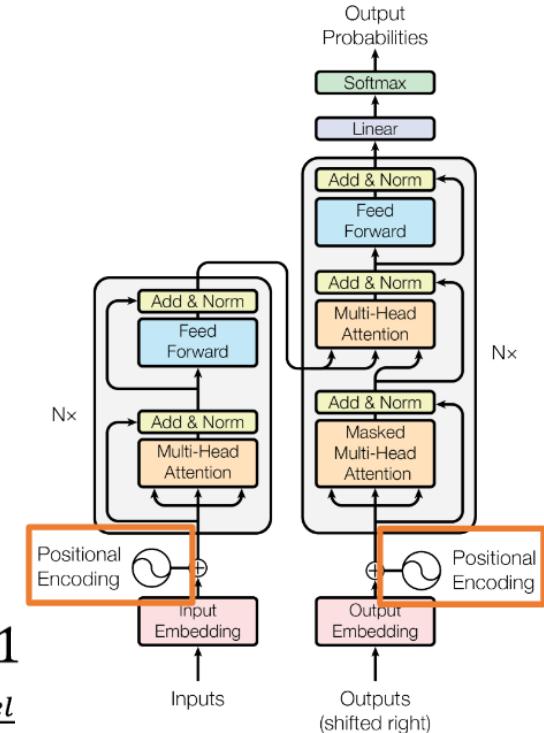
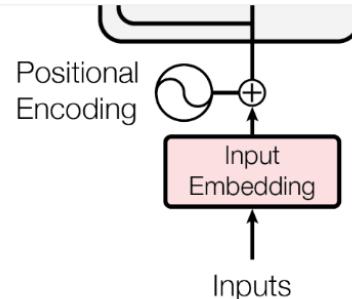


Figure 1: The Transformer - model architecture.

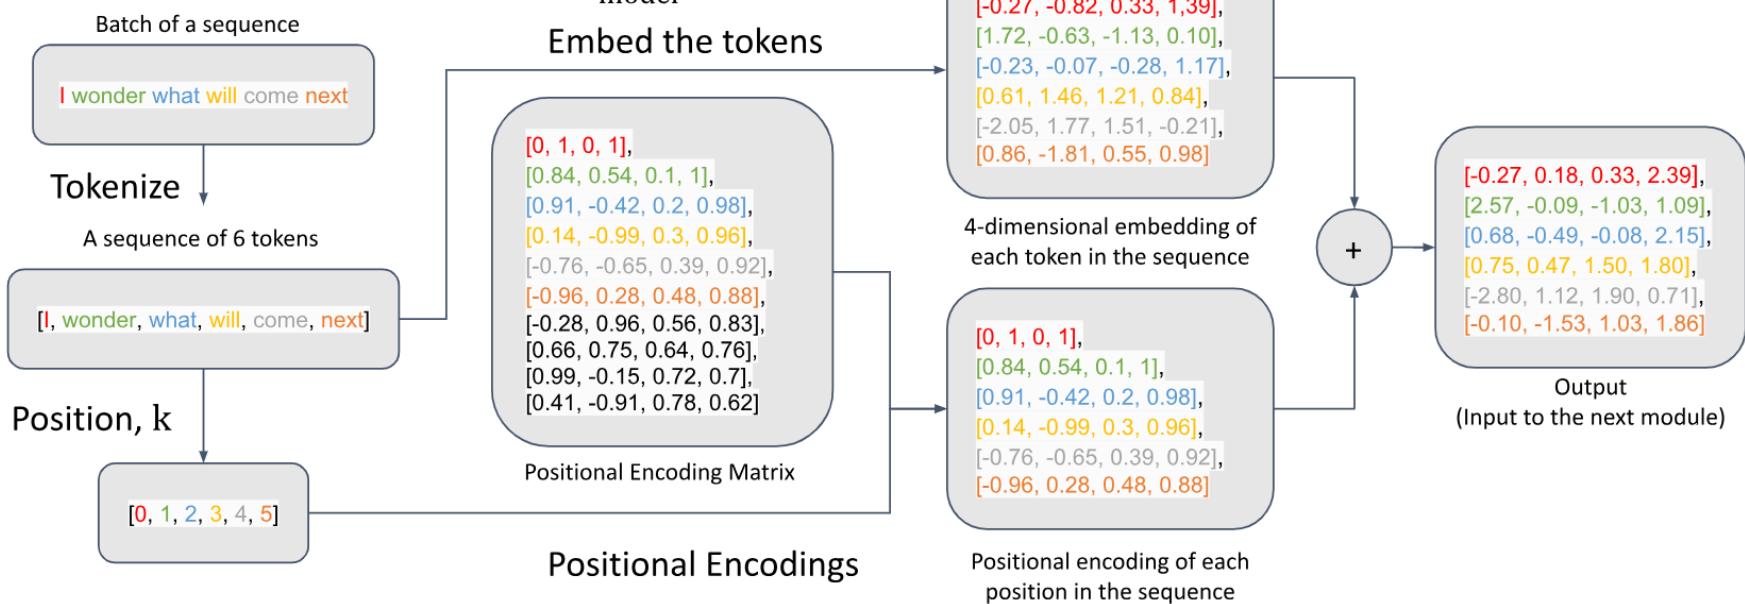


Positional Encoding (PE)



- Arguments

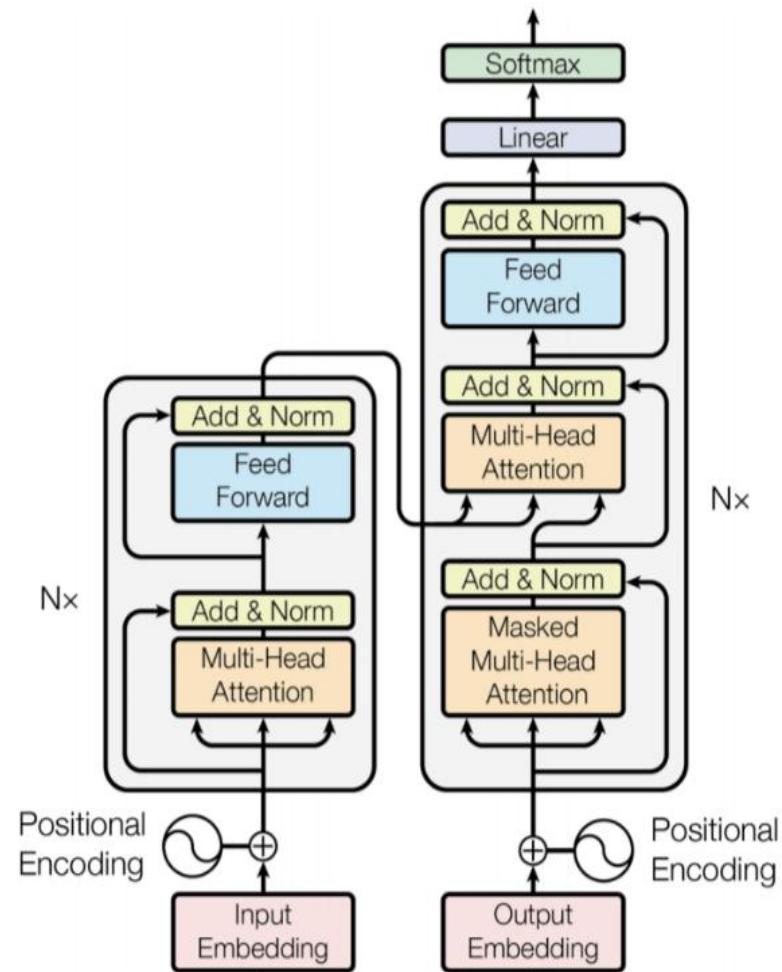
- $L=10, n=100, d_{\text{model}}=4$





Transformer Model Architecture

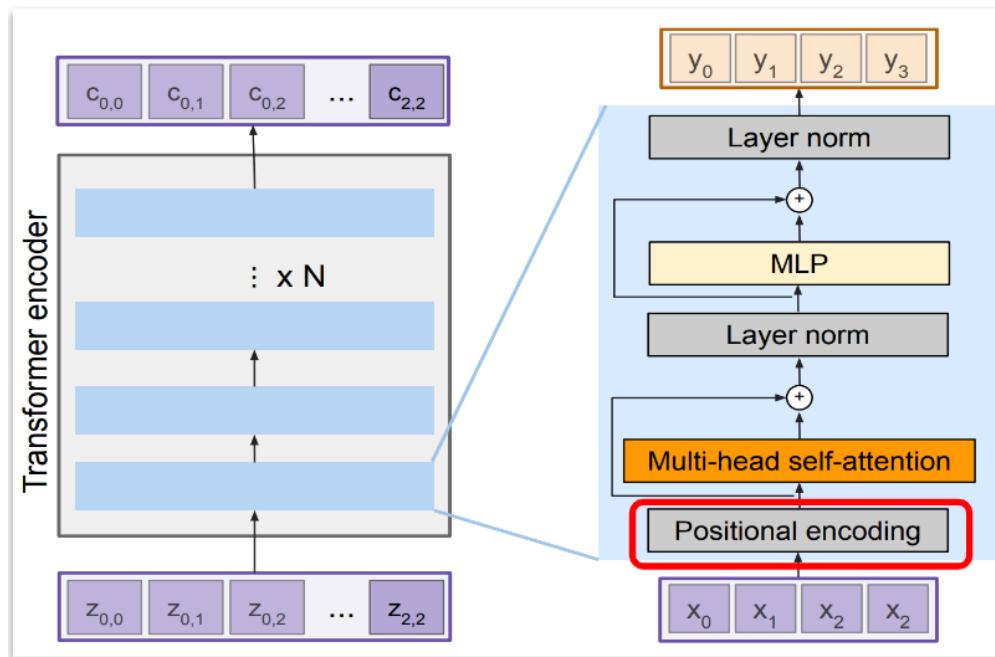
- Each **encoder** block has two sub-layers
 - Multi-head self-attention
 - A position-wise fully connected feed-forward
- Each **decoder** block has an additional third sub-layer
 - The third is a masked multi-head attention over the output of the encoder stack
- A **residual connection** is added around each of the two sub-layers
- The decoder yields the output sequence of symbols one element at a time





Transformer Model Architecture

- Transformer encoder



img src: [Stanford CS 231n](#)

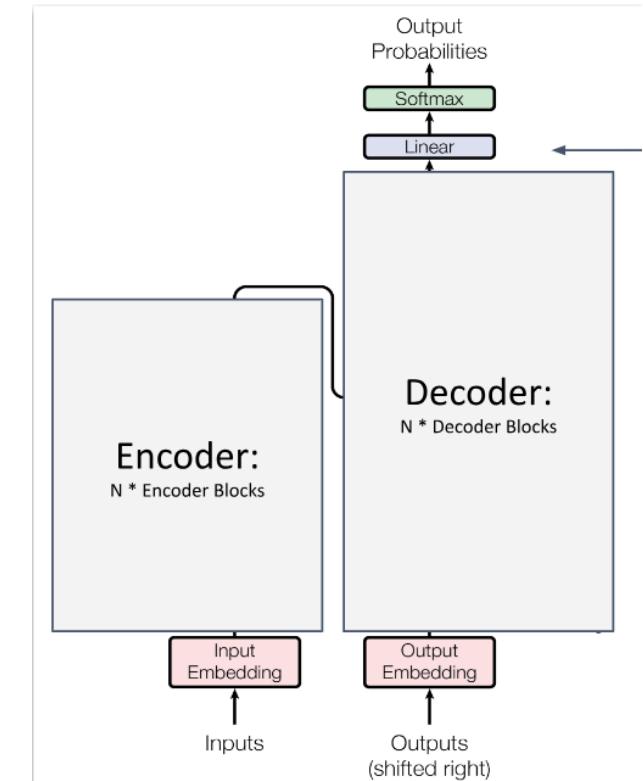
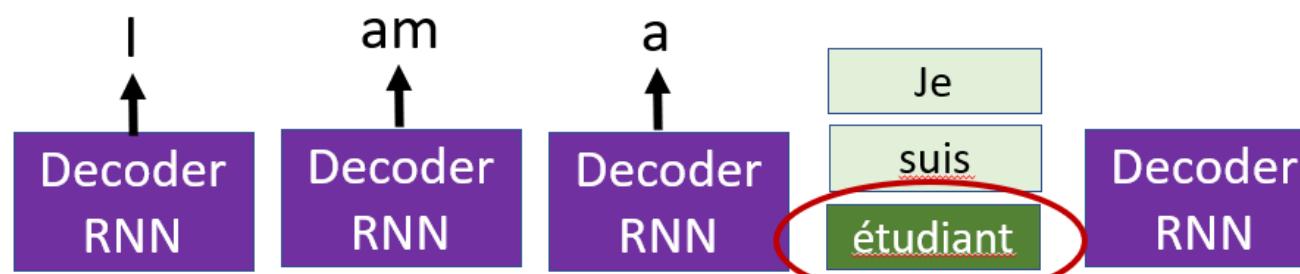


Figure 1: The Transformer - model architecture.



Bottleneck of Sequence-to-Sequence Model

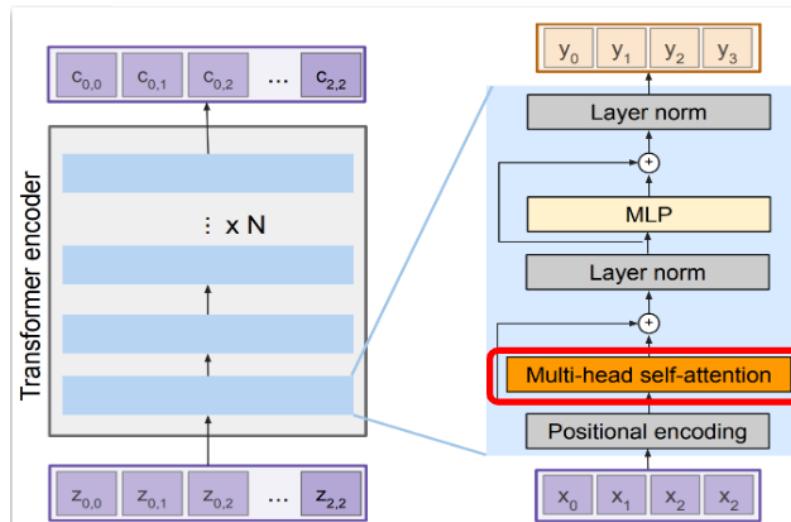
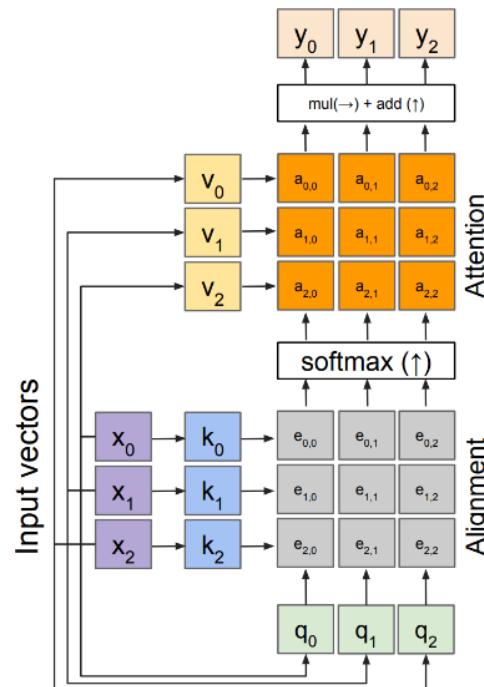
- It is challenging for the model to deal with long sentences
- **Attention**
 - The encoder passes all the hidden states to the decoder
 - The attention enables the decoder to focus on the word before it generates the English translation
 - This ability amplifies the signal from the relevant part of the input sentence





Transformer Model Architecture

- Encoder – Multi-head self-attention



- $Q = \mathbf{X} \cdot \mathbf{W}^Q$
- $K = \mathbf{X} \cdot \mathbf{W}^K$
- $V = \mathbf{X} \cdot \mathbf{W}^V$

Q: query
K: key
V: value

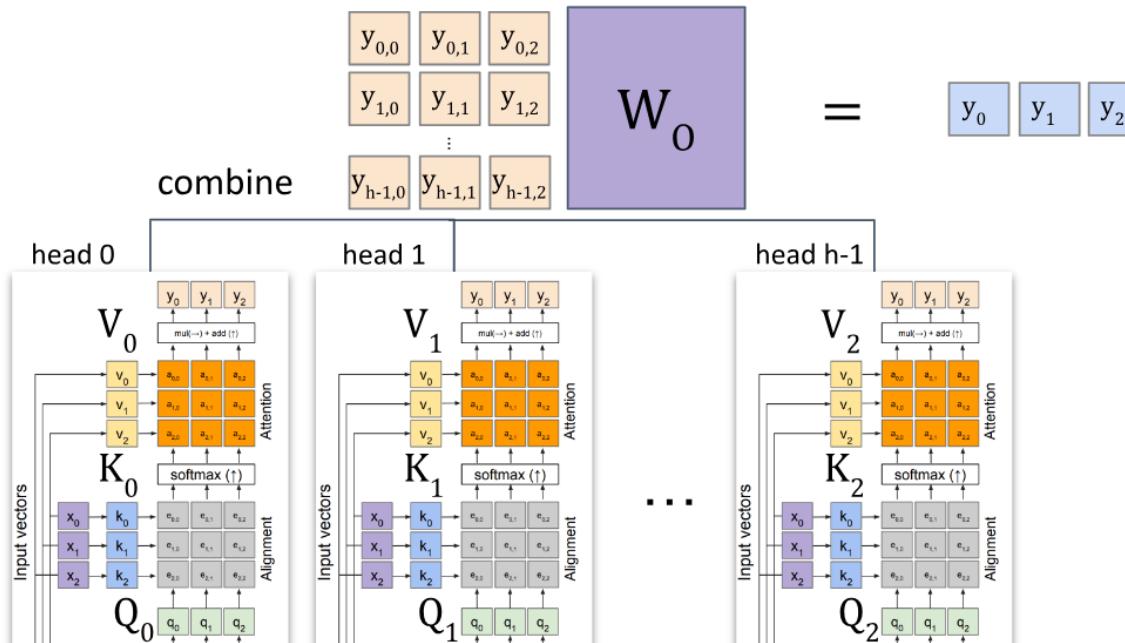
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



Transformer Model Architecture

- Encoder – Multi-head self-attention

$$\begin{aligned} Q_i &= \mathbf{X} \cdot W_i^Q \\ K_i &= \mathbf{X} \cdot W_i^K \\ V_i &= \mathbf{X} \cdot W_i^V \end{aligned}$$



Multi-head Self-attention

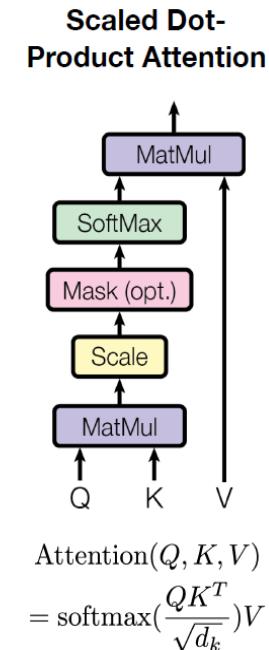
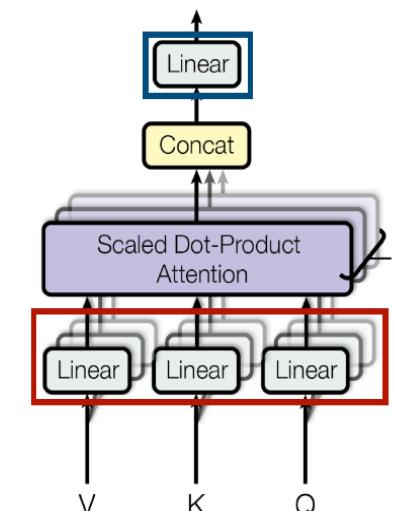


Multi-Head Self-Attention (MHSA)

- **Project Q, K, and V with h different learned linear projections**
- Perform the **scaled dot-product attention** function on each of Q, K, V **in parallel**
- **Concatenate** the output values
- **Project** the output values again, resulting in the final values

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

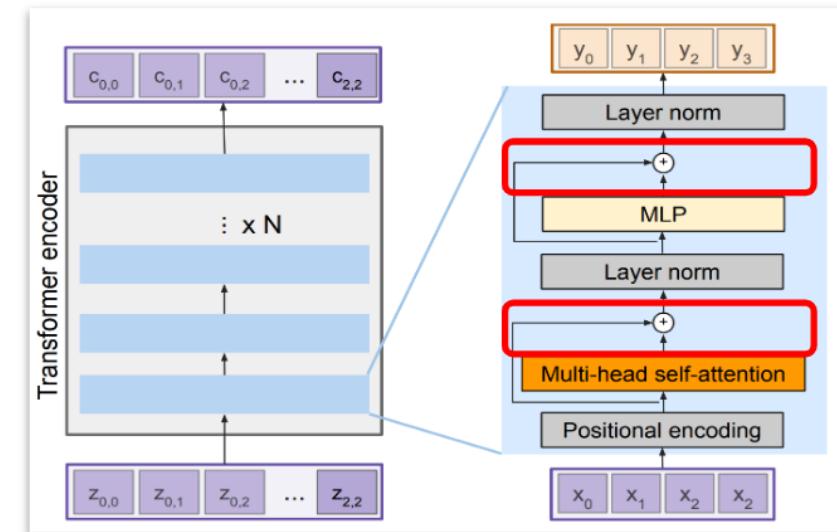




Transformer Model Architecture

- Encoder – Residual connection
 - Address exploding/vanishing gradients
 - Equation
 - $F(x)$: output of previous layer
 - x : input of previous layer

$$y = x + F(x)$$





Transformer Model Architecture

- Encoder – Layer normalization
 - For each input vector $x = (x_1 \dots x_m)$
 - Calculate Mean, variance
 - Normalize the vector x
 - Small ϵ -> avoid dividing 0
 - Scale and shift

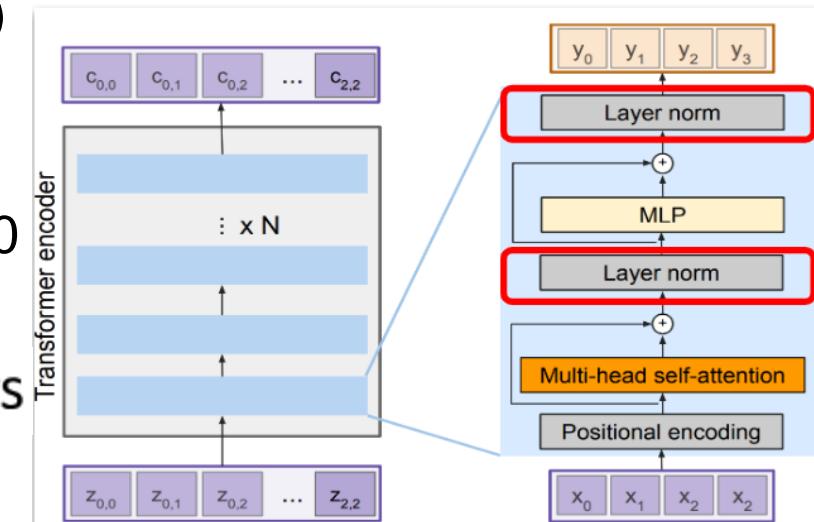
γ, β are learnable parameters

$$\mu = \frac{\sum_i x_i}{m}$$

$$\hat{x}_i = \frac{(x_i - \mu)}{\sqrt{\sigma^2 + \epsilon}}$$

$$\sigma^2 = \frac{\sum_i (x_i - \mu)^2}{m}$$

$$y_i = \gamma \hat{x}_i + \beta$$





Transformer Model Architecture

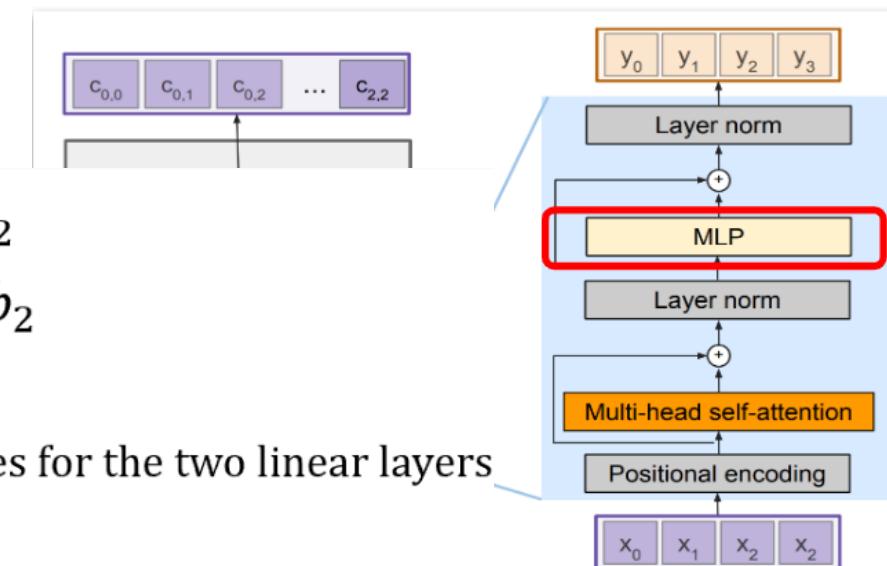
- Encoder – Feed Forward Network (FFN)
 - Contains 2 linear transformations and 1 ReLU

$$\begin{aligned} FFN(x) &= \text{ReLU}(W_1x + b_1)W_2 + b_2 \\ &= \max(0, W_1x + b_1)W_2 + b_2 \end{aligned}$$

x : output of the previous layer

$W_{1(d_{model} \times d_{ff})}, W_{2(d_{ff} \times d_{model})}$: weight matrices for the two linear layers

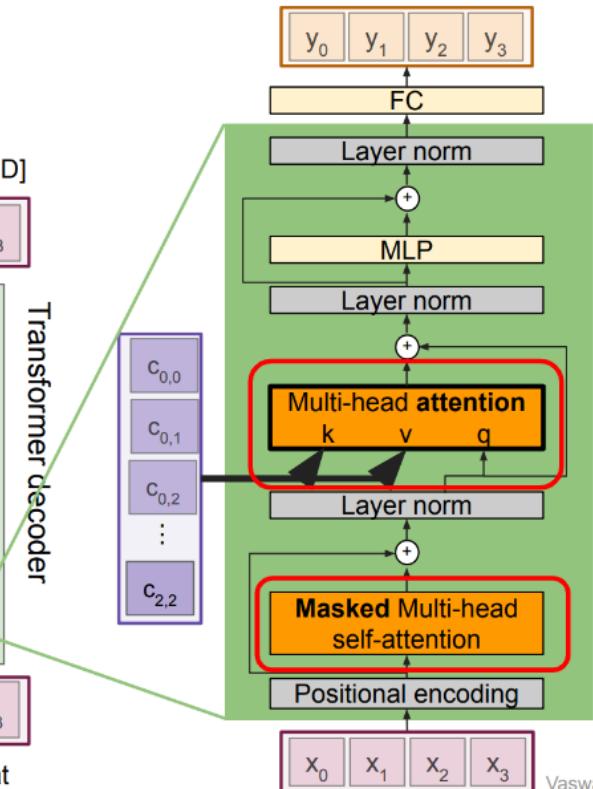
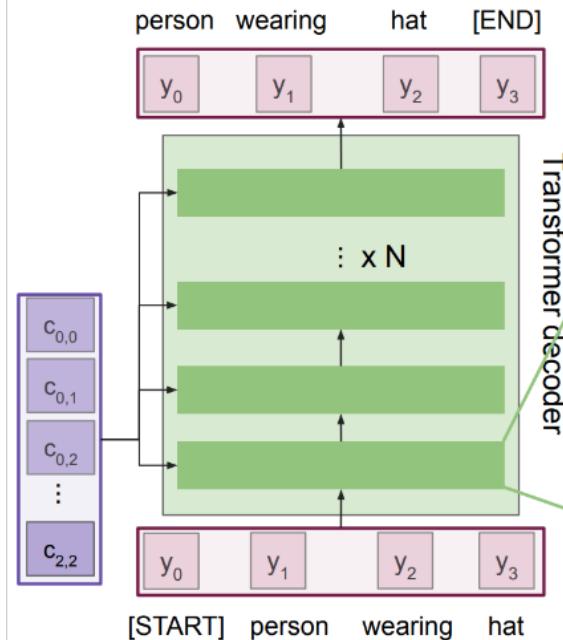
b_1, b_2 : biases for the two linear layers





Transformer Model Architecture

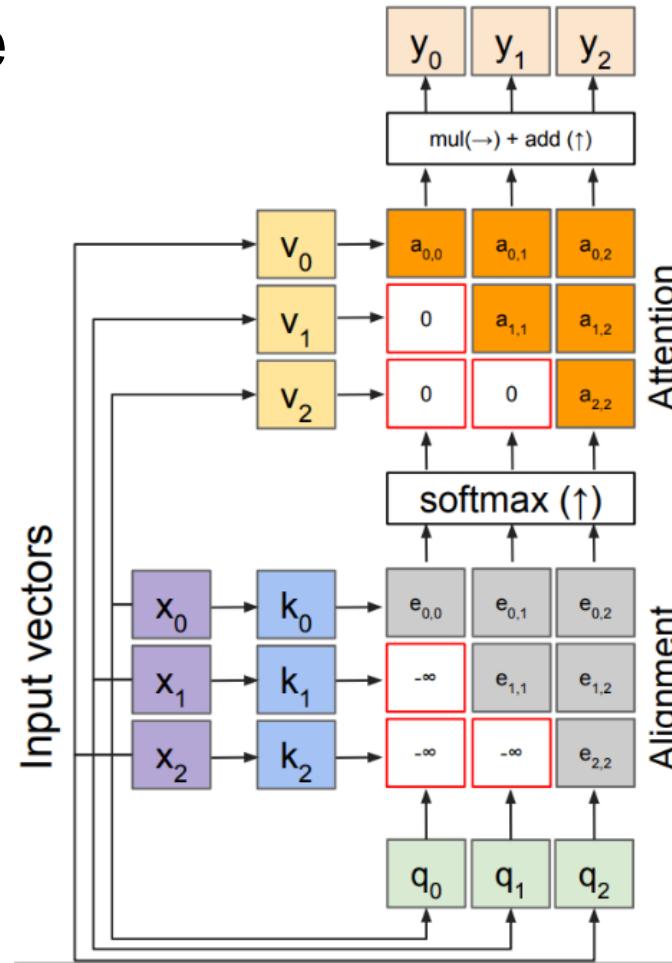
The Transformer Decoder block





Transformer Model Architecture

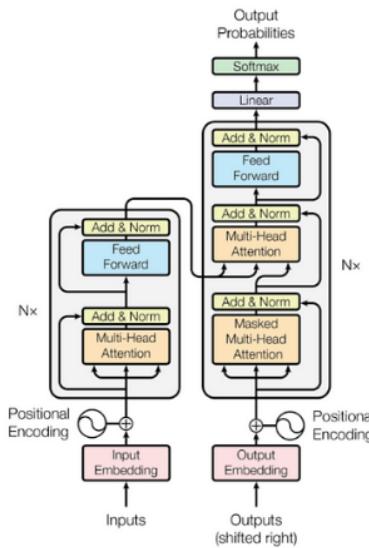
- Decoder – Mask self-attention
 - The output of a certain position can only depend on the words on the previous positions
 - Set alignment scores of successive position to negative infinity





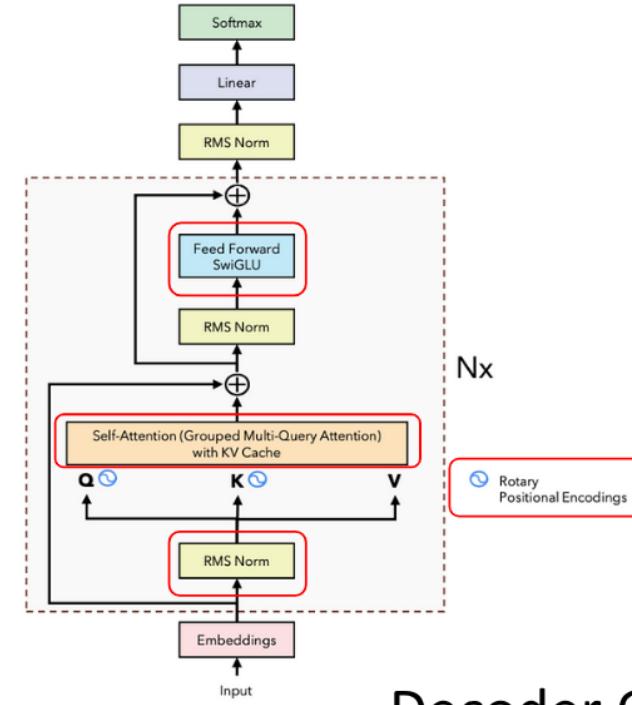
LLaMA Model Architecture

Transformer vs LLaMA



Transformer

("Attention is all you need")

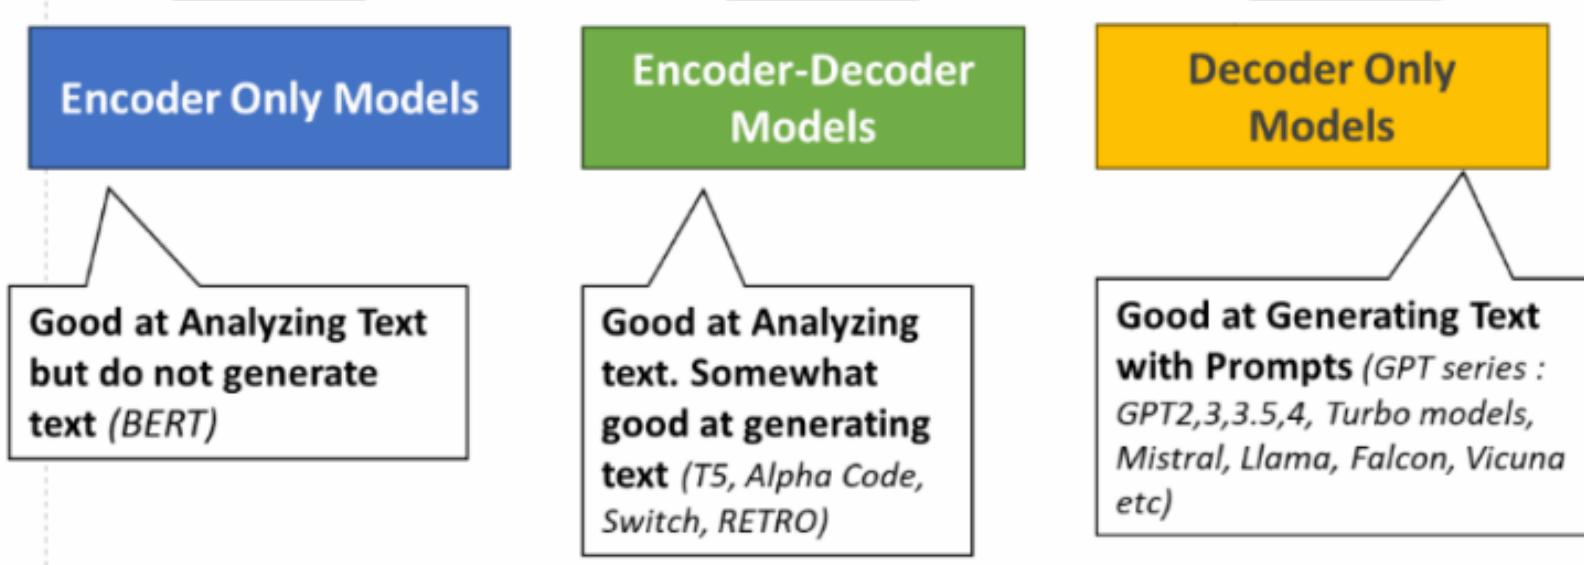


Decoder Only!



LLaMA Model Architecture

- Why LLaMA is decoder-only model?





LLaMA Model Architecture

- Difference between Transformer and LLaMA
 - Decoder-only model
 - Pre-Norm (root mean square (RMS) norm)
 - Rotary positional embedding (RoPE)
 - KV cache
 - Grouped multi query attention
 - SwiGLU activation function rather than ReLU in FFN



LLaMA Model Architecture

- Rotary positional embedding (RoPE)
 - Combine absolute and relative encoding
 - Absolute positional embedding
 - Assigns a unique vector to each position and doesn't scale well to capture relative position
 - Relative embeddings
 - Focuses on the distance between tokens
 - Enhance the model's understanding of token relationship
 - Rotational mechanism
 - Each position in the sequence is represented by a rotation in the embedding space

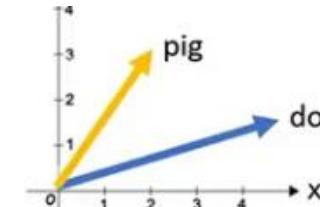


LLaMA Model Architecture

- Rotary positional embedding (RoPE)
 - RoPE applies a rotation to the word vector
 - The equation incorporates a rotation matrix that rotates a vector by an angle of $M\theta$, where M is the absolute position in the sentence.
 - This rotation is applied to the query and key vectors in the self-attention

$$f_{\{q,k\}}(\mathbf{x}_m, m) = \begin{pmatrix} \cos m\theta & -\sin m\theta \\ \sin m\theta & \cos m\theta \end{pmatrix} \begin{pmatrix} W_{\{q,k\}}^{(11)} & W_{\{q,k\}}^{(12)} \\ W_{\{q,k\}}^{(21)} & W_{\{q,k\}}^{(22)} \end{pmatrix} \begin{pmatrix} \mathbf{x}_m^{(1)} \\ \mathbf{x}_m^{(2)} \end{pmatrix}$$

The **pig** chased the **dog**





LLaMA Model Architecture

- Root Mean Square (RMS) normalization
 - Calculate the RMS of the input vectors rather than the mean and variance
 - Efficient normalization
 - No subtracting mean before squaring
 - No shifting in implementation

$$\text{RMS}(x) = \sqrt{\frac{1}{m} \sum_i x_i^2}$$

$$\hat{x}_i = \frac{x_i}{\text{RMS}(x)+\varepsilon}$$

$$y_i = \gamma \hat{x}_i$$

$$\mu = \frac{\sum_i x_i}{m}$$

$$\sigma^2 = \frac{\sum_i (x_i - \mu)^2}{m}$$

$$\hat{x}_i = \frac{(x_i - \mu)}{\sqrt{\sigma^2 + \varepsilon}}$$

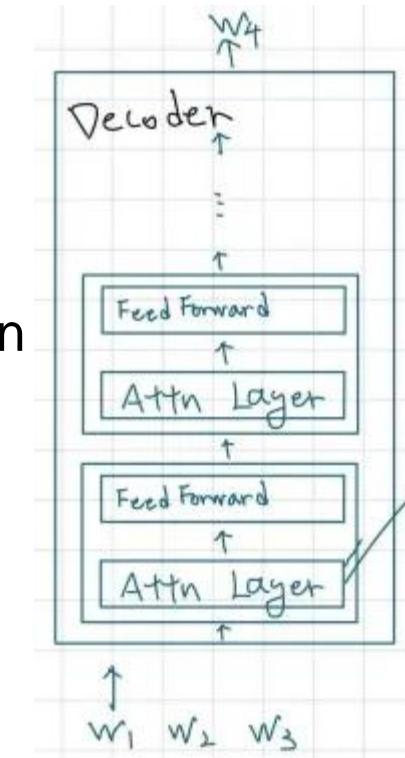
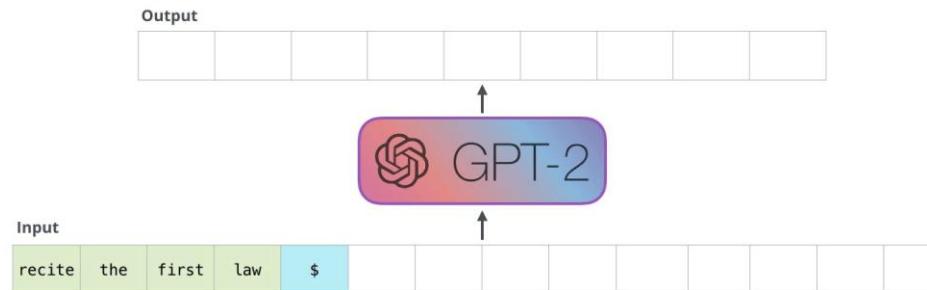
$$y_i = \gamma \hat{x}_i + \beta$$

Layer norm



LLaMA Model Architecture

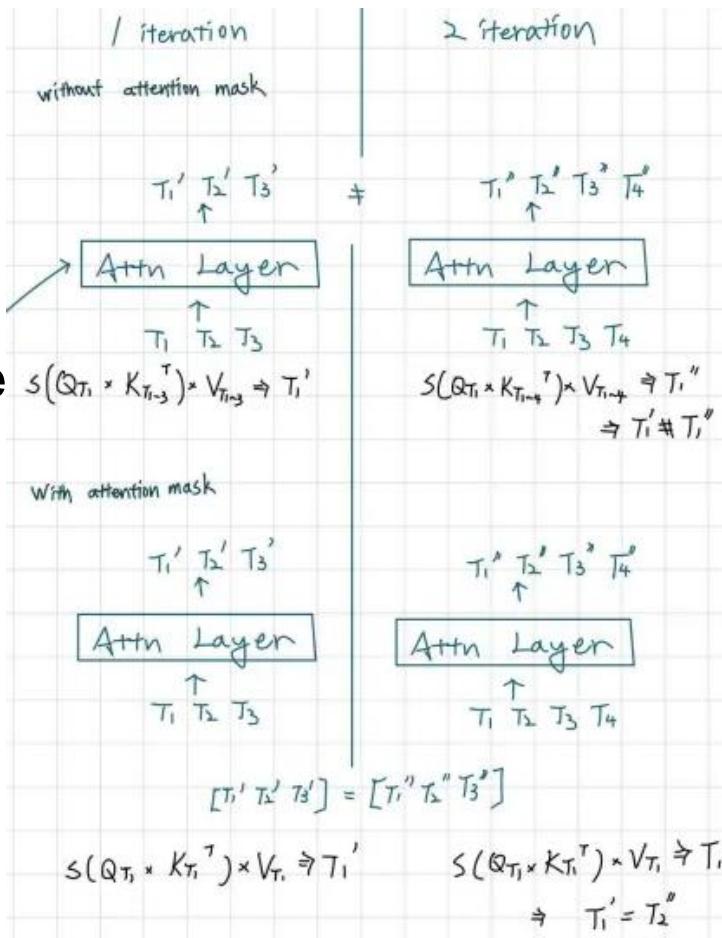
- LLM decoder
 - The decoder works in an auto-regressive fashion
 - Given an input, the model predict the next token
 - Taking the combined input in the next step





LLaMA Model Architecture

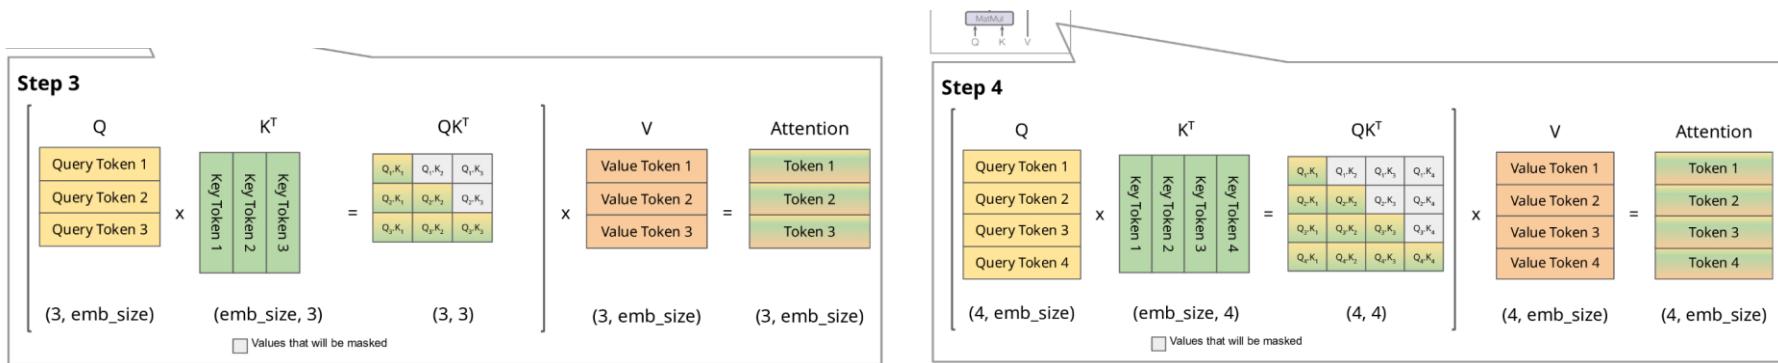
- Sequence mask (in decoder)
 - The decoder cannot see the message in the coming future
 - Use the mask to enable the decoder to only rely on the previous outputs to do the inference → training the decoder





LLaMA Model Architecture

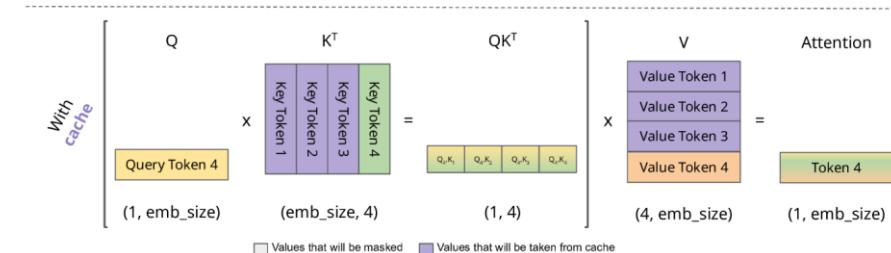
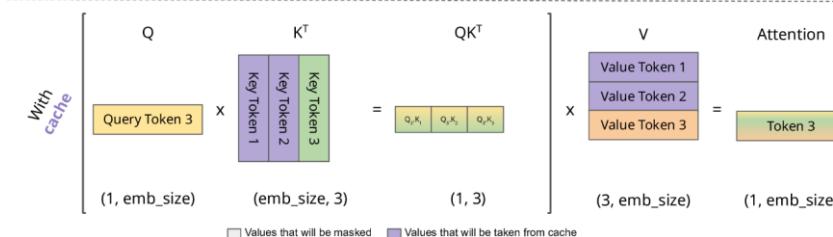
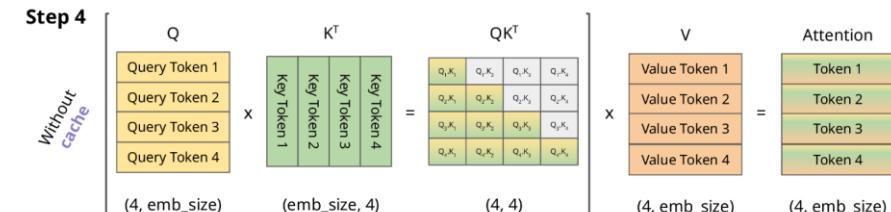
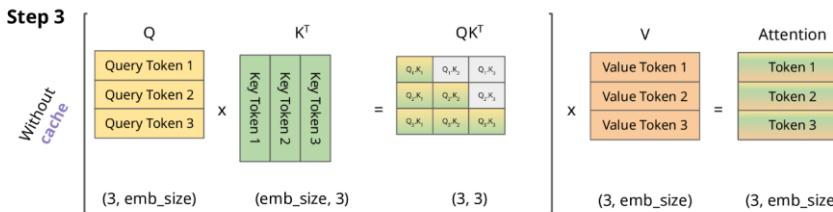
- The scaled dot-product attention
 - The attention of a token only depends on its preceding tokens
 - At each generation step we are recalculating the same previous token attention, when we actually just want to calculate the attention for the new token





LLaMA Model Architecture

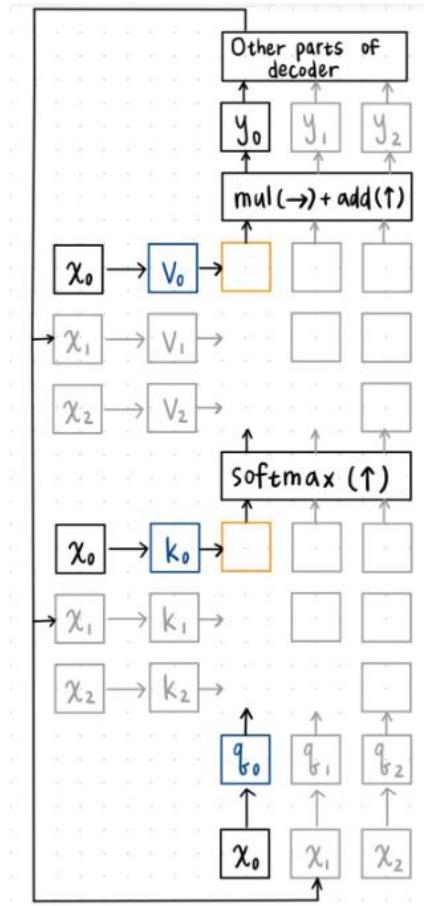
- Key-Value (KV) cache
 - By caching the previous Keys and Values, we can focus on only calculating the attention for the new token.





LLaMA Model Architecture

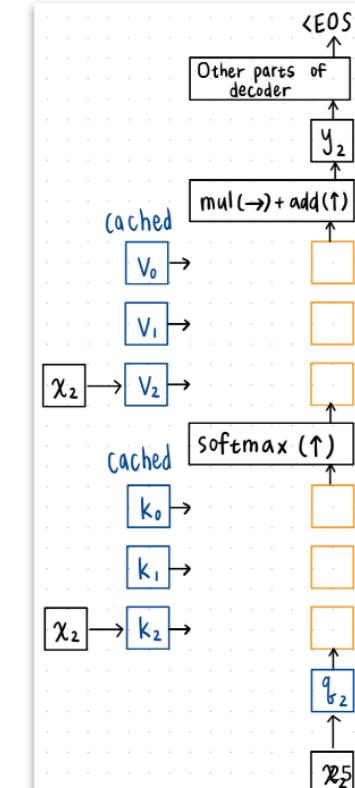
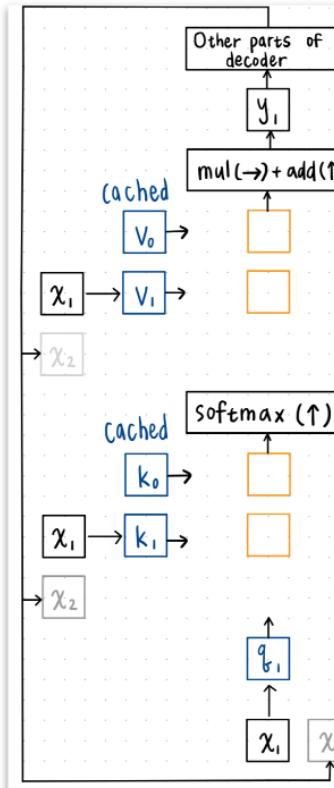
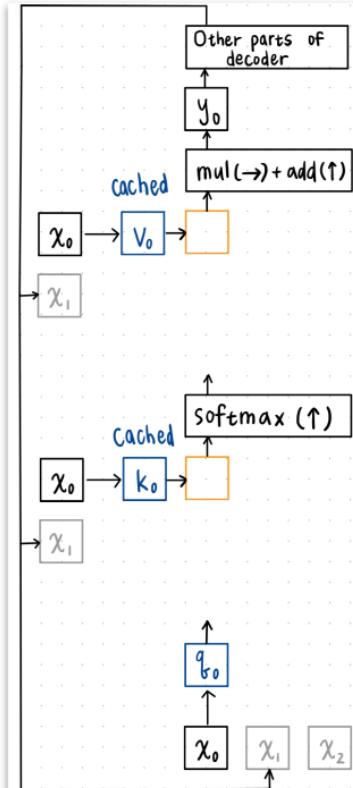
- Key-Value (KV) cache
 - LLM models can generate only one token at a time
 - Each new prediction is dependent on the previous context
 - To predict token number 1000 in the generation, you need information from the previous 999 tokens
 - Optimize the sequential generation process by storing previous calculations to reuse in subsequent tokens, so they don't need to be computed again.





LLaMA Model Architecture

- KV cache





LLaMA Model Architecture

- Key-Value (KV) cache
 - The matrices obtained with KV caching are way smaller, which leads to faster matrix multiplications
 - The downside of the KV cache is
 - When the length of sequences is becoming long
 - Needs the large memory to cache the Key and Value states

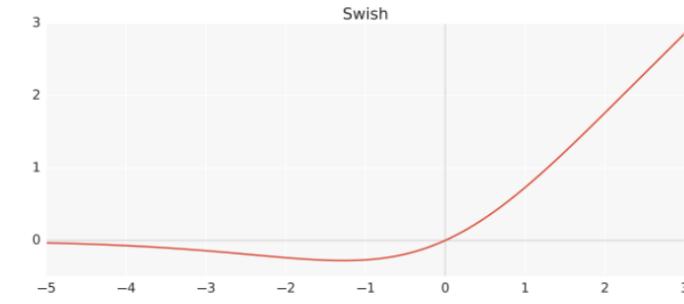


LLaMA Model Architecture

- SwiGLU (Swish and Gated Linear Unit)
 - LLU such as PALM and LLAMA use SwiGLU in FFN rather than the usual ReLU -> SwiGLU tackles minus value better than ReLU

$$\text{FFN}_{\text{SwiGLU}}(x, W, V, W_2) = (\text{Swish}_1(xW) \otimes xV)W_2,$$

$$\text{where } \text{Swish}_\beta(x) = x \cdot \text{sigmoid}(\beta x) = \frac{x}{1+e^{-\beta x}}$$



img src: [medium@neuralnets](#)



Takeaway Questions

- What's problem the “Attention” aiming to solve?
 - (A) Gradient vanishing
 - (B) Message passing in the long sequence of data
 - (C) Over-fitting
- What are benefits of the “Transformer” ?
 - (A) Large hidden layer
 - (B) The amount of computation is small
 - (C) More data parallelism



Takeaway Questions

- How does the “self-attention” help the encoder?
 - (A) Looking at other words in the input sentence
 - (B) Memorizing the more messages within a network
 - (C) Focus on a specific word