



# Accelerator Architectures for Machine Learning (AAML)

## Lecture 3: Quantization

Tsung Tai Yeh  
Department of Computer Science  
National Yang-Ming Chiao Tung University



# Acknowledgements and Disclaimer

- Slides was developed in the reference with  
Joel Emer, Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, ISCA 2019 tutorial  
Efficient Processing of Deep Neural Network, Vivienne Sze, Yu-Hsin Chen,  
Tien-Ju Yang, Joel Emer, Morgan and Claypool Publisher, 2020  
Yakun Sophia Shao, EE290-2: Hardware for Machine Learning, UC  
Berkeley, 2020  
CS231n Convolutional Neural Networks for Visual Recognition, Stanford  
University, 2020
- 6.5940, TinyML and Efficient Deep Learning Computing, MIT
- NVIDIA, Precision and performance: Floating point and IEEE 754  
Compliance for NVIDIA GPUs, TB-06711-001\_v8.0, 2017



# Outline

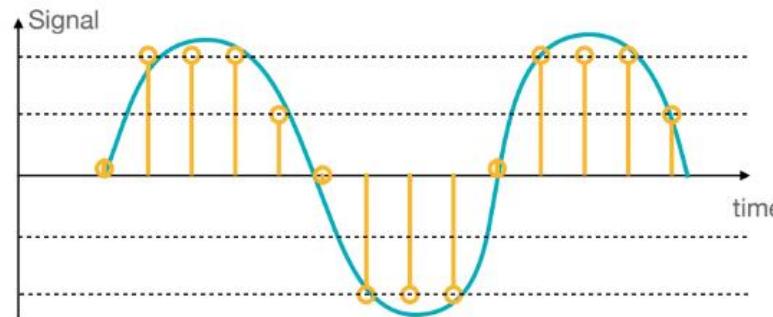
- K-Means-based Quantization
- Linear Quantization
- Binary and Ternary Quantization



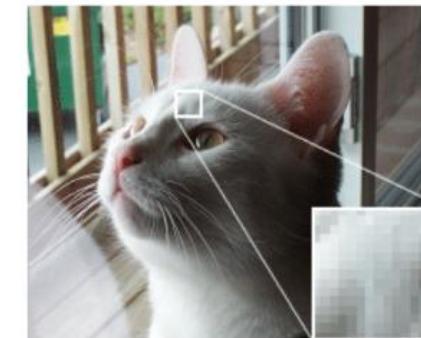
# What is Quantization ?

- **Quantization**
  - A process that reduces the precision of a digital signal by converting high-precision data into a lower-precision format

— Continuous Signal    —○— Quantized Signal



Original Image



16-Color Image

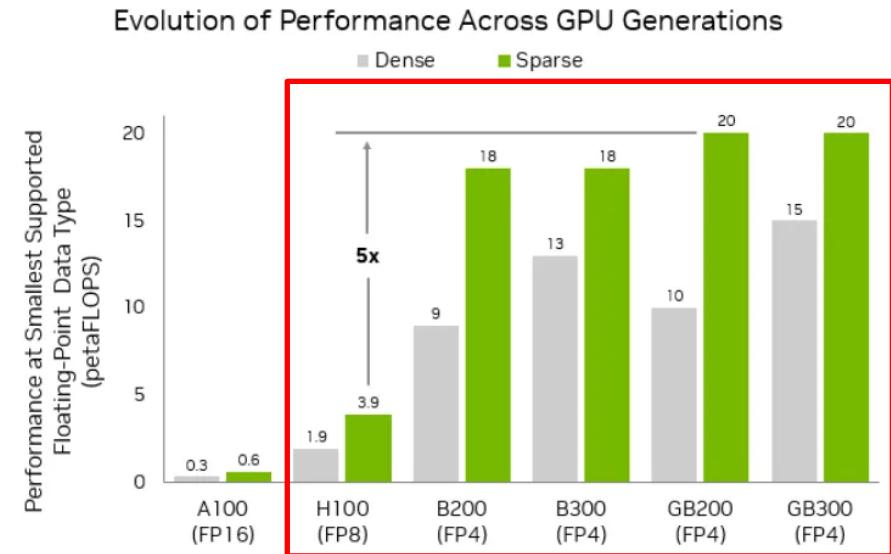


Images are in the public domain.



# Benefits of Quantization

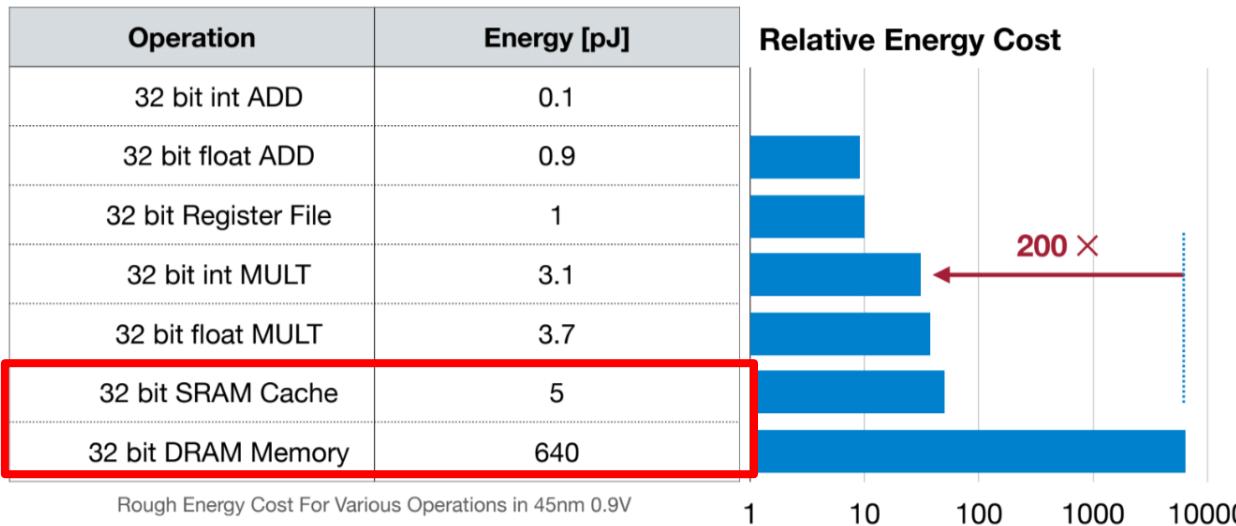
- **Reduced memory burden**
  - Reduce pressure on memory bandwidth which can improve output token throughput
- **Simplified compute operations**
  - Improve overall end-to-end latency performance as a result of simplified attention layer computations





# Memory is Expensive !!

- Data movement -> Move memory reference -> More energy



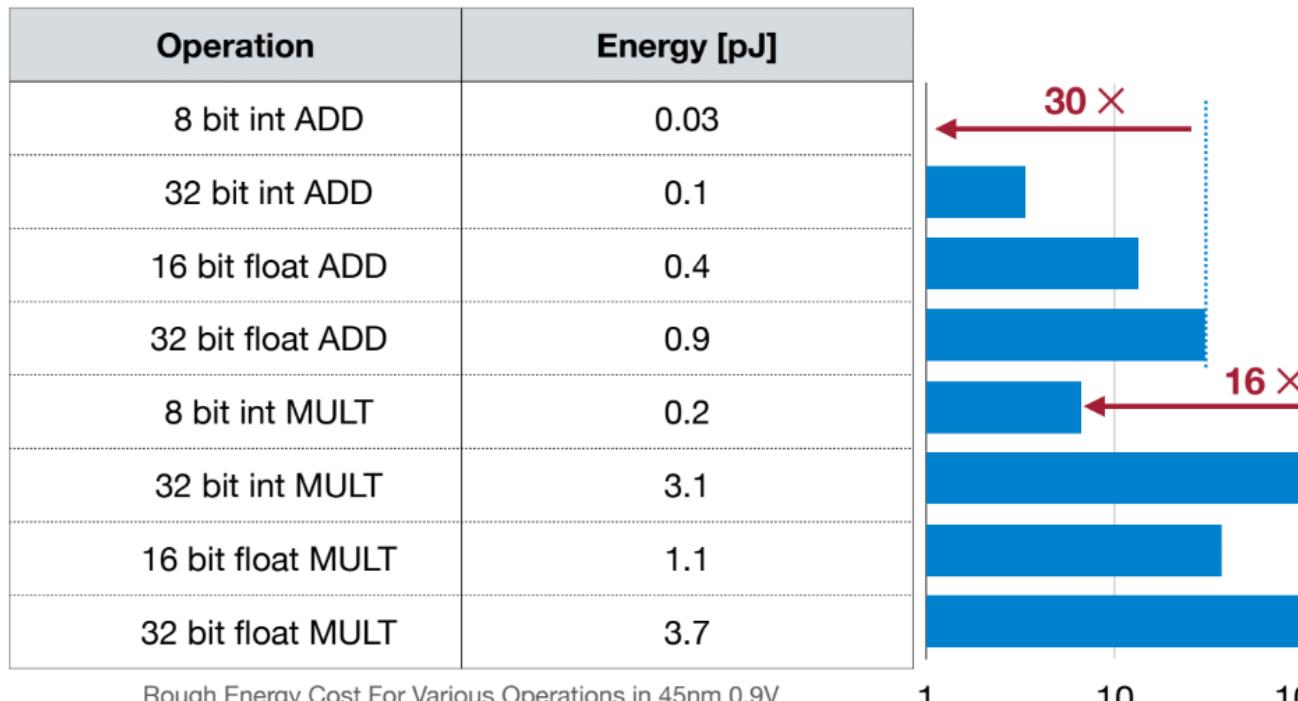
1  = 200 ✕+

This image is in the public domain



# Low Bit-Width Operations are Cheap

- Less Bit-Width -> Less energy





# Energy and Area Cost

Could we make the deep learning efficient by lowering the precision of data ?

| Operation           | Energy (pJ) | Area(um <sup>2</sup> ) |
|---------------------|-------------|------------------------|
| 8b Add              | 0.03        | 36                     |
| 16b Add             | 0.05        | 67                     |
| 32b Add             | 0.1         | 137                    |
| 16b FP Add          | 0.4         | 1360                   |
| 32b FP Add          | 0.9         | 4184                   |
| 16b FP Mult         | 1.1         | 1640                   |
| 32b FP Mult         | 3.7         | 7700                   |
| 32b SRAM Read (8KB) | 5           |                        |
| 32b DRAM Read       | 640         |                        |

173X

4.7X



# Numeric Data Types

- Fixed-point number



Integer . Fraction

"Decimal" Point



$$-2^3 + 2^2 + 2^1 + 2^0 + 2^{-1} + 2^{-2} + 2^{-3} + 2^{-4} = 3.0625$$

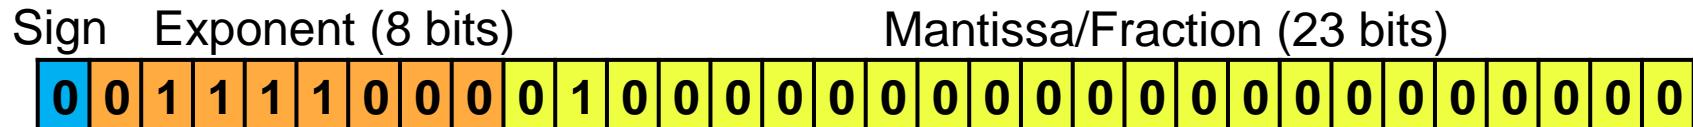


$$(-2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0) \times 2^{-4} = 49 \times 0.0625 = 3.0625$$



# IEEE 765 Single Precision Float Point

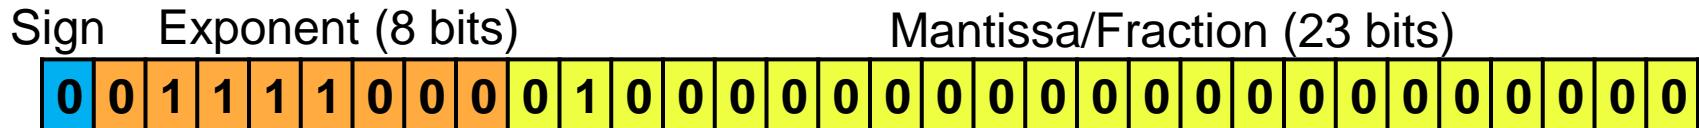
- **Sign** determines the sign of the number
- **Exponent** (8 bit) represent -127 (all 0s) and +128 (all 1s)
- **Significand** (23 fraction bits), total precision is 24 bits (23 + 1 implicit leading bit)  $\log_{10}(2^{24}) \approx 7.225$  digital bit



$$value = (-1)^{sign} \times 2^{(e-127)} \times \left(1 + \sum_{i=1}^{23} b_{(23-i)} 2^{-i}\right)$$



# IEEE 765 FP32



$$value = (-1)^{sign} \times 2^{(e-127)} \times (1 + \sum_{i=1}^{23} b_{(23-i)} 2^{-i})$$

$$\text{Sign} = b31 = 0; (-1)^0 = 1$$

$$e = 120; 2^{(120 - 127)} = 2^{-7}$$

$$1.b_{22}b_{21}\dots b_0 = \left(1 + \sum_{i=1}^{23} b_{(23-i)} 2^{-i}\right) = 1 + 2^{-2} = 1.25$$

$$\text{Value} = 1 \times 2^{-7} \times 1.25 = 0.009765625$$



# Numeric Data Type

- **Question:** What is the decimal “11.375” in FP32 format ?

11.375

$$\equiv 11 + 0.375$$

$$= (1011)_2 + (0.011)_2$$

$$= (1011.011)_2$$

$$= (1.011011)_2 \times 2^3$$

$$0.375 \times 2 = 0.750 = 0 + 0.750 \Rightarrow b_{-1} = 0$$

$$0.750 \times 2 = 1.500 = 1 + 0.500 \Rightarrow b_{-2} = 1$$

$$0.500 \times 2 = 1.000 = 1 + 0.000 \Rightarrow b_{-3} = 1$$

- The exponent is 3 and biased form

$$= (3 + 127) = 130 = 1000\ 0010$$

## Sign

## Exponent (8 bits)

## Mantissa/Fraction (23 bits)





# Floating-Point Number

- Exponent Width -> Range; Fraction Width-> Precision

IEEE 754 Single Precision 32-bit Float (IEEE FP32)



IEEE Half Precision 16-bit Float (IEEE FP16)



Brain Float (BF16)



Nvidia TensorFloat (TF32)



AMD 24-bit Float (AMD FP24)



Exponent  
(bits)

8

5

8

8

Fraction  
(bits)

23

10

7

10

Total  
(bits)

32

16

16

19

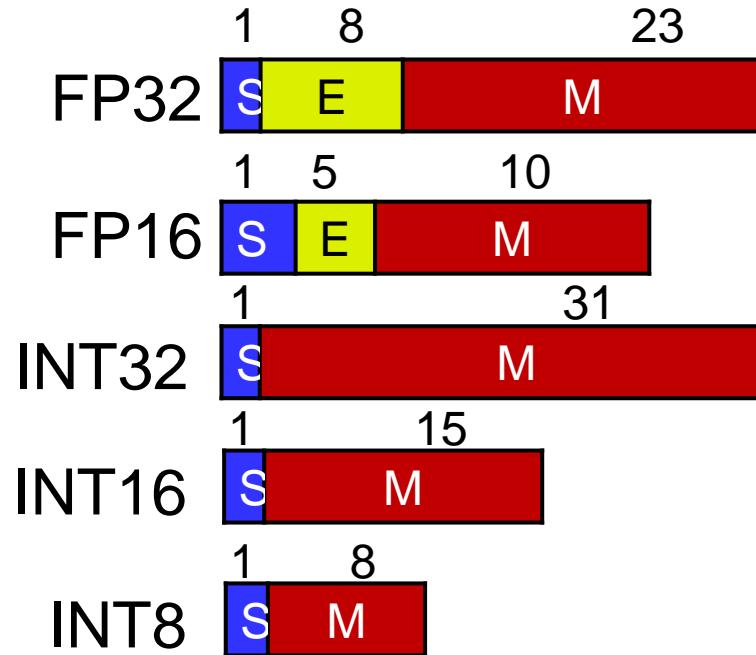
24

7

16



# Number Representation



## Range

1.2E-38 to 3.4E+38

6.1E-5 to 6.6E+4

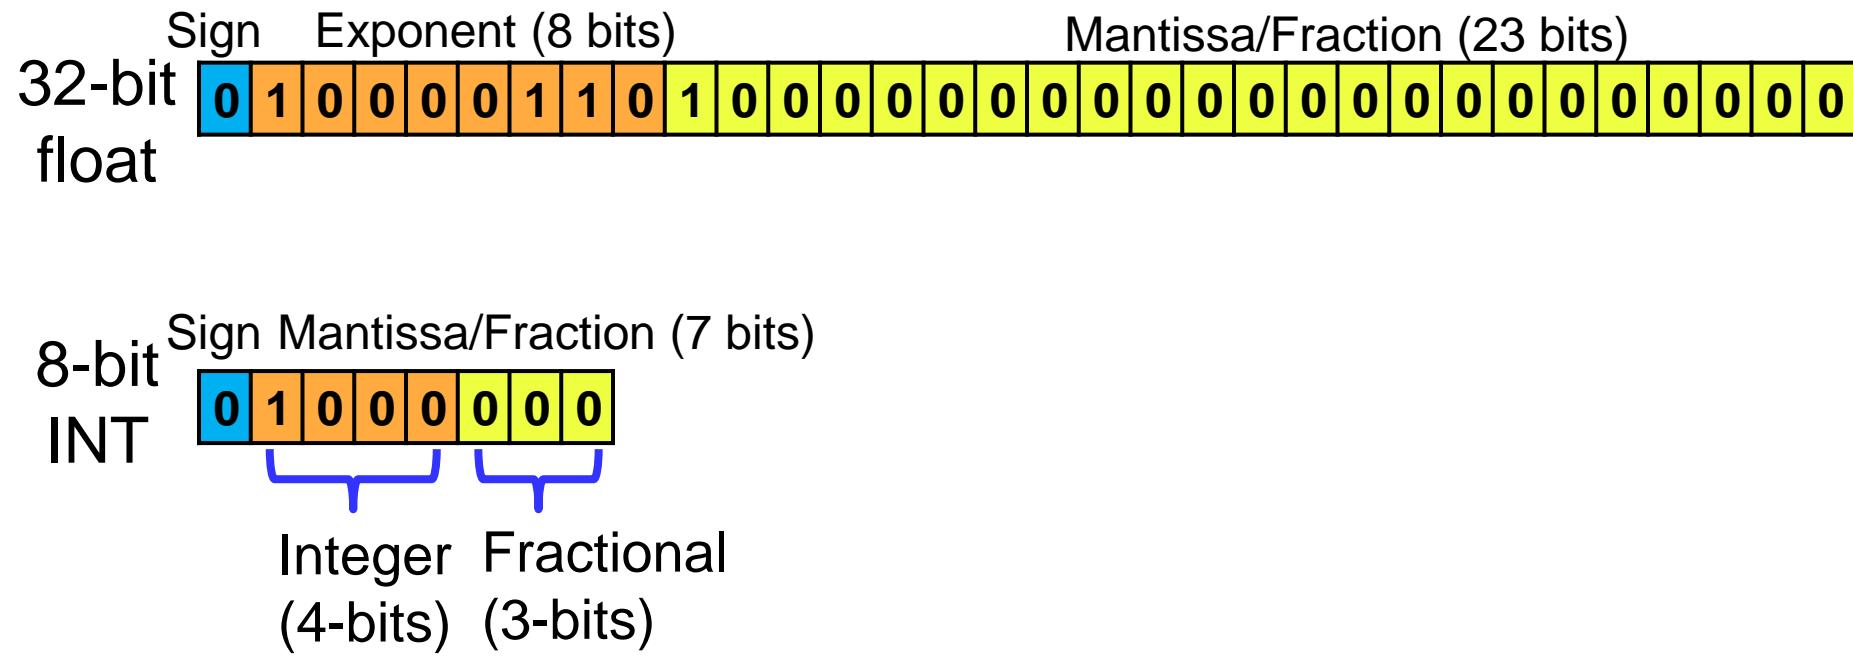
2147483648 to 2147483647

-32,768 to 32,767

-128 ~ 127



# Reduced Bit Width





# FP32 vs FP16 vs BF16

- **FP32 – single precision**
  - With 6-9 significant decimal digits precision
- **FP16 – half precision**
  - Uses in some neural network applications
  - With 4 significant decimal digits precision
- **BF16**
  - A truncated FP32
  - Allow for fast conversion to and from an FP32
  - With 3 significant decimal digits

(a) fp32: Single-precision IEEE Floating Point Format



Range:  $\sim 1e^{-38}$  to  $\sim 3e^{38}$

(b) fp16: Half-precision IEEE Floating Point Format



Range:  $\sim 5.96e^{-8}$  to  $65504$

(c) bfloat16: Brain Floating Point Format



Range:  $\sim 1e^{-38}$  to  $\sim 3e^{28}$

| Format | Bits | Exponent | Fraction |
|--------|------|----------|----------|
| FP32   | 32   | 8        | 23       |
| FP16   | 16   | 5        | 10       |
| BF16   | 16   | 8        | 7        |



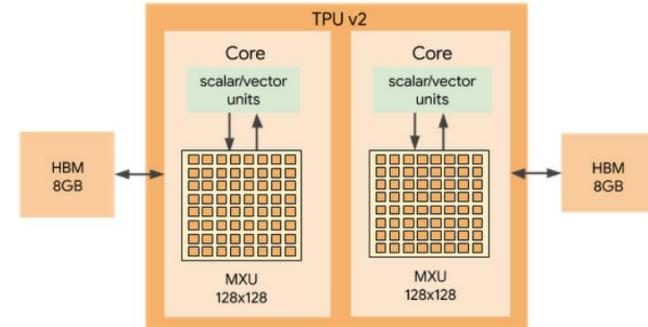
# Choosing bFloat16

- **Motivation**

- The physical size of a hardware multiplier scales with the square of the mantissa width
- Mantissa bit length – FP32-> 23 bits, FP16-> 10 bits, BF16:->7 bits

- **BF16**

- 8 X smaller than an FP32 multiplier
- Has the same exponent size as FP32
- No require special handling (loss scaling) in the FP16 conversion
- XLA compiler's automatic format conversion
- In side the MXU, multiplications are performed in BF16 format
- Accumulations are performed in full FP32 precision

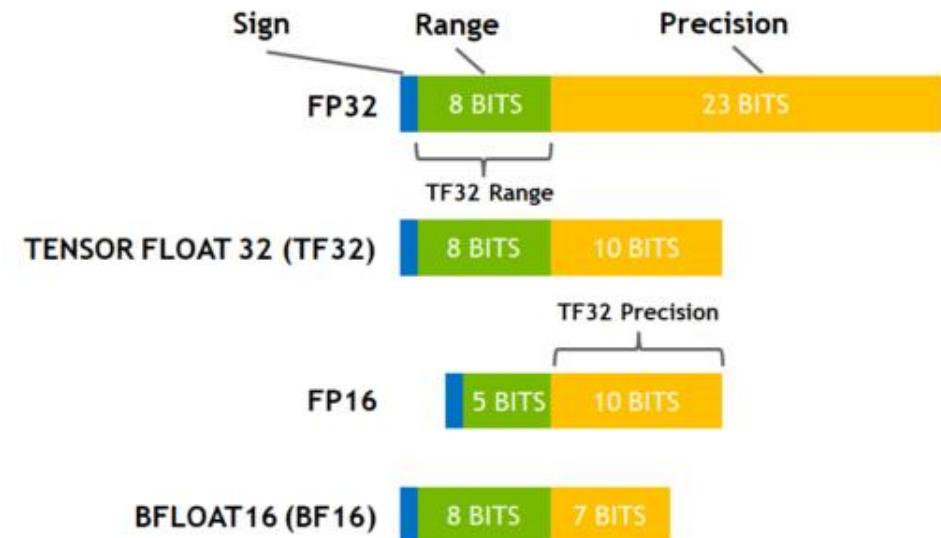




# Nvidia's TF32

- **Nvidia's TF32**

- 19-bit (BF19)
- 1-bit sign, 8-bit exponent  
10-bit fraction
- Fuse BF16 and FP16
  - BF16: 8-bit exponent +
  - FP16: 10-bit fraction
- Nvidia A100 Tensor Core
  - TF32: 156 TFLOPS
  - FP16/BF16: 312 TFLOPS



<https://reurl.cc/Omo1dv>



# FP8 and Tesla CFloat

- **FP8 (1-5-2)**
  - Large loss in MobileNet v2
  - Hybrid FP8 (HFP8)
    - Use FP(1-4-3) in forward
    - Use FP(1-5-2) in backward
- **Tesla Dojo Cffloat (configurable float)**
  - Configurable exponent and mantissa
  - Use software to choose appropriate Cffloat format
    - CF16
    - CF8 (1-4-3), CF8 (1-5-2)

c. Trans-precision Inference Accuracy  
of FP32 models in FP8 1-5-2 precision

| FP32 Model                    | Baseline | FP8 1-5-2    |
|-------------------------------|----------|--------------|
| MobileNet_v2<br>ImageNet      | 71.81    | <b>52.51</b> |
| ResNet50<br>ImageNet          | 76.44    | 75.31        |
| DenseNet121<br>ImageNet       | 74.76    | 73.64        |
| MaskRCNN<br>COCO <sup>†</sup> | 33.58    | 32.83        |
|                               | 29.27    | 28.65        |

<sup>†</sup> Box and Mask average precision

<https://proceedings.neurips.cc/paper/2019/file/65fc9fb4897a89789352e211ca2d398f-Paper.pdf>



# Nvidia's NVFP4

- **Nvidia's NVFP4**

- 1 sign bit, 2 exponent bits, and 1 mantissa bit (E2M1)
- The value in the format ranges approximately -6 to 6
- The values in the range could include 0.0, 0.5, 1.0, 1.5, 2, 3, 4, 6 (same for the negative range)

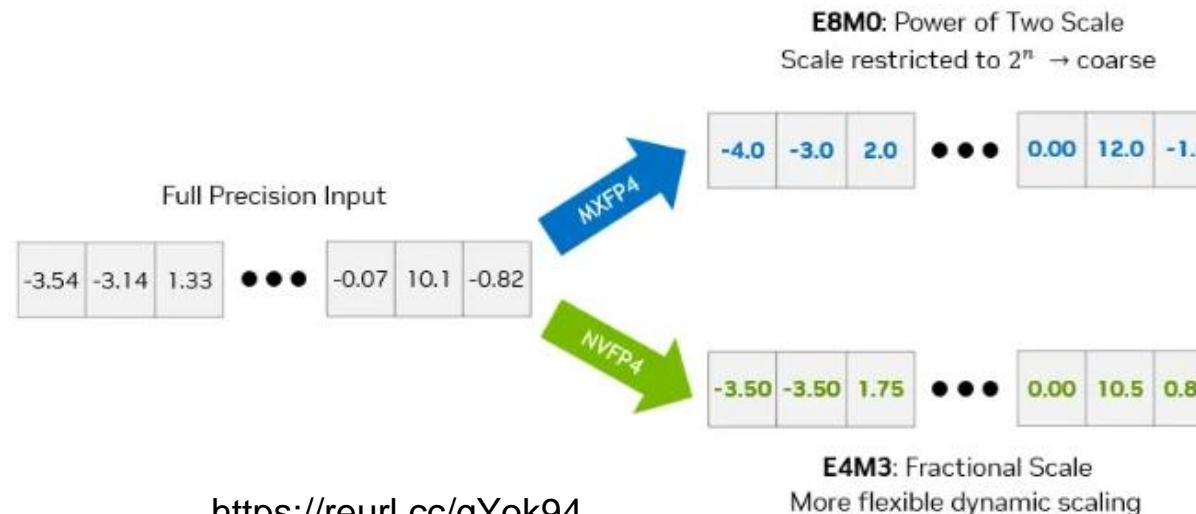
Table 1: Numbers represented in FP4-E2M1 with NaN and Inf (IEEE 754 standard) and Numbers represented in FP4-E2M1 without NaN and Inf (Our design).

| UINT4           | 0 | 1   | 2 | 3   | 4 | 5 | 6   | 7   | 8  | 9    | 10 | 11   | 12 | 13 | 14  | 15  |
|-----------------|---|-----|---|-----|---|---|-----|-----|----|------|----|------|----|----|-----|-----|
| FP4 w/ NaN&Inf  | 0 | 0.5 | 1 | 1.5 | 2 | 3 | Inf | NaN | -0 | -0.5 | -1 | -1.5 | -2 | -3 | Inf | NaN |
| FP4 w/o NaN&Inf | 0 | 0.5 | 1 | 1.5 | 2 | 3 | 4   | 6   | -0 | -0.5 | -1 | -1.5 | -2 | -3 | -4  | -6  |



# Nvidia's NVFP4

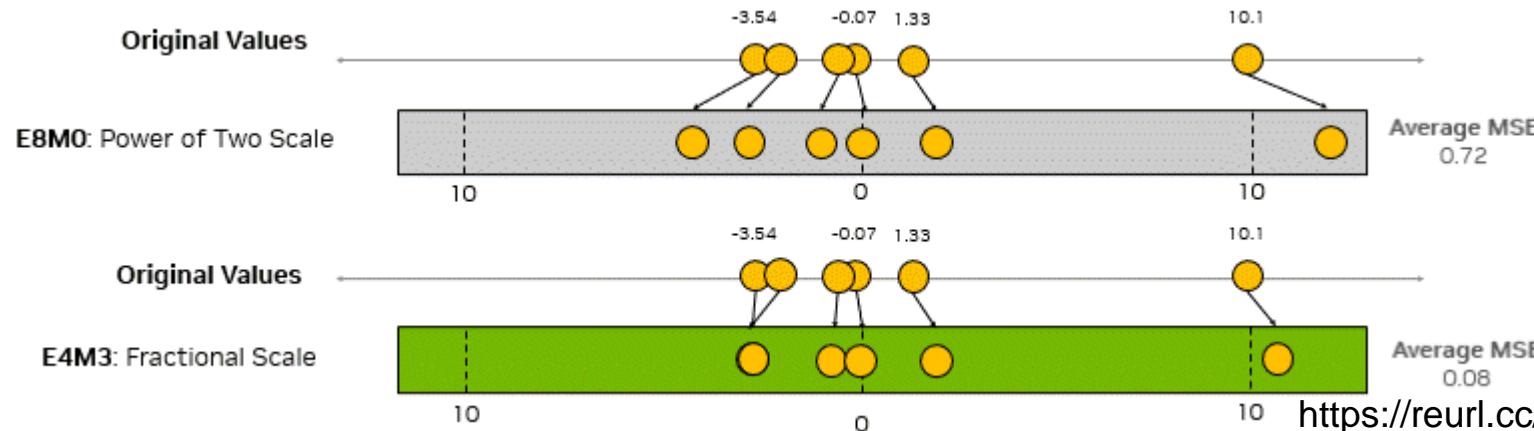
- **High-precision scaling**
  - NVFP4 encodes blocks using E4M3 FP8 precision
  - Enables non-power-of-two scaling factors with fractional precision





# Nvidia's NVFP4

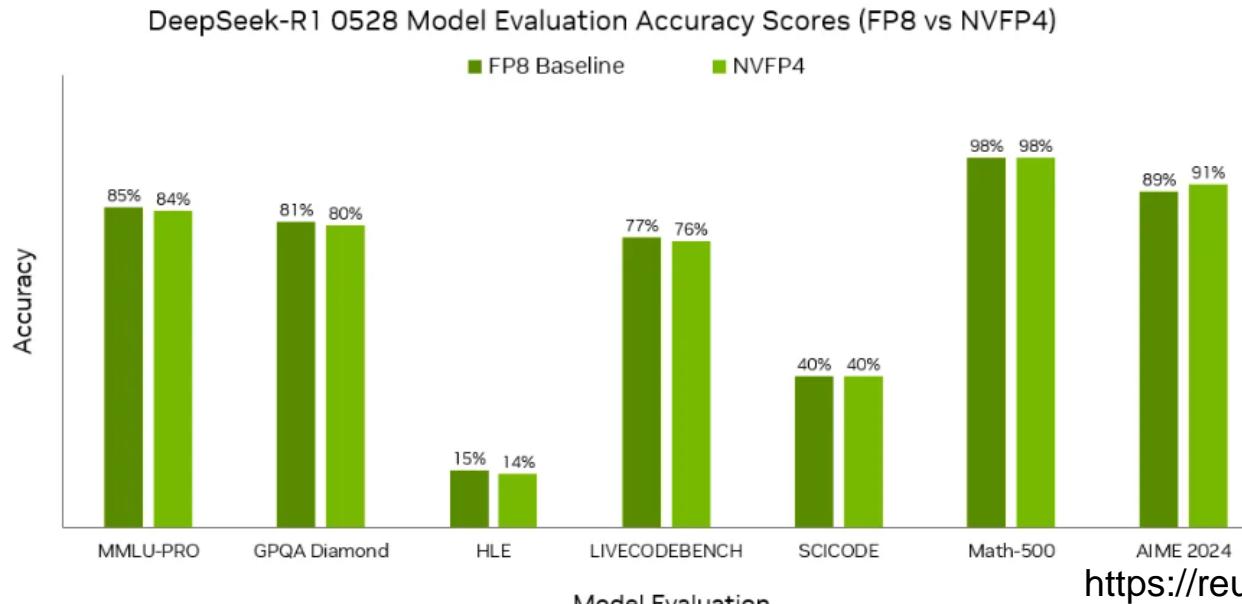
- **High-precision scaling**
  - E8M0 = Snaps the scale factor to nearest  $2^n$
  - E4M3 = Finds one scale factor that makes the block errors collectively as small as possible





# Nvidia's NVFP4

- **Quantize model weights to 4-bits**
  - The analysis showcases the 1% or less accuracy degradation





# How to Determine Bit Width on DNN ?

- For accuracy, DNN operations decide bit width to achieve sufficient precision
- Which DNN operations affect the accuracy ?
  - **For inference:** weights, activations, and partial sums
  - **For training:** weights, activations, partial sums, gradients, and weight update
    - post-training quantization (PTQ)
      - A model compression technique that converts a pre-trained, full-precision model into a lower-precision model without needing to retrain or fine-tune it



# Takeaway Questions

- What are advantages to use BF16 instead of FP16 ?
  - (A) Fast conversion from FP32
  - (B) Get more precise value
  - (C) Represent few different values
- What are benefits to use lower precision data type on neural network ?
  - (A) Reduce the latency of DNN models
  - (B) Save the memory space
  - (C) Lower the power consumption of the accelerator



# K-Means-based Weight Quantization

- **Storage**
  - Integer Weights; Floating-Point Codebook
- **Computation**
  - Floating-Point Arithmetic

weights  
(32-bit float)

|       |       |       |       |
|-------|-------|-------|-------|
| 2.09  | -0.98 | 1.48  | 0.09  |
| 0.05  | -0.14 | -1.08 | 2.12  |
| -0.91 | 1.92  | 0     | -1.03 |
| 1.87  | 0     | 1.53  | 1.49  |

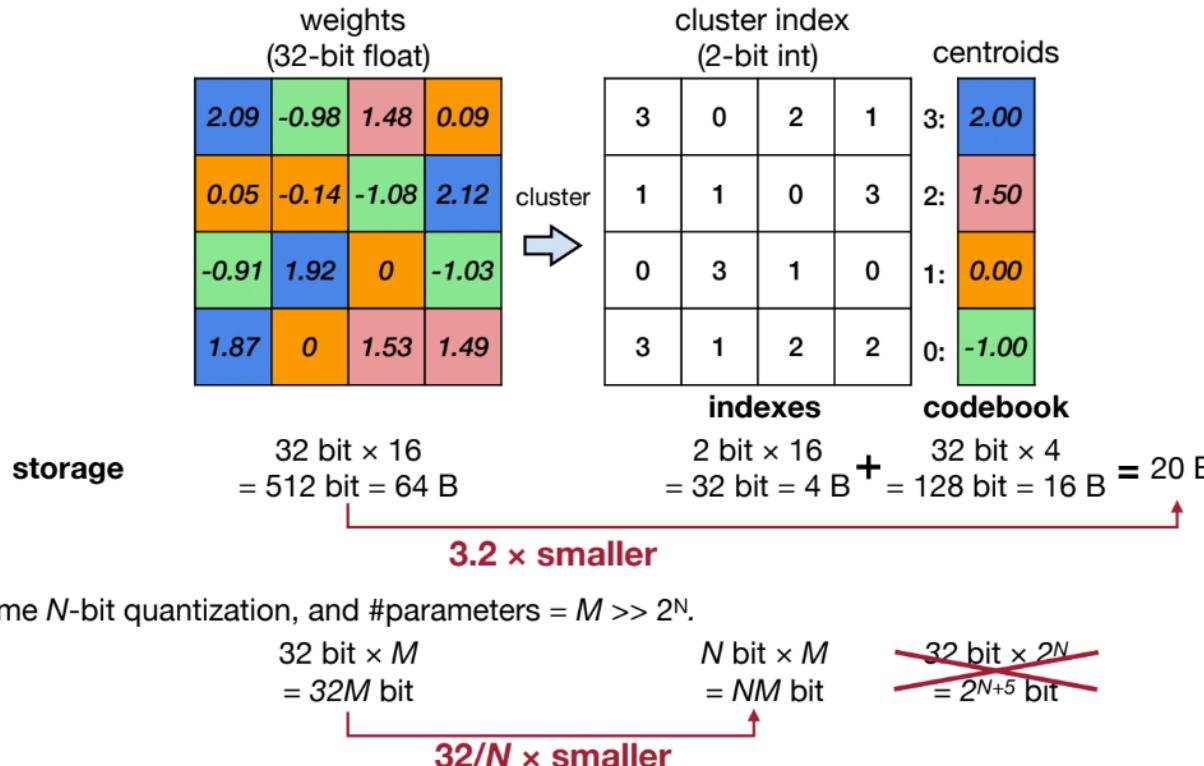
~~2.09, 2.12, 1.92, 1.87~~



2.0



# K-Means-based Weight Quantization

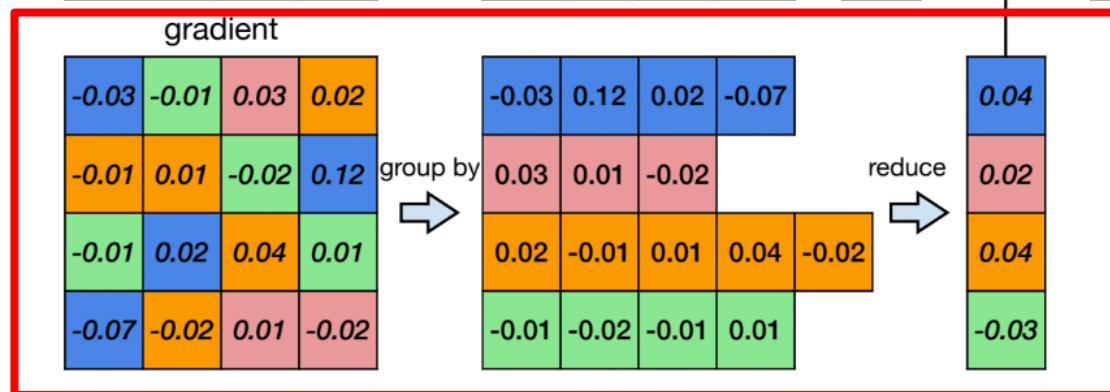
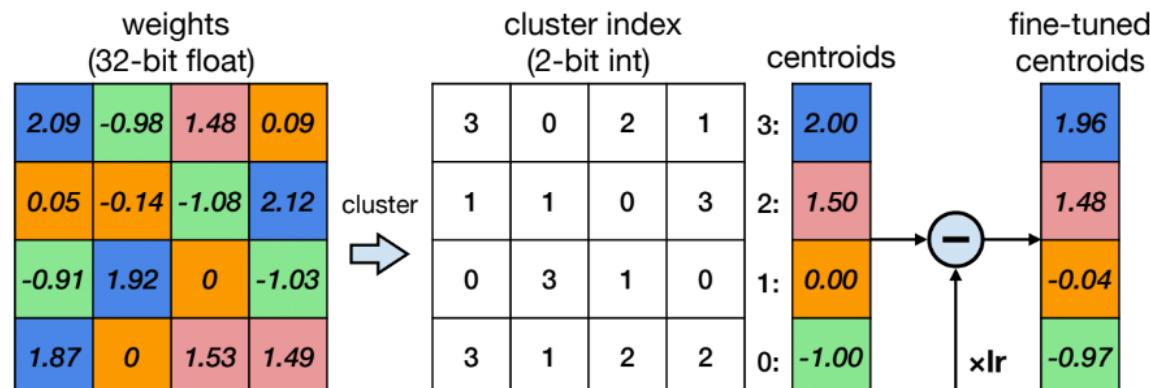


Assume  $N$ -bit quantization, and #parameters  $= M \gg 2^N$ .



# K-Means-based Weight Quantization

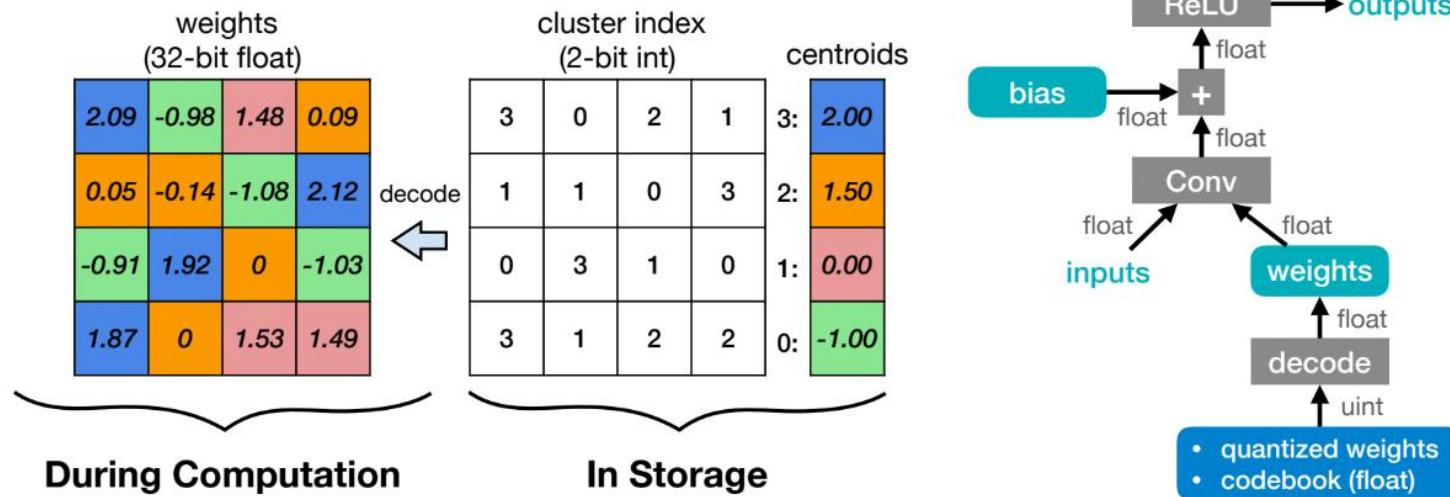
- **Fine-tuning Quantized Weights**
  - Reduce the quantization error





# K-Means-based Weight Quantization

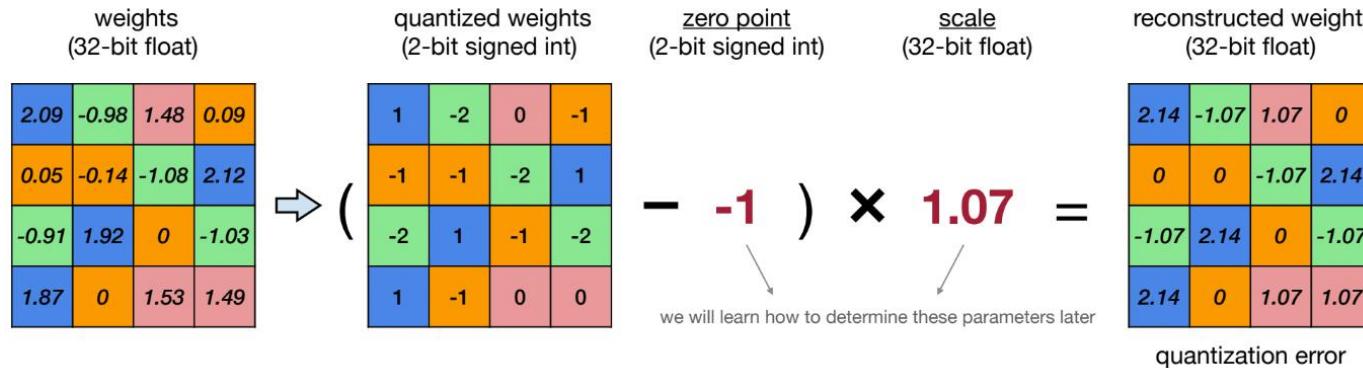
- Weights are decompressed using a lookup table during runtime inference
- Only saves storage cost of a neural network model
- All the computation and memory access are still floating-point





# What is Linear Quantization ?

- An affine mapping of integers to real numbers
- **Storage:** Integer Weights; **Computation:** Integer Arithmetic



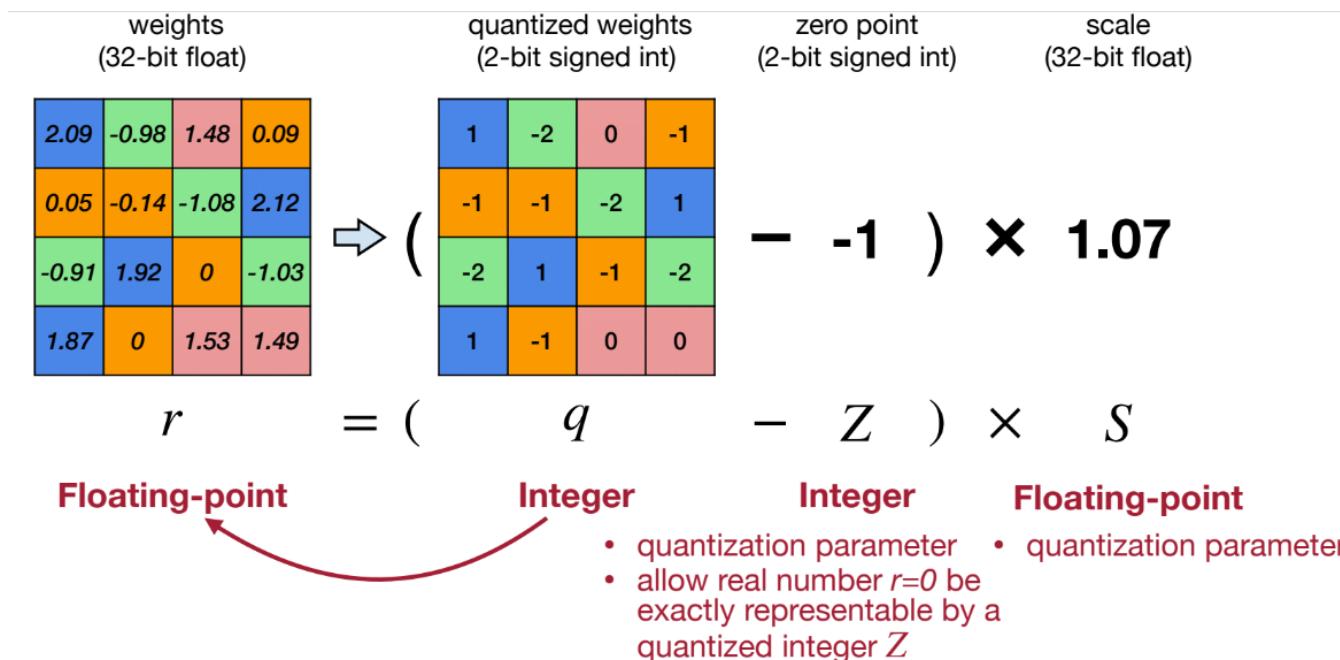
| Binary | Decimal |
|--------|---------|
| 01     | 1       |
| 00     | 0       |
| 11     | -1      |
| 10     | -2      |

|       |       |       |       |
|-------|-------|-------|-------|
| -0.05 | 0.09  | 0.41  | 0.09  |
| 0.05  | -0.14 | -0.01 | -0.02 |
| 0.16  | -0.22 | 0     | 0.04  |
| -0.27 | 0     | 0.46  | 0.42  |



# Linear Quantization

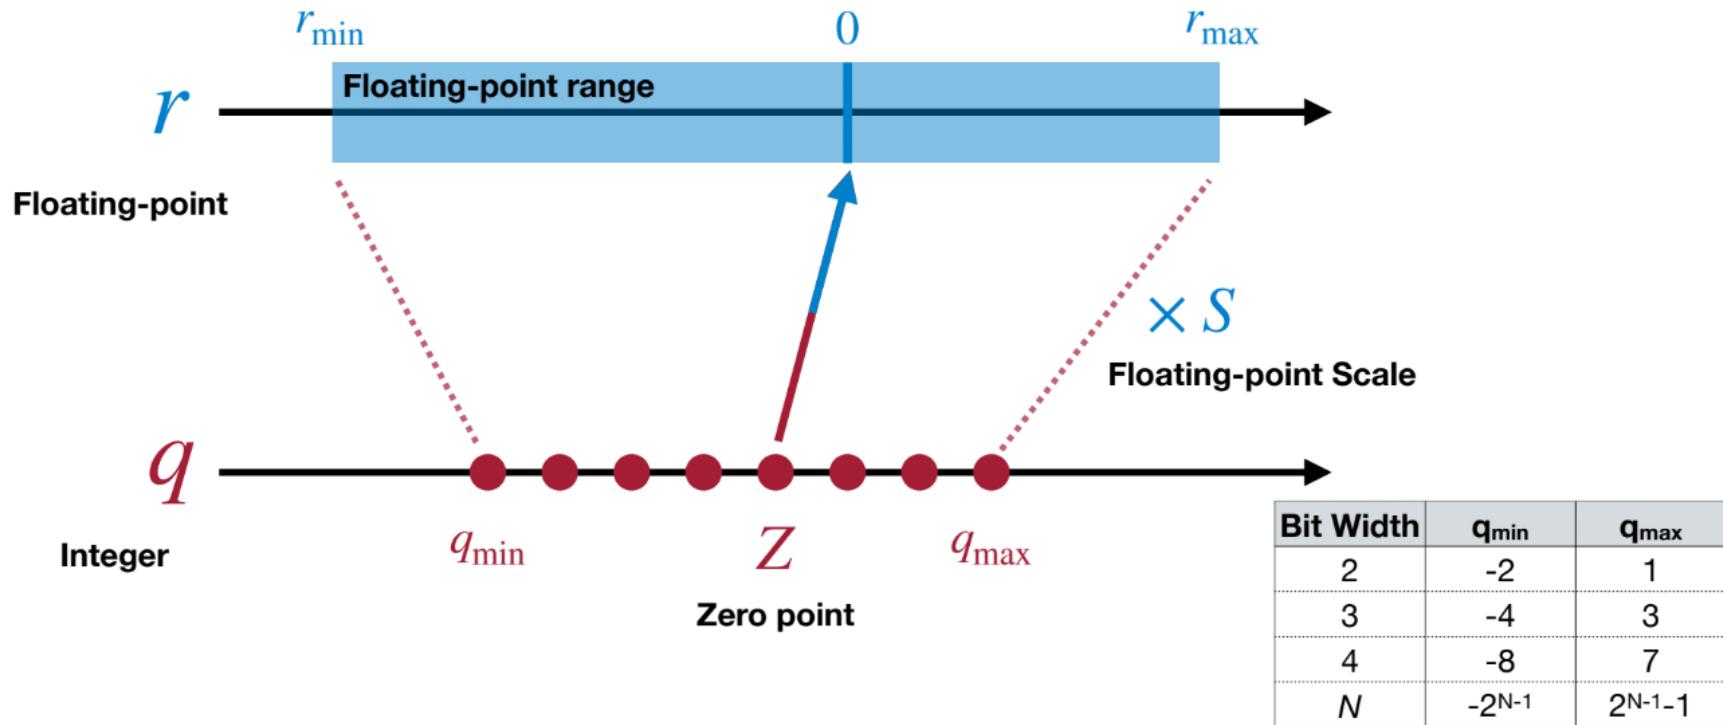
- An affine mapping of integers to real numbers ( $r = S(q - Z)$ )





# Linear Quantization

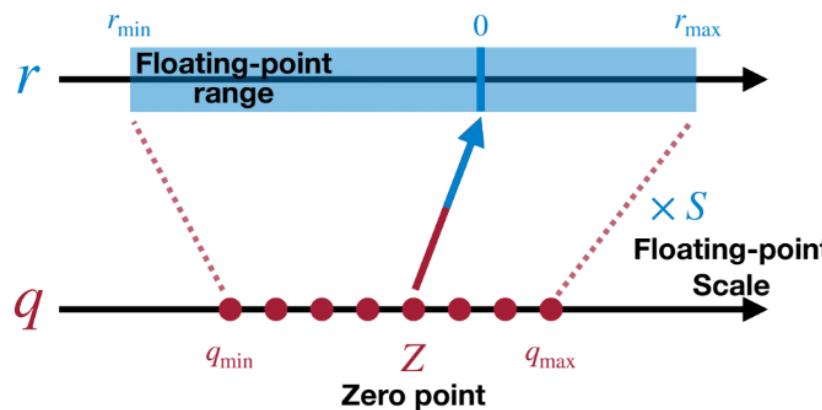
- An affine mapping of integers to real numbers ( $r = S(q - Z)$ )





# Scale of Linear Quantization

- An affine mapping of integers to real numbers ( $r = S(q - Z)$ )



$$r_{\max} = S (q_{\max} - Z) \quad \text{---}$$
$$r_{\min} = S (q_{\min} - Z)$$

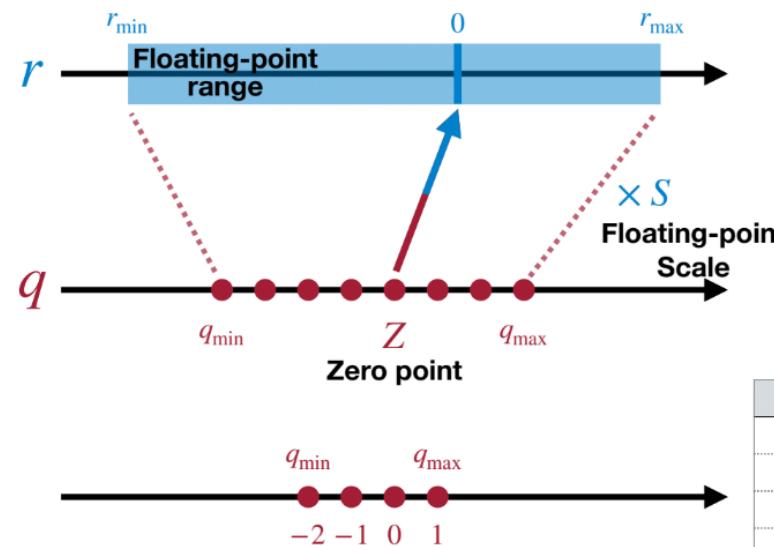
$$r_{\max} - r_{\min} = S (q_{\max} - q_{\min})$$

$$S = \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}}$$



# Scale of Linear Quantization

- An affine mapping of integers to real numbers ( $r = S(q - Z)$ )



|       |       |       |       |
|-------|-------|-------|-------|
| 2.09  | -0.98 | 1.48  | 0.09  |
| 0.05  | -0.14 | -1.08 | 2.12  |
| -0.91 | 1.92  | 0     | -1.03 |
| 1.87  | 0     | 1.53  | 1.49  |

$$S = \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}}$$

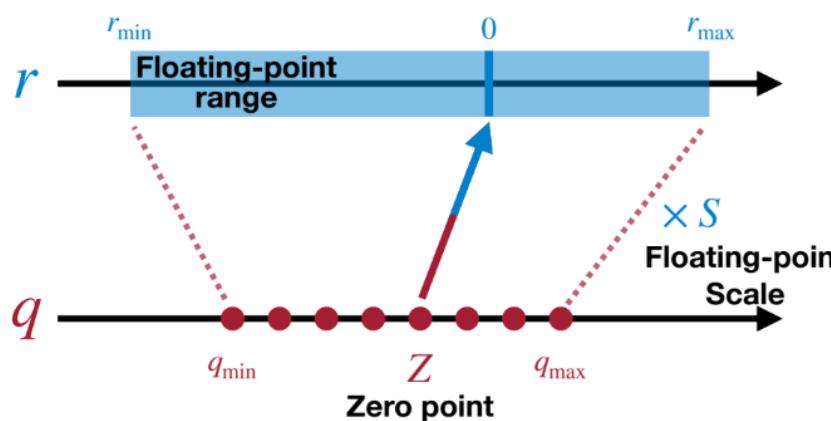
$$= \frac{2.12 - (-1.08)}{1 - (-2)}$$
$$= 1.07$$

| Binary | Decimal |
|--------|---------|
| 01     | 1       |
| 00     | 0       |
| 11     | -1      |
| 10     | -2      |



# Zero Point of Linear Quantization

- An affine mapping of integers to real numbers ( $r = S(q - Z)$ )



$$r_{\min} = S (q_{\min} - Z)$$



$$Z = q_{\min} - \frac{r_{\min}}{S}$$

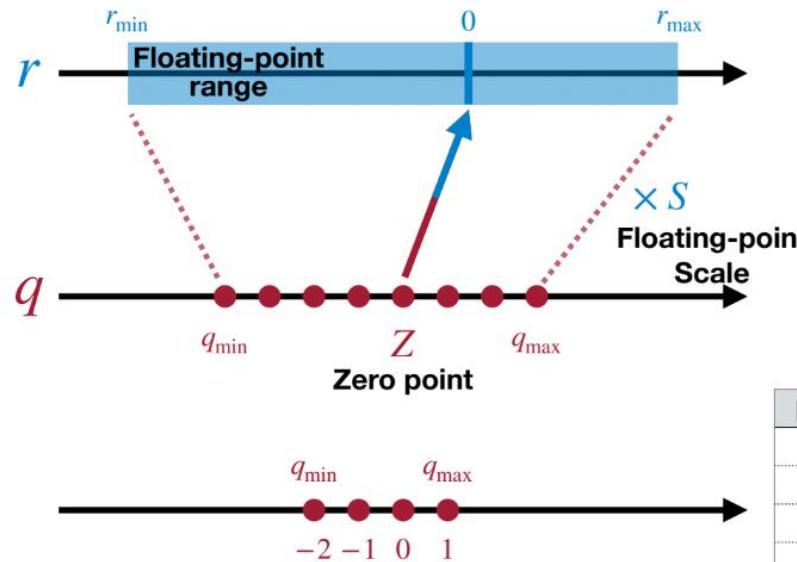


$$Z = \text{round} \left( q_{\min} - \frac{r_{\min}}{S} \right)$$



# Zero Point of Linear Quantization

- An affine mapping of integers to real numbers ( $r = S(q - Z)$ )



|       |       |       |       |
|-------|-------|-------|-------|
| 2.09  | -0.98 | 1.48  | 0.09  |
| 0.05  | -0.14 | -1.08 | 2.12  |
| -0.91 | 1.92  | 0     | -1.03 |
| 1.87  | 0     | 1.53  | 1.49  |

$$Z = q_{\min} - \frac{r_{\min}}{S}$$

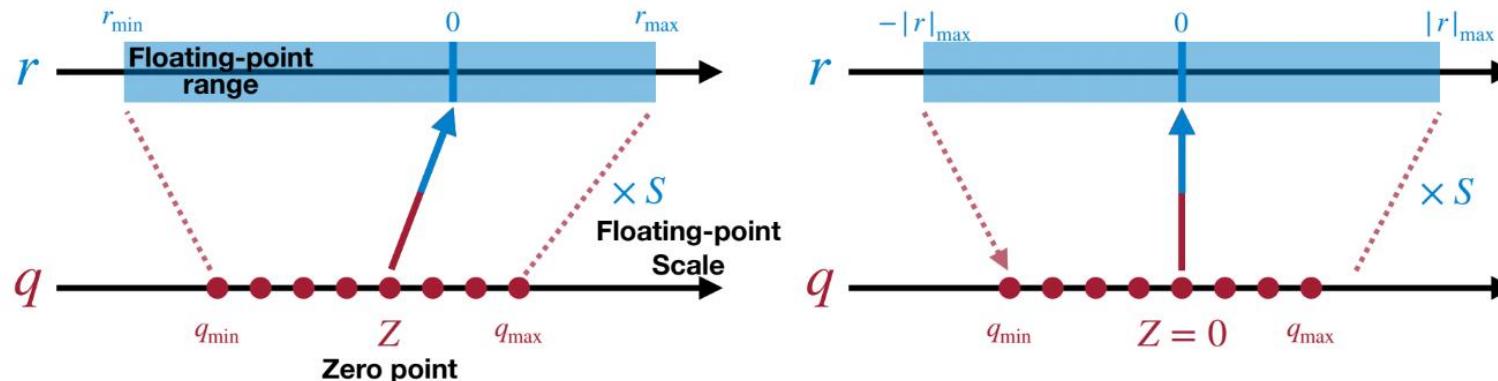
$$\begin{aligned} &= \text{round}\left(-2 - \frac{-1.08}{1.07}\right) \\ &= -1 \end{aligned}$$

| Binary | Decimal |
|--------|---------|
| 01     | 1       |
| 00     | 0       |
| 11     | -1      |
| 10     | -2      |



# Asymmetric Linear Quantization

- Full range mode



$$S = \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}}$$

$$S = \frac{r_{\min}}{q_{\min} - Z} = \frac{-|r|_{\max}}{q_{\min}} = \frac{|r|_{\max}}{2^{N-1}}$$

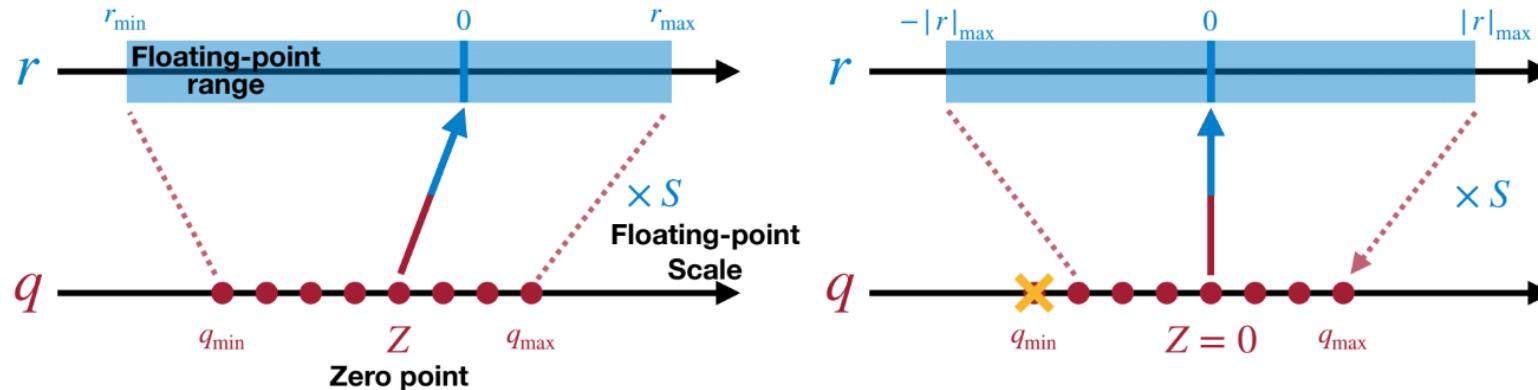
- use full range of quantized integers
- example: PyTorch's native quantization, ONNX

| Bit Width | $q_{\min}$ | $q_{\max}$  |
|-----------|------------|-------------|
| 2         | -2         | 1           |
| 3         | -4         | 3           |
| 4         | -8         | 7           |
| N         | $-2^{N-1}$ | $2^{N-1}-1$ |



# Symmetric Linear Quantization

- Restricted range mode



$$S = \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}}$$

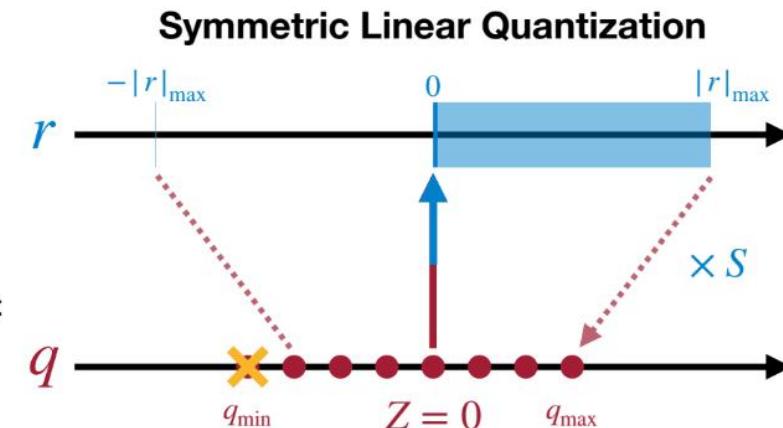
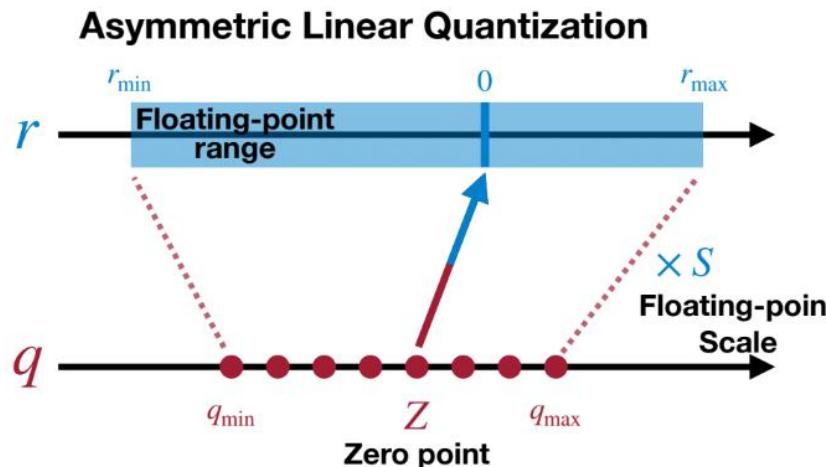
$$S = \frac{r_{\max}}{q_{\max} - Z} = \frac{|r|_{\max}}{q_{\max}} = \frac{|r|_{\max}}{2^{N-1} - 1}$$

| Bit Width | $q_{\min}$ | $q_{\max}$  |
|-----------|------------|-------------|
| 2         | -2         | 1           |
| 3         | -4         | 3           |
| 4         | -8         | 7           |
| N         | $-2^{N-1}$ | $2^{N-1}-1$ |

- example: TensorFlow, NVIDIA TensorRT, Intel DNNL



# Asymmetric vs. Symmetric



- The quantized range is fully used.
- The implementation is more complex, and zero points require additional logic in hardware.

- The quantized range will be wasted for biased float range.
  - Activation tensor is non-negative after ReLU, and thus symmetric quantization will lose 1 bit effectively.
- The implementation is much simpler.

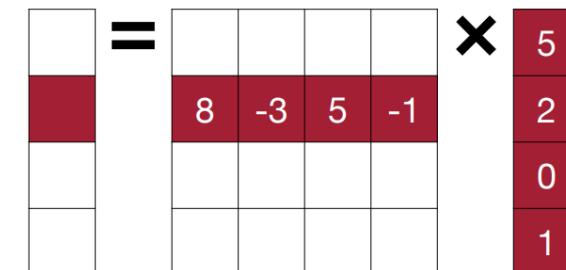
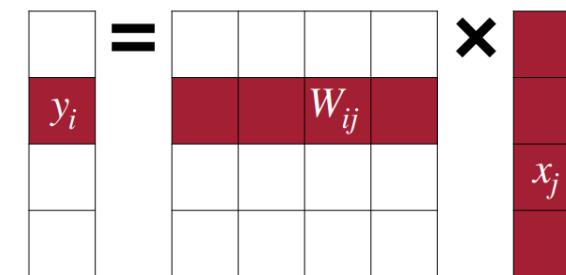


# Binary/Ternary Quantization

- Could we push the quantization precision to 1 bit?

$$y_i = \sum_j W_{ij} \cdot x_j$$
$$= 8 \times 5 + (-3) \times 2 + 5 \times 0 + (-1) \times 1$$

| input        | weight       | operations | memory     | computation |
|--------------|--------------|------------|------------|-------------|
| $\mathbb{R}$ | $\mathbb{R}$ | $+ \times$ | $1 \times$ | $1 \times$  |
|              |              |            |            |             |
|              |              |            |            |             |





# Binary/Ternary Quantization

- If weights are quantized to +1 and -1

$$y_i = \sum_j W_{ij} \cdot x_j$$
$$= 5 - 2 + 0 - 1$$

| input | weight | operations | memory    | computation |
|-------|--------|------------|-----------|-------------|
| R     | R      | + ×        | 1x        | 1x          |
| R     | B      | + -        | ~32x less | ~2x less    |

$$\begin{matrix} & = & \begin{matrix} & & & \\ & & & \\ \textcolor{darkred}{8} & -3 & 5 & -1 \\ & & & \\ & & & \end{matrix} & \times & \begin{matrix} 5 \\ 2 \\ 0 \\ 1 \end{matrix} \end{matrix}$$

$$\begin{matrix} & = & \begin{matrix} & & & \\ & & & \\ \textcolor{darkred}{1} & -1 & 1 & -1 \\ & & & \\ & & & \end{matrix} & \times & \begin{matrix} 5 \\ 2 \\ 0 \\ 1 \end{matrix} \end{matrix}$$



# Binarization

- **Deterministic Binarization**

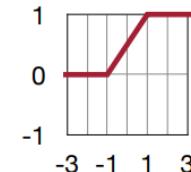
- directly computes the bit value based on a threshold, usually 0, resulting in a sign function.

$$q = \text{sign}(r) = \begin{cases} +1, & r \geq 0 \\ -1, & r < 0 \end{cases}$$

- **Stochastic Binarization**

- use global statistics or the value of input data to determine the probability of being -1 or +1
  - e.g., in Binary Connect (BC), probability is determined by hard sigmoid function  $\sigma(r)$

$$q = \begin{cases} +1, & \text{with probability } p = \sigma(r) \\ -1, & \text{with probability } 1 - p \end{cases}, \quad \text{where } \sigma(r) = \min(\max(\frac{r+1}{2}, 0), 1)$$



- harder to implement as it requires the hardware to generate random bits when quantizing.



# Minimizing Quantization Error in Binarization

| weights<br>(32-bit float) |       |       |       |
|---------------------------|-------|-------|-------|
| 2.09                      | -0.98 | 1.48  | 0.09  |
| 0.05                      | -0.14 | -1.08 | 2.12  |
| -0.91                     | 1.92  | 0     | -1.03 |
| 1.87                      | 0     | 1.53  | 1.49  |

$$\mathbf{W}^B = \text{sign}(\mathbf{W})$$

$$\alpha = \frac{1}{n} \|\mathbf{W}\|_1$$

$$\begin{array}{l} \mathbf{W} \\ \downarrow \quad \quad \quad \downarrow \\ \mathbf{W}^B \quad \quad \quad \alpha \mathbf{W}^B \end{array}$$

| binary weights<br>(1-bit) |    |    |    |
|---------------------------|----|----|----|
| 1                         | -1 | 1  | 1  |
| 1                         | -1 | -1 | 1  |
| -1                        | 1  | 1  | -1 |
| 1                         | 1  | 1  | 1  |

| scale<br>(32-bit float) |    |    |    |
|-------------------------|----|----|----|
| 1                       | -1 | 1  | 1  |
| 1                       | -1 | -1 | 1  |
| -1                      | 1  | 1  | -1 |
| 1                       | 1  | 1  | 1  |

| AlexNet-based Network       | ImageNet Top-1 Accuracy Delta |
|-----------------------------|-------------------------------|
| BinaryConnect               | -21.2%                        |
| Binary Weight Network (BWN) | 0.2%                          |

$$\|\mathbf{W} - \mathbf{W}^B\|_F^2 = 9.28$$

$$\times 1.05 = \frac{1}{16} \|\mathbf{W}\|_1$$

$$\|\mathbf{W} - \alpha \mathbf{W}^B\|_F^2 = 9.24$$



# Binary Net

- **Binary Connect**

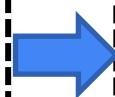
- Weights  $\{-1, 1\}$  (Bipolar binary),  
Activation 32-bit float
- Accuracy loss: 19 % on AlexNet

- **Binarized Neural Networks**

- Weights  $\{-1, 1\}$ , Activations  $\{-1, 1\}$
- Both of operands are binary, the multiplication turns into an XNOR
- Accuracy loss: 29.8 % on AlexNet

for each  $i$  in width:

$$C += A[\text{row}][i] * B[i][\text{col}]$$



for each  $i$  in width:

$$C += \text{popcount}(\text{XNOR}(A[\text{row}][i], B[i][\text{col}]))$$

| XNOR |   |     |
|------|---|-----|
| A    | B | Out |
| 0    | 0 | 1   |
| 1    | 0 | 0   |
| 0    | 1 | 0   |
| 1    | 1 | 1   |

Popcount (110010001) = 4



# Case Study: Binary Multiplication

- $A = 10010, B = 01111$  (0 is really -1 here)
- **Dot product:**
  - $A * B = (1 * -1) + (-1 * 1) + (-1 * 1) + (1 * 1) + (-1 * 1) = -3$
- $P = \text{XNOR}(A, B) = 00010, \text{popcount}(P) = 1$
- Result =  $2 * P - N$ , where  $N$  is the total number of bits
- $2 * P - N = 2 * 1 - 5 = -3$



# XNOR-Net

- If both activations and weights are binarized

$$\begin{aligned}y_i &= \sum_j W_{ij} \cdot x_j \\&= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1 \\&= 1 + (-1) + (-1) + (-1) = -2\end{aligned}$$

$$\begin{array}{c|c} \text{ } & = \\ \text{ } & \times \end{array} \begin{array}{c|c|c|c|c} & & & & \\ \text{ } & 8 & -3 & 5 & -1 \\ & & & & \\ & & & & \\ & & & & \end{array} \begin{array}{c|c} 5 \\ 2 \\ 0 \\ 1 \end{array}$$

$$\begin{array}{c|c} \text{ } & = \\ \text{ } & \times \end{array} \begin{array}{c|c|c|c|c} & & & & \\ \text{ } & 1 & -1 & 1 & -1 \\ & & & & \\ & & & & \\ & & & & \end{array} \begin{array}{c|c} 1 \\ 1 \\ -1 \\ 1 \end{array}$$



# XNOR-Net

- If both activations and weights are binarized

$$\begin{aligned}y_i &= \sum_j W_{ij} \cdot x_j \\&= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1 \\&= 1 + (-1) + (-1) + (-1) = -2\end{aligned}$$

| W  | X  | Y=WX |
|----|----|------|
| 1  | 1  | 1    |
| 1  | -1 | -1   |
| -1 | -1 | 1    |
| -1 | 1  | -1   |

| b <sub>W</sub> | b <sub>X</sub> | XNOR(b <sub>W</sub> , b <sub>X</sub> ) |
|----------------|----------------|--|
| 1              | 1              | 1                                      |
| 1              | 0              | 0                                      |
| 0              | 0              | 1                                      |
| 0              | 1              | 0                                      |



# XNOR-Net

- If both activations and weights are binarized

$$y_i = \sum_j W_{ij} \cdot x_j$$

$$\begin{aligned} &= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1 \\ &= 1 + (-1) + (-1) + (-1) = -2 \end{aligned}$$



$$\begin{aligned} &= 1 \text{ xnor } 1 + 0 \text{ xnor } 1 + 1 \text{ xnor } 0 + 0 \text{ xnor } 1 \\ &= 1 + 0 + 0 + 0 = 1 \end{aligned}$$



| W  | X  | Y=WX |
|----|----|------|
| 1  | 1  | 1    |
| 1  | -1 | -1   |
| -1 | -1 | 1    |
| -1 | 1  | -1   |

| b <sub>w</sub> | b <sub>x</sub> | XNOR(b <sub>w</sub> , b <sub>x</sub> ) |
|----------------|----------------|--|
| 1              | 1              | 1                                      |
| 1              | 0              | 0                                      |
| 0              | 0              | 1                                      |
| 0              | 1              | 0                                      |



# XNOR-Net

- If both activations and weights are binarized

$$\begin{aligned}y_i &= \sum_j W_{ij} \cdot x_j \\&= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1 \\&= 1 + (-1) + (-1) + (-1) = -2\end{aligned}$$

$$\begin{aligned}y_i &= -n + 2 \cdot \sum_j W_{ij} \text{ xnor } x_j \\&= 1 \text{ xnor } 1 + 0 \text{ xnor } 1 + 1 \text{ xnor } 0 + 0 \text{ xnor } 1 \\&= 1 + 0 + 0 + 0 = \boxed{\begin{matrix} 1 & \times 2 \\ + & \\ -4 & \end{matrix}} = -2\end{aligned}$$

Assuming  $-1 \quad -1 \quad -1 \quad -1 \rightarrow -4$

| W  | X  | Y=WX |
|----|----|------|
| 1  | 1  | 1    |
| 1  | -1 | -1   |
| -1 | -1 | 1    |
| -1 | 1  | -1   |

| b <sub>w</sub> | b <sub>x</sub> | XNOR(b <sub>w</sub> , b <sub>x</sub> ) |
|----------------|----------------|--|
| 1              | 1              | 1                                      |
| 1              | 0              | 0                                      |
| 0              | 0              | 1                                      |
| 0              | 1              | 0                                      |



# XNOR-Net

- If both activations and weights are binarized

$$\begin{aligned}y_i &= -n + 2 \cdot \sum_j W_{ij} \text{ xnor } x_j \quad \rightarrow \quad y_i = -n + \text{popcount}(W_i \text{ xnor } x) \ll 1 \\&= -4 + 2 \times (1 \text{ xnor } 1 + 0 \text{ xnor } 1 + 1 \text{ xnor } 0 + 0 \text{ xnor } 1) \\&= -4 + 2 \times (1 + 0 + 0 + 0) = -2\end{aligned}$$

→ **popcount: return the number of 1**

| W  | X  | Y=WX |
|----|----|------|
| 1  | 1  | 1    |
| 1  | -1 | -1   |
| -1 | -1 | 1    |
| -1 | 1  | -1   |

| b <sub>w</sub> | b <sub>x</sub> | XNOR(b <sub>w</sub> , b <sub>x</sub> ) |
|----------------|----------------|--|
| 1              | 1              | 1                                      |
| 1              | 0              | 0                                      |
| 0              | 0              | 1                                      |
| 0              | 1              | 0                                      |



# XNOR-Net

- If both activations and weights are binarized

$$y_i = -n + \text{popcount}(W_i \text{ xnor } x) \ll 1$$

$$= -4 + \text{popcount}(1010 \text{ xnor } 1101) \ll 1$$

$$= -4 + \text{popcount}(1000) \ll 1 = -4 + 2 = -2$$

$$\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} = \begin{array}{c|c|c|c|c} & & & & \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \\ 8 & -3 & 5 & -1 & \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \end{array} \times \begin{array}{c} 5 \\ 2 \\ 0 \\ 1 \end{array}$$

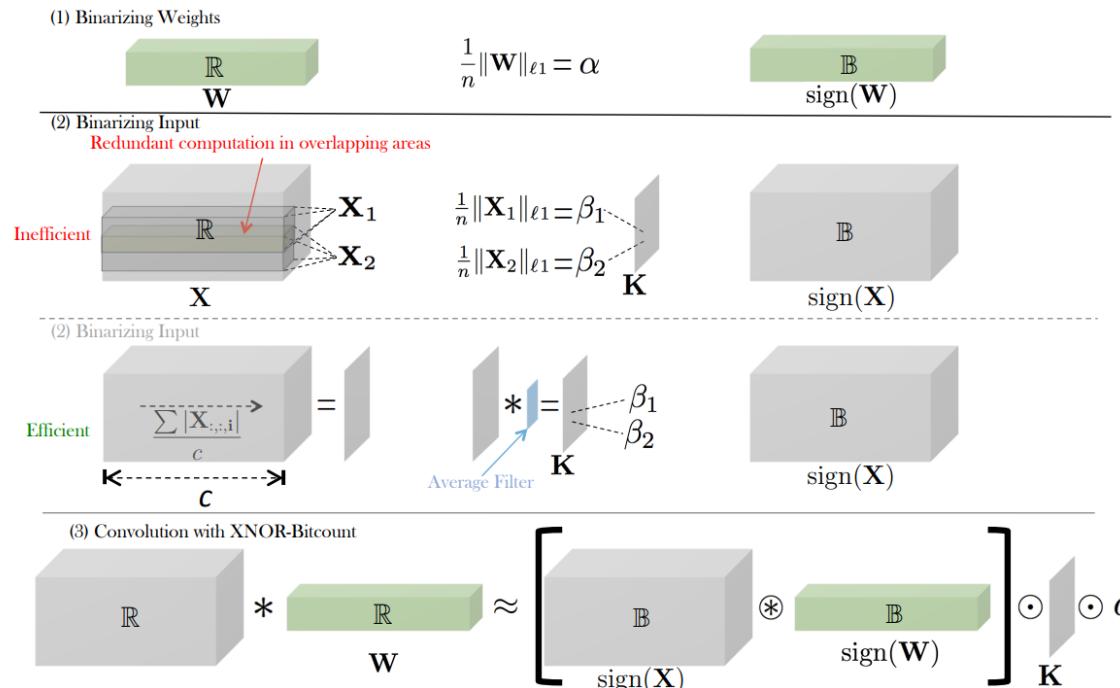
| input | weight | operations        | memory    | computation |
|-------|--------|-------------------|-----------|-------------|
| R     | R      | + ×               | 1×        | 1×          |
| R     | B      | + -               | ~32× less | ~2× less    |
| B     | B      | xnor,<br>popcount | ~32× less | ~58× less   |

$$\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} = \begin{array}{c|c|c|c|c} & & & & \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \\ 1 & -1 & 1 & -1 & \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \\ \text{---} & \text{---} & \text{---} & \text{---} & \text{---} \end{array} \times \begin{array}{c} 1 \\ 1 \\ -1 \\ 1 \end{array}$$



# XNOR-Net

- Minimizing quantization error in binarization





# XNOR-Net

| Neural Network | Quantization | Bit-Width |    | ImageNet<br>Top-1 Accuracy<br>Delta |
|----------------|--------------|-----------|----|-------------------------------------|
|                |              | W         | A  |                                     |
| AlexNet        | BWN          | 1         | 32 | 0.2%                                |
|                | BNN          | 1         | 1  | -28.7%                              |
|                | XNOR-Net     | 1         | 1  | -12.4%                              |
| GoogleNet      | BWN          | 1         | 32 | -5.80%                              |
|                | BNN          | 1         | 1  | -24.20%                             |
| ResNet-18      | BWN          | 1         | 32 | -8.5%                               |
|                | XNOR-Net     | 1         | 1  | -18.1%                              |

\* BWN: Binary Weight Network with scale for weight binarization

\* BNN: Binarized Neural Network without scale factors

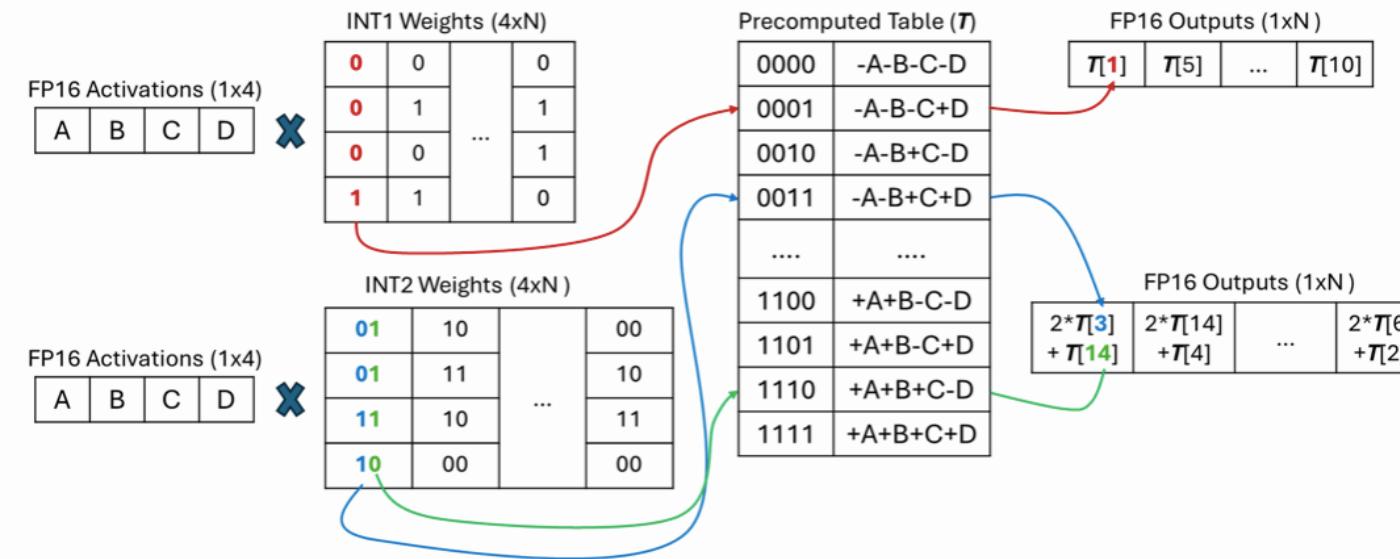
\* XNOR-Net: scale factors for both activation and weight binarization



# BitNet

- FP16 activation and 1.58 bit weights – Transformer-based model
- Lookup table (LUT) calculations

<https://www.arxiv.org/pdf/2407.00088>





# What do we Learn from Quantization?

- **Quantization** can improve DNN computational throughput while maintaining accuracy
- Layers on DNN models can be offered with **different bit widths**
- Varying bit width requires **the support of the hardware**
- **No systematic approach** to figure out the proper bit width in layers of DNN models
- What else ?



# Takeaway Questions

- What are purposes of data quantization ?
  - (A) Constrain the value of inputs to a set of discrete values
  - (B) Create more values
  - (C) Improve the degree of parallelism on DNN training
- Why training requires large bit width ?
  - (A) The training needs to compute more data
  - (B) Avoid the value underflow and overflow
  - (C) Gradient and weight update have a larger range