



Accelerator Architectures for Machine Learning (AAML)

Lecture 5: Systolic Accelerator

Tsung Tai Yeh
Department of Computer Science
National Yang-Ming Chiao Tung University



Acknowledgements and Disclaimer

- Slides was developed in the reference with
Joel Emer, Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, ISCA 2019
tutorial
Efficient Processing of Deep Neural Network, Vivienne Sze, Yu-Hsin
Chen, Tien-Ju Yang, Joel Emer, Morgan and Claypool Publisher, 2020
Yakun Sophia Shao, EE290-2: Hardware for Machine Learning, UC
Berkeley, 2020
CS231n Convolutional Neural Networks for Visual Recognition,
Stanford University, 2020
CS224W: Machine Learning with Graphs, Stanford University, 2021



Outline

- Systolic Array Architecture
 - Google Tensor Processing Unit (TPU)
- Dataflow
 - Weight-stationary
 - Output-stationary
 - Input-stationary



Systolic DNN Accelerator



A Golden Age in Microprocessor Design

- A great leap in microprocessor speed $\sim 10^6$ X faster over 40 years
- Architectural innovations
 - Width: 8->16->32->64 bits (~8X)
 - Instruction level parallelism (ILP)
 - Multicore: 1 processor to 16 cores
 - Clock rate: 3 – 4000 MHz (~1000 X through technology & architecture)
- IC technology makes it possible
 - **Moore's Law:** growth in transistor count (2X every 1.5 years)
 - **Dennard Scaling:** power/transistor shrinks at the same rate as transistors are added



Current Situation

- **Technology**
 - End of Dennard scaling: power becomes the key constraint
 - Slowdown of Moore's Law: transistor cost
- **Architectural Designs**
 - Inefficiency to exploit instruction level parallelism in the uniprocessor era, 2004
 - Amdahl's Law and its implications end



What's Left ?

- Transistors not getting much better
- Power budget not getting much higher
- One inefficient processor/chip to N efficient processors/chip
- Only path left is **Domain Specific Architectures**
 - Just do a few tasks, but extremely well



Lessons from DSA

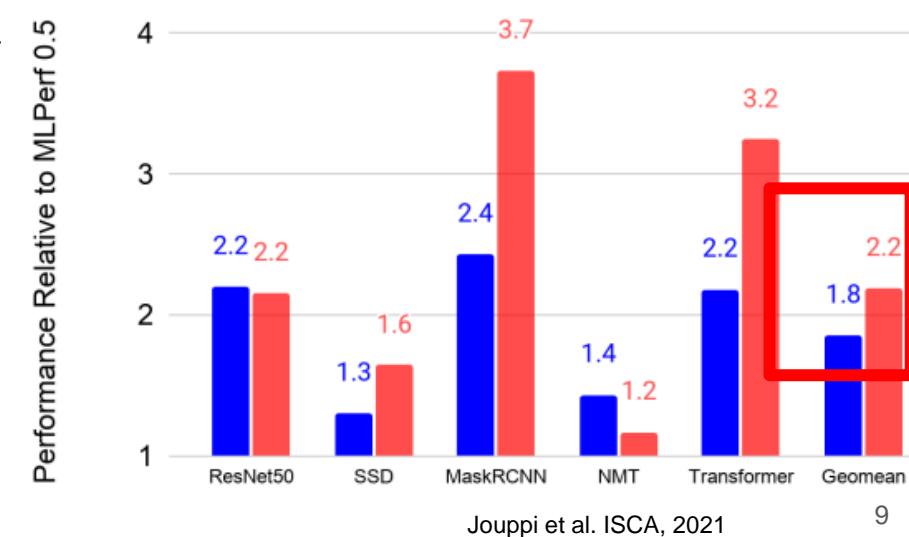
- **Logic, wires, SRAM & DRAM improve unequally**
 - **SRAM access improved only 1.3X – 2.4 X → SRAM density is scaling slowly**
 - **DRAM access improved 6.3X**
 - Packaging innovations
 - **High Bandwidth Memory (HBM)**
 - HBM is more energy-efficient than GDDR6 or DDR DRAM
 - **Logic improves much faster than wires and SRAM**

Operation		Picojoules per Operation		
		45 nm	7 nm	45 / 7
+	Int 8	0.03	0.007	4.3
	Int 32	0.1	0.03	3.3
	BFloat 16	--	0.11	--
	IEEE FP 16	0.4	0.16	2.5
	IEEE FP 32	0.9	0.38	2.4
×	Int 8	0.2	0.07	2.9
	Int 32	3.1	1.48	2.1
	BFloat 16	--	0.21	--
	IEEE FP 16	1.1	0.34	3.2
	IEEE FP 32	3.7	1.31	2.8
SRAM	8 KB SRAM	10	7.5	1.3
	32 KB SRAM	20	8.5	2.4
	1 MB SRAM ¹	100	14	7.1
	GeoMean ¹	--	--	2.6
DRAM		Circa 45 nm	Circa 7 nm	
	DDR3/4	1300 ²	1300 ²	1.0
	HBM2	--	250-450 ²	--
	GDDR6	--	350-480 ²	--



Lessons from DSA

- **Leverage prior compiler optimization**
 - Many DSAs rely on VLIW including TPUs
 - XLA (Accelerated Linear Algebra) compiler
 - XLA raises the TPU by 2.2 X compared to the same compiler 20 months ago
 - C compilers improve general purpose code 1 – 2% annually
 - Good compilers are critical to a DSA's success





Lessons from DSA

- **Some inference applications need floating point arithmetic**
 - **Quantized arithmetic** grants area and power savings
 - But may reduce quality, delayed deployment and some apps don't work well when quantized
- **Production inference needs multi-tenancy**
 - **Sharing can lower costs and reduce latency** if applications use many models
 - Multi-tenancy suggests **fast DRAM for DSAs**, since all weights can't fit in SRAM



Lessons from DSA

- DNN workloads evolve with DNN breakthroughs**

- MLP drops (65% to 25%)
- BERT appeared in 2018, yet it's already 28% of the workload
- A transformer encode + LSTM decoder (RNN0) + a wave RNN (RNN1) is 29%
- The importance of **programmability and flexibility** for inference DSAs to track DNN progress

Name	Avg. Size (MB)	Max Size (MB)	Multi-tenancy?	Avg. Number of Programs (StdDev), Range	% Use 2016/2020
MLP0	580	2500	Yes	27 (± 17), 1-93	61%-25%
MLP1	90	N.A.	Yes	5 (± 0.3), 1-5	
CNN0	60	454	No	1	5%-18%
CNN1	120	680	Yes	6 (± 10), 1-34	
RNN0	1300	1300	Yes	13 (± 3), 1-29	0%-29%
RNN1	120	400	No	1	
BERT0	3000	3000	Yes	9 (± 2), 1-14	0%-28%
BERT1	90	N.A.	Yes	5 (± 0.3), 1-5	



Lessons from DSA

- **DNNs grow ~1.5X per year in memory and compute**
 - DNNs grow as fast as Moore's Law
 - This rate suggests architects should provide headroom so DSAs can remain useful over their full lifetime

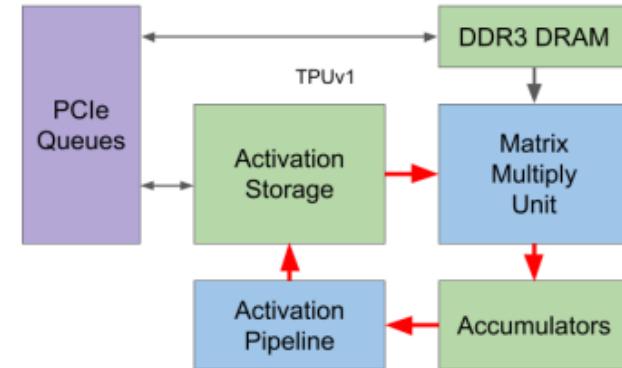
<i>Model</i>	<i>Annual Memory Increase</i>	<i>Annual FLOPS Increase</i>
CNN1	0.97	1.46
MLP1	1.26	1.26
CNN0	1.63	1.63
MLP0	2.16	2.16



Tensor Processing Unit (TPU)

- **TPU v1**

- Google's first DNN DSA
- Handle **inference (serving)**
- The **systolic array MXU has 64K 8-bit integer Multiply Accumulate (MAC) units**
- The CPU exchanges over PCIe
 - Model inputs and outputs
 - instructions
- Perf/Watt compared to GPUs and CPUs
 - 30 – 80 X higher



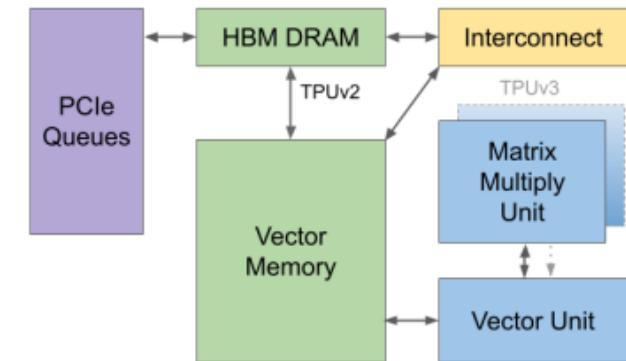
Feature	TPUv1
Peak TFLOPS / Chip	92 (8b int)
First deployed (GA date)	Q2 2015
DNN Target	Inference only
Network links x Gbits/s / Chip	--
Max chips / supercomputer	--
Chip Clock Rate (MHz)	700
Idle Power (Watts) Chip	28
TDP (Watts) Chip / System	75 / 220
Die Size (mm ²)	< 330
Transistors (B)	3
Chip Technology	28 nm
Memory size (on-/off-chip)	28MB / 8GB
Memory GB/s / Chip	34
MXU Size / Core	1 256x256
Cores / Chip	1
Chips / CPUHost	4



Tensor Processing Unit (TPU)

- **TPU v2**

- Addresses **training**
- Merge activation storage and the accumulators into **a single vector memory**
- A more programmable **vector unit**
- Support **Bfloat16** with 16 K MAC units (1/4 of the TPUs's size)
- The MXU was attached to the vector unit as **a matrix co-processor**
- High **HBM DRAM** bandwidth keeps TPUs's core well utilized
- TPUs's fetches its own **322-bit VLIW instructions from a local memory** rather than the host memory





Tensor Processing Unit (TPU)

- **TPUv2**

- Add a **chip-to-chip interconnect fabric (ICI)** enable up to 256 chips
- **Two TensorCores per chip**
- Prevent the excessive latency
 - Two small cores per chip vs.
 - A single large full-chip core

- **TPUv3**

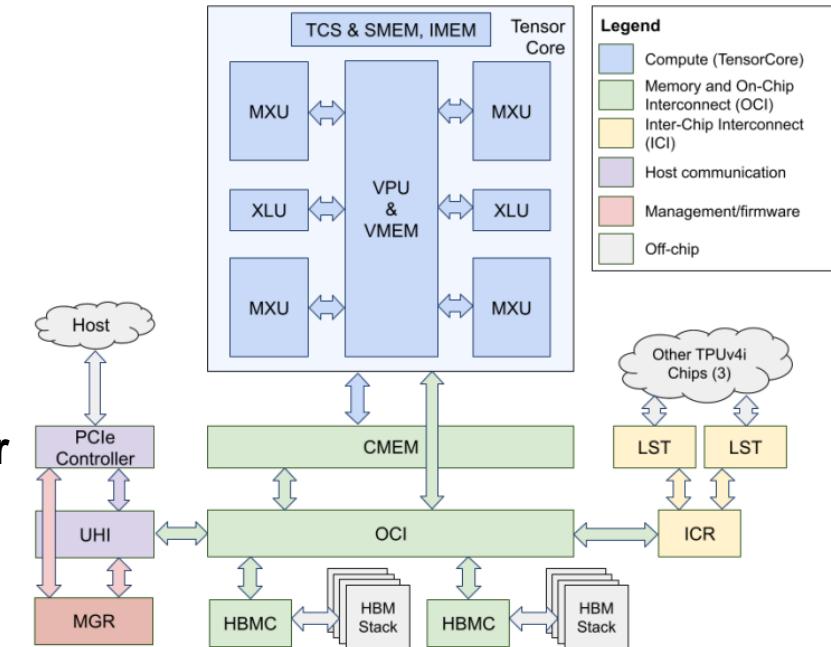
- Has 2X the number of MXUs and HBM capacity
- 1024 chips

Feature	TPUv1	TPUv2
Peak TFLOPS / Chip	92 (8b int)	46 (bf16)
First deployed (GA date)	Q2 2015	Q3 2017
DNN Target	Inference only	Training & Inf.
Network links x Gbits/s / Chip	--	4 x 496
Max chips / supercomputer	--	256
Chip Clock Rate (MHz)	700	700
Idle Power (Watts) Chip	28	53
TDP (Watts) Chip / System	75 / 220	280 / 460
Die Size (mm ²)	< 330	< 625
Transistors (B)	3	9
Chip Technology	28 nm	16 nm
Memory size (on-/off-chip)	28MB / 8GB	32MB / 16GB
Memory GB/s / Chip	34	700
MXU Size / Core	1 256x256	1 128x128
Cores / Chip	1	2
Chips / CPUHost	4	4



Tensor Processing Unit (TPU)

- **TPUv4i** (i means inference)
 - Add **128 MB common memory**
 - A large data structure don't fit in vector memory
 - **Tensor DMA engine**
 - Fully decode and execute TensorCore DMA instructions
 - Enable **512B-granular 4D tensor** memory transfers between any pair of architectural memories
 - **Unified DMA engine** across local, remote and host transfer





Tensor Processing Unit (TPU)

- **TPUv4i**
 - **Custom on-chip interconnect (OCI)**
 - The increase of memory bandwidth and the number of components
 - A point-to-point approach becomes too expensive -> significant routing resources/die area
 - A shared OCI connects all components on the die
 - **Wider data path**
 - 512B native access size instead of 64B cache lines
 - HBM bandwidth per core is 1.3X increased over TPUv3
 - NUMA memory system – use (spatial locality and bisection bandwidth)
 - Physically partitioned into four 128B-wide groups to optimize HBM accesses



Tensor Processing Unit (TPU)

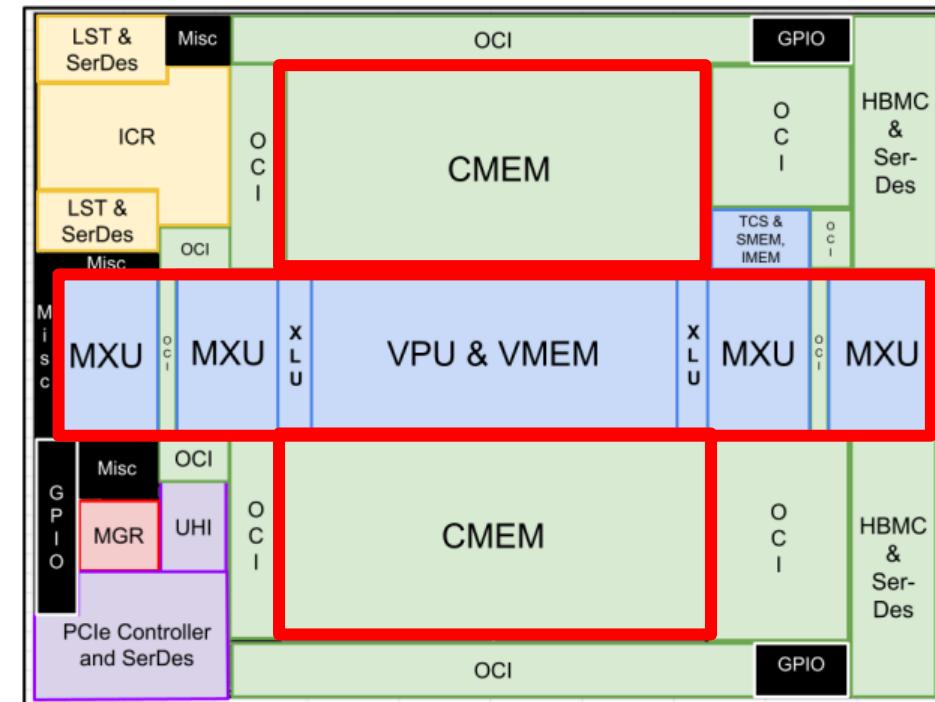
- **TPUv4i**
 - **Arithmetic unit**
 - The **VLIW instruction needs extra fields** to handle the four MXUs and CMEM scratchpad memory -> 25% wider than TPUv3
 - Sums groups of four multiplication results together
 - Adds them to previous partial sum with a series of 32 two-input adders
 - **A four-input floating point adder**
 - Cuts the critical path through the systolic array
 - The four-input adder saves 40% area and 25% power to a series 128 two-input adders



Tensor Processing Unit (TPU)

- **TPUv4i**

- The die is < 400 mm²
- **CMEM is 28% of the area**
- OCI blocks are filled the space in the abutted floorplan
- The die dimensions and overall layout are dominated by the TensorCore, CMEM, and SerDes





Tensor Processing Unit (TPU)

Jouppi et al. ISCA, 2021

Feature	TPUv1	TPUv2	TPUv3	TPUv4i	NVIDIA T4
Peak TFLOPS / Chip	92 (8b int)	46 (bf16)	123 (bf16)	138 (bf16/8b int)	65 (ieee fp16)/130 (8b int)
First deployed (GA date)	O2 2015	O3 2017	O4 2018	O1 2020	O4 2018
DNN Target	Inference only	Training & Inf	Training & Inf	Inference only	Inference only
Network links x Gbits/s / Chip	--	4 x 496	<u>4 x 656</u>	<u>2 x 400</u>	--
Max chips / supercomputer	--	256	1024	--	--
Chip Clock Rate (MHz)	700	700	940	1050	585 / (Turbo 1590)
Idle Power (Watts) Chip	28	<u>53</u>	<u>84</u>	<u>55</u>	36
TDP (Watts) Chip / System	75 / 220	<u>280 / 460</u>	<u>450 / 660</u>	<u>175 / 275</u>	70 / 175
Die Size (mm ²)	< 330	<u>< 625</u>	<u>< 700</u>	<u>< 400</u>	545
Transistors (B)	3	<u>9</u>	<u>10</u>	<u>16</u>	14
Chip Technology	28 nm	<u>16 nm</u>	16 nm	<u>7 nm</u>	12 nm
Memory size (on-/off-chip)	28MB / 8GB	<u>32MB / 16GB</u>	32MB / 32GB	<u>144MB / 8GB</u>	18MB / 16GB
Memory GB/s / Chip	34	<u>700</u>	<u>900</u>	<u>614</u>	320 (if ECC is disabled)
MXU Size / Core	1 256x256	1 128x128	2 128x128	4 128x128	8 8x8
Cores / Chip	1	<u>2</u>	2	<u>1</u>	40
Chips / CPUHost	4	4	4	<u>8</u>	8



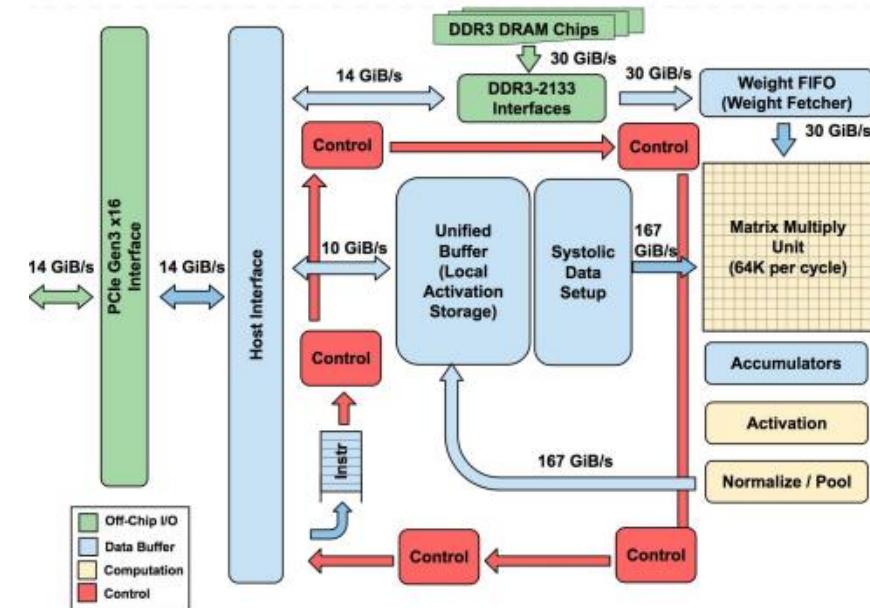
TPU Instruction Set Architectures

- TPU instruction follows the **CISC** fashion
- Average clock cycles per instructions > 10
- **No** program counter and branch instruction
- In-order issue
- SW controls buffer, pipeline synchronization
- A dozen instructions overall, five key ones
 - Read_Host_Memory
 - Read_Weights
 - MatrixMultiply/Convole
 - Activate
 - Write_Host_Memory



TPU Microarchitecture

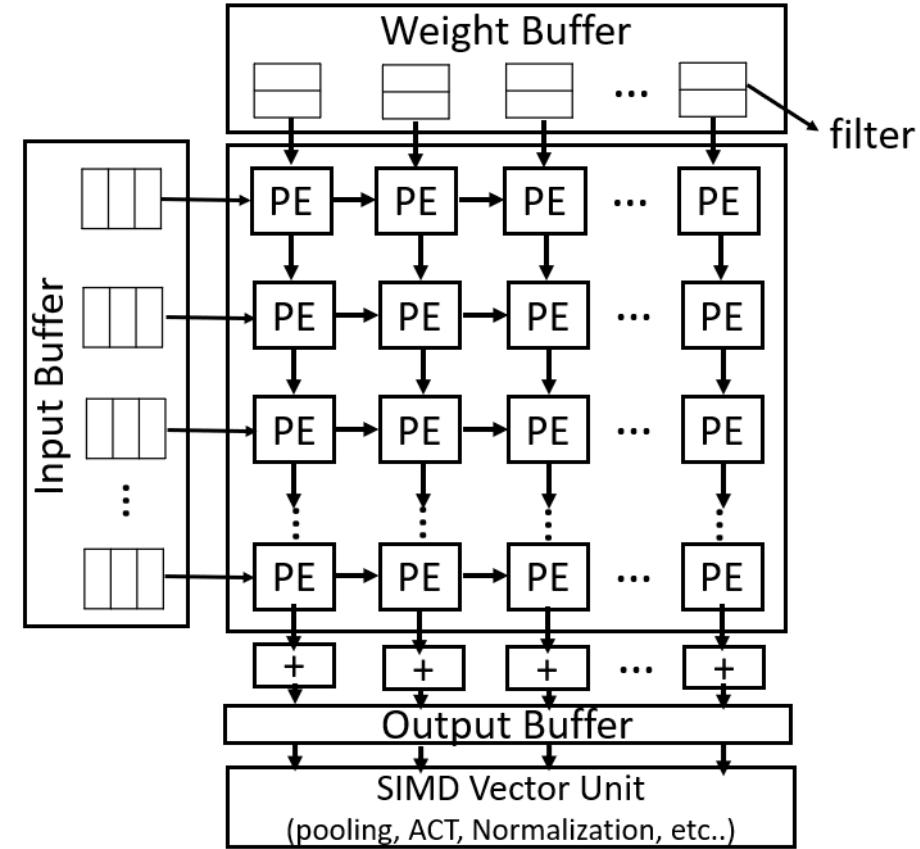
- **4-stage overlapped execution**,
1 instruction type/ stage
- Execute other instructions while
MM is busy
- Read_Weight doesn't wait for
weights fetched from DRAM
- The MM unit uses **not-ready**
signal to indicate data aren't
available in unified and Weight
FIFO buffer





TPU Micro-architecture

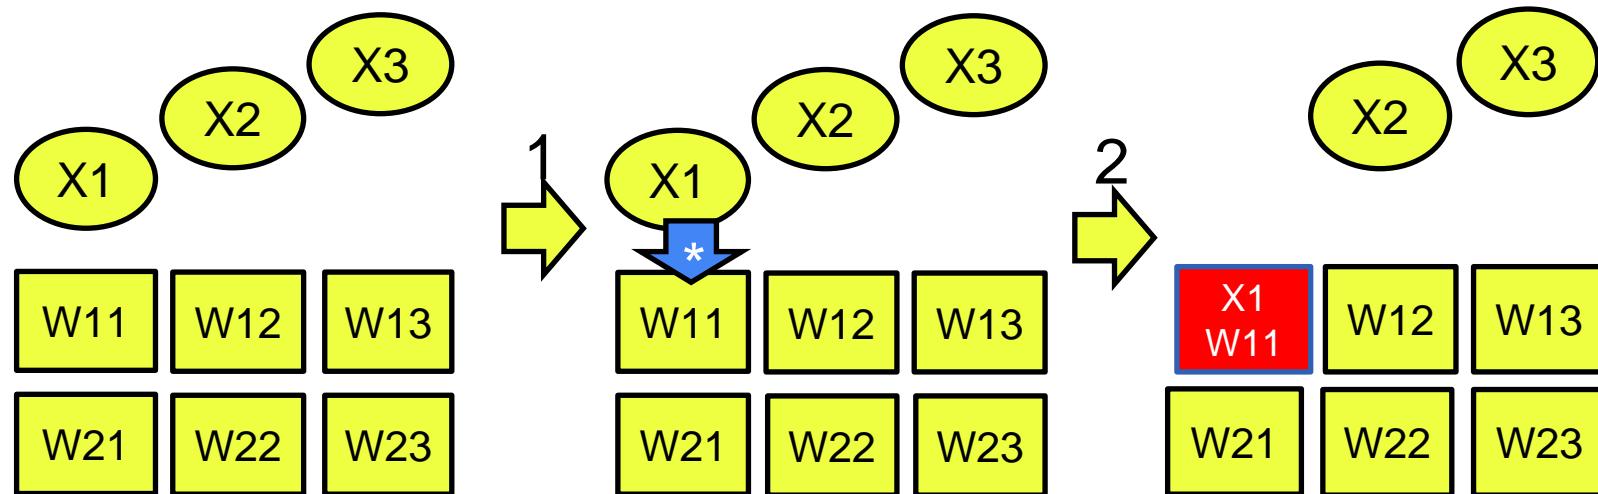
- Each PE performs Multiply-and Accumulate (MAC) operation
- The unified memory buffer is decomposed into input, weight, and output buffer
- Each weight buffer stores weights of a filter
- At each cycle, inputs are pushed in the PE horizontally
- Partial sums flow vertically





Systolic Execution in TPU

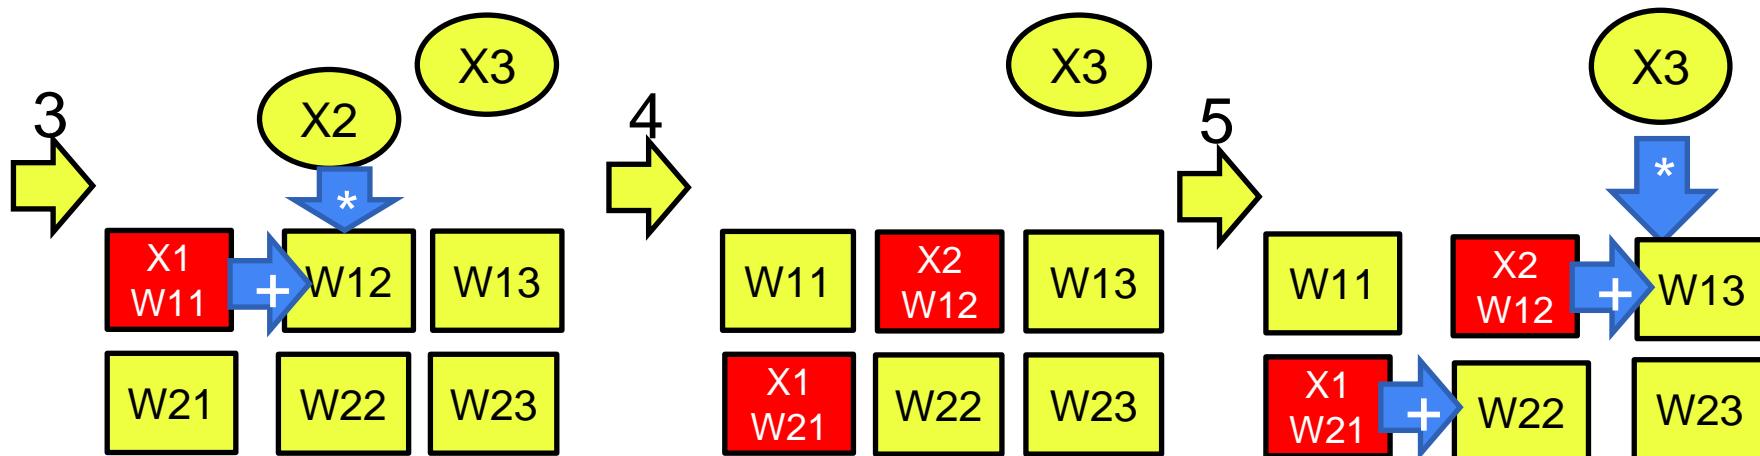
- Reading a large SRAM is much more expansive than arithmetic
- Using systolic execution to reduce R/W of the unified buffer





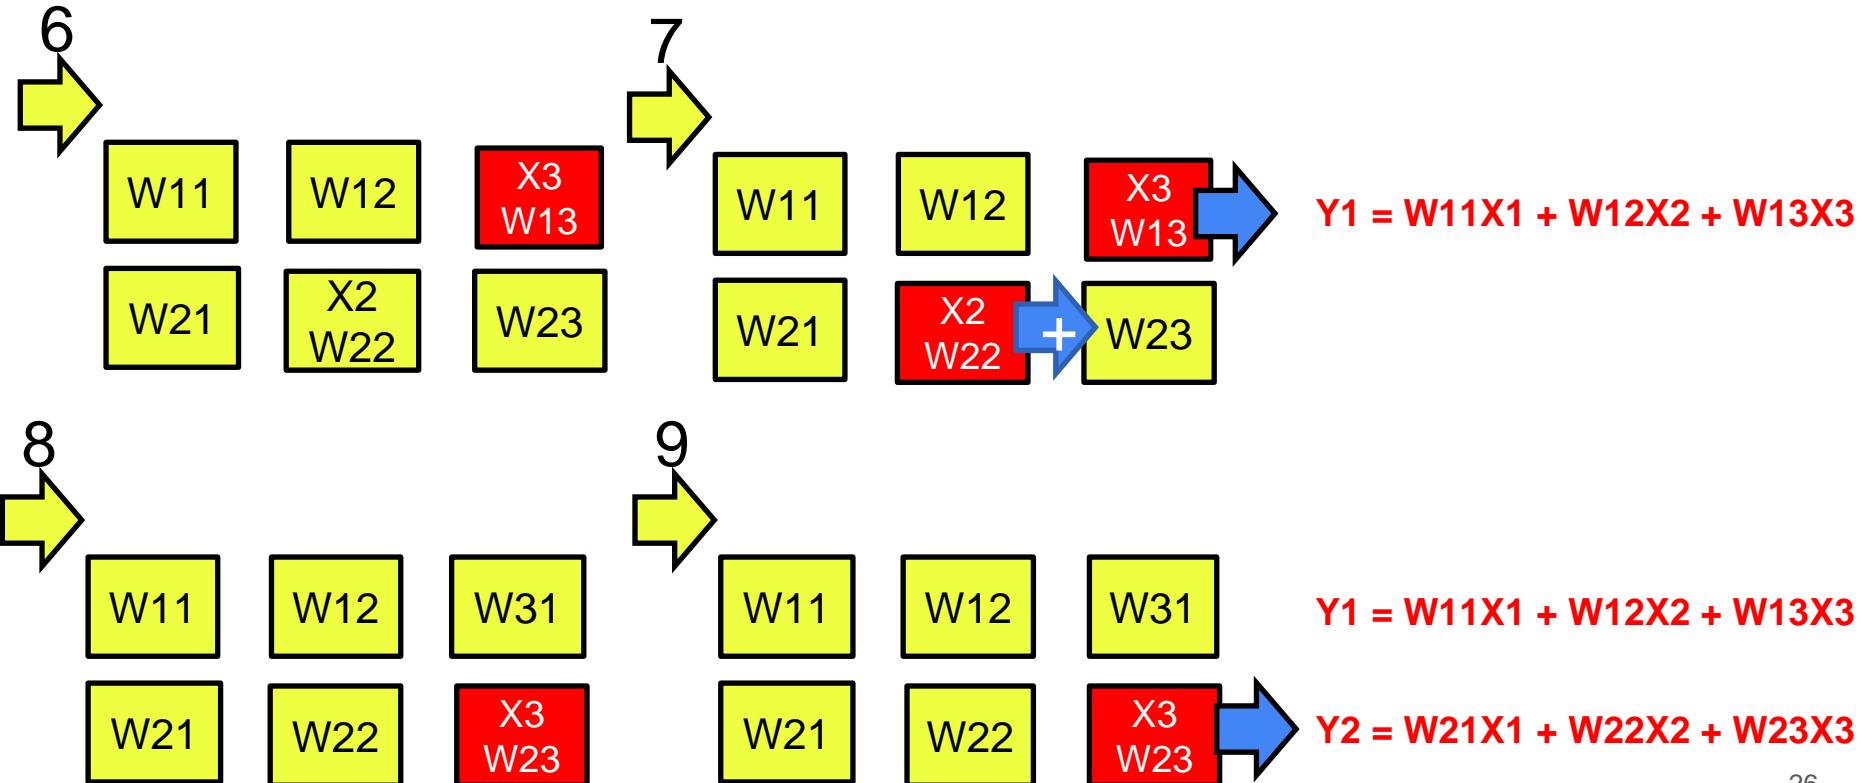
Systolic Execution in TPU

- Reuse input values
- Relies on data from different directions arriving at each array at regular interval to do the calculation





Systolic Execution in TPU





TPU Case Study

- How to map input feature map and filter (weight) to TPU ?
- Suppose the size of the input feature map is 4×4 , and the size of filter is 2×2 .

$m \times m$

a_0	a_1	a_2	a_3
a_4	a_5	a_6	a_7
a_8	a_9	a_{10}	a_{11}
a_{12}	a_{13}	a_{14}	a_{15}

Input

$$\begin{matrix} & \\ * & \end{matrix} \quad \begin{matrix} k \times k \\ \text{Filter} \end{matrix} =$$

$(m-k+1)(m-k+1)$

C_0	C_1	C_2
C_3	C_4	C_5
C_6	C_7	C_8

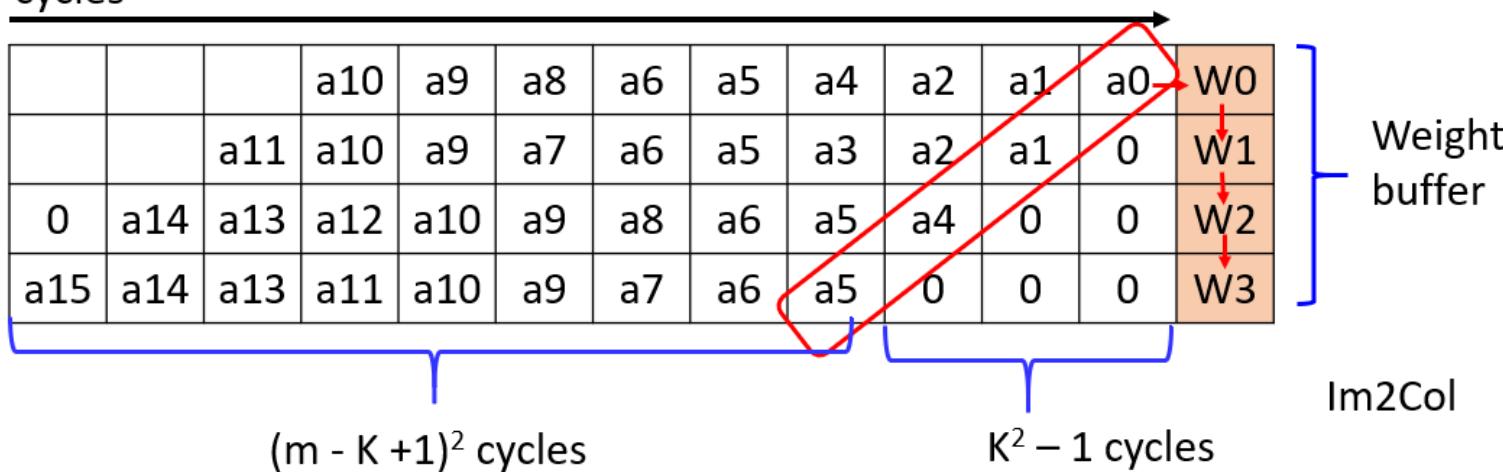
Output



TPU Case Study

- How to map input feature map and filter to TPU ?
- How many cycles takes to complete the CONV of one feature map with 2×2 filter, # of filter = 1 ?
 - $(m - K + 1)^2 + K^2 - 1 + (\# \text{ of filter} - 1)$

cycles



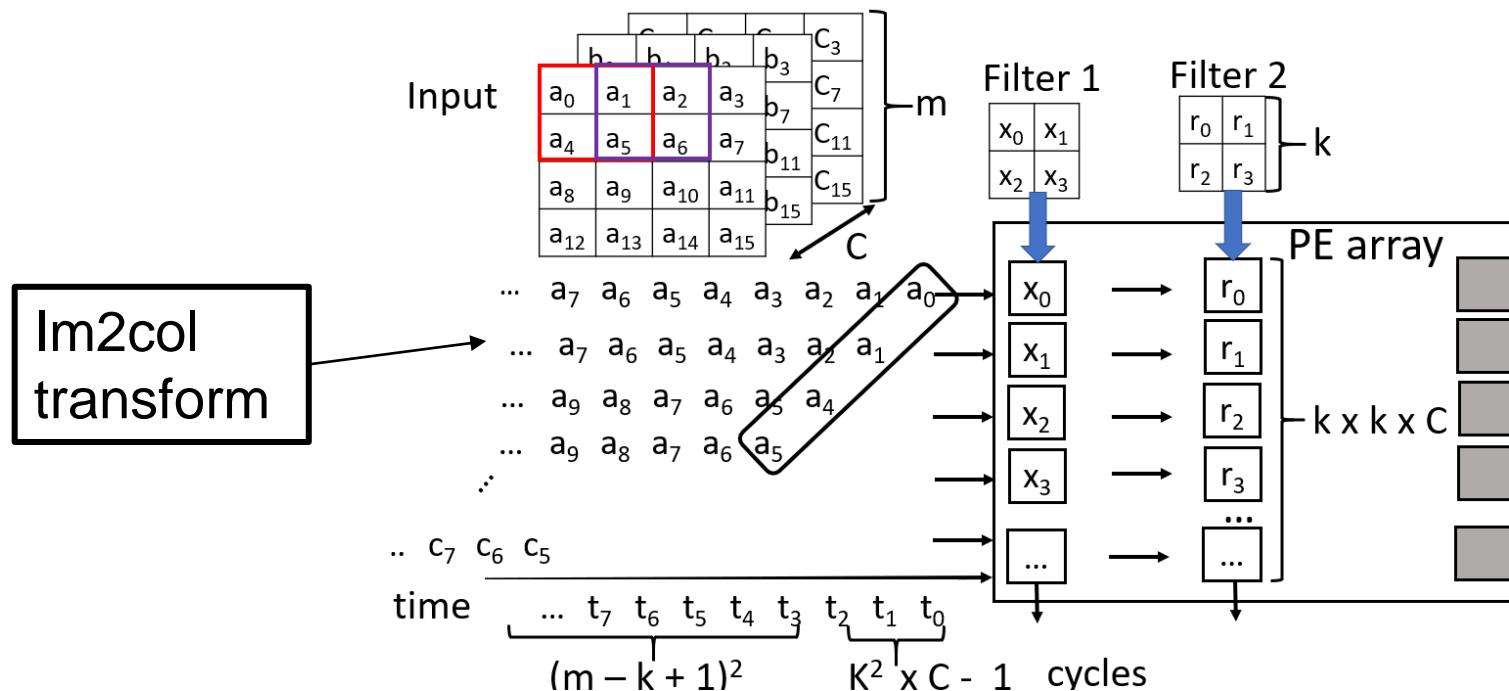
$m \times m$			
a_0	a_1	a_2	a_3
a_4	a_5	a_6	a_7
a_8	a_9	a_{10}	a_{11}
a_{12}	a_{13}	a_{14}	a_{15}

Input



TPU Case Study

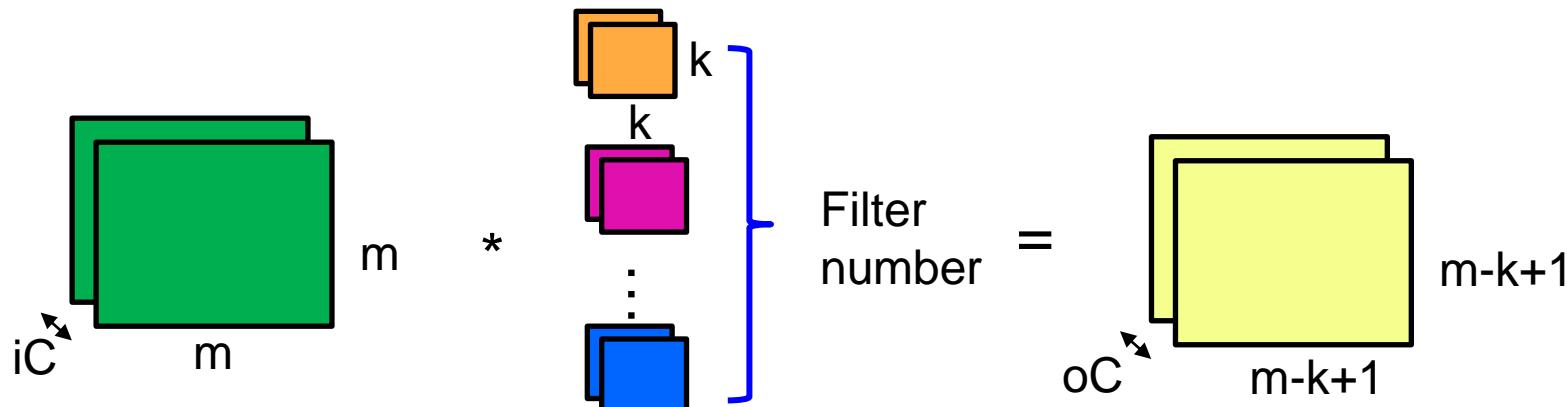
- The CONV weight stationary data flow





TPU Case Study

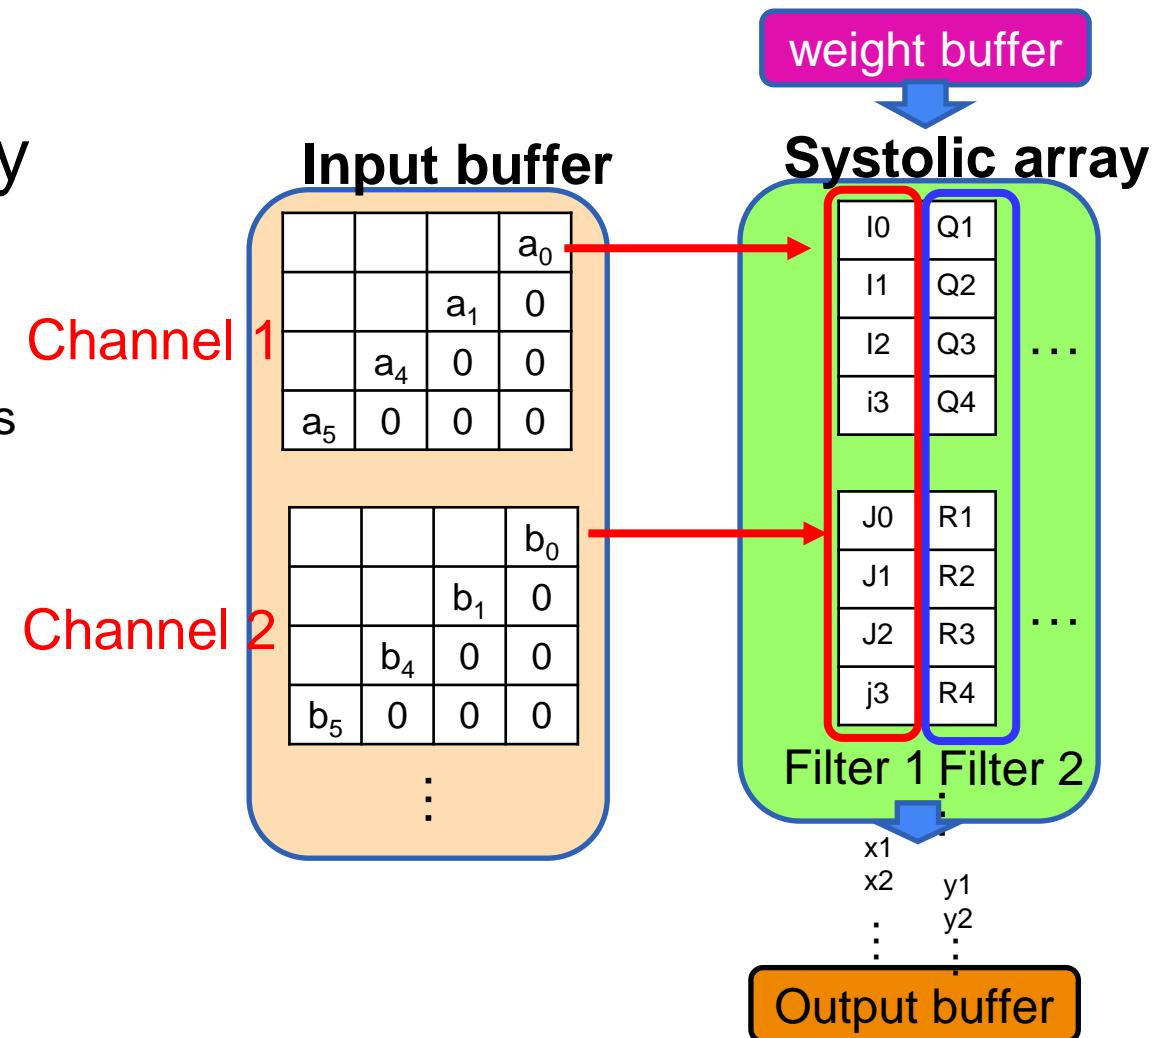
- In real-world model, a DNN model often has multiple channels and filters
- How many ops take to complete a CONV in the systolic array ?
 - $(m - k + 1) \times (m - k + 1) \times (k \times k \times iC \times oC)$





TPU Case Study

- How to map CONV to the systolic array ?
- Systolic array contains multiple PEs
- Each filter element is placed on the local buffer of each PE





TPU Case Study

- How many cycles takes to complete a CONV ?
 - Systolic array size: 128×128
 - Kernel size: 2×2
 - Input channel: 256
 - Input size: 10×10
 - The number of filter: 16
1. 128×128 systolic array can execute $\text{floor}(128/(2 \times 2)) = 32$ channels
 2. The systolic array needs to take $\text{ceil}(256/32) = 8$ times
 3. Each input takes $(10 - 2 + 1)^2 + (16 - 1) = 96$ cycles
 4. Total = $96 \times 8 + (2^2 \times 32 - 1) = 895$ cycles



Takeaway Questions

- How does TPU reduce the energy consumption ?
 - (A) Employ the weight stationery data flow
 - (B) Increase the clock frequency of PEs
 - (C) Increase the number of PEs
- Given a DNN layer with 2×2 filter with a single channel, how many cycles will take before activate the first row of the systolic array?
 - (A) 3
 - (B) 4
 - (C) 5

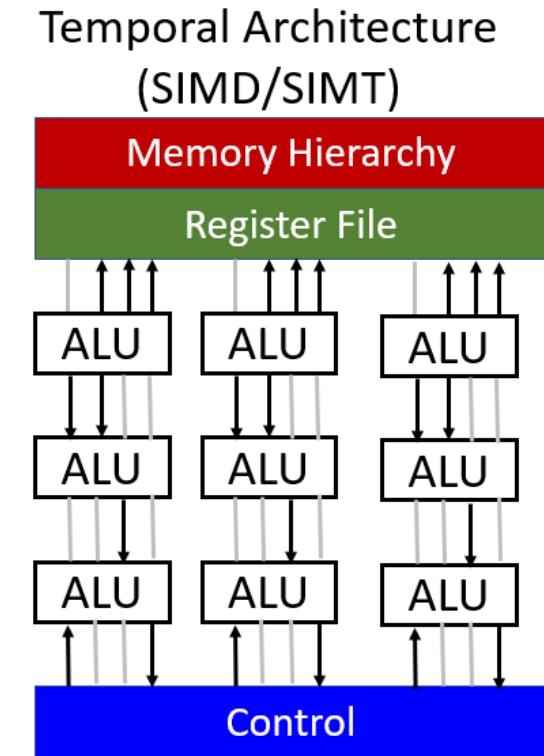


Dataflow DNN Accelerator



Design Aspects of Temporal Accelerator (TA)

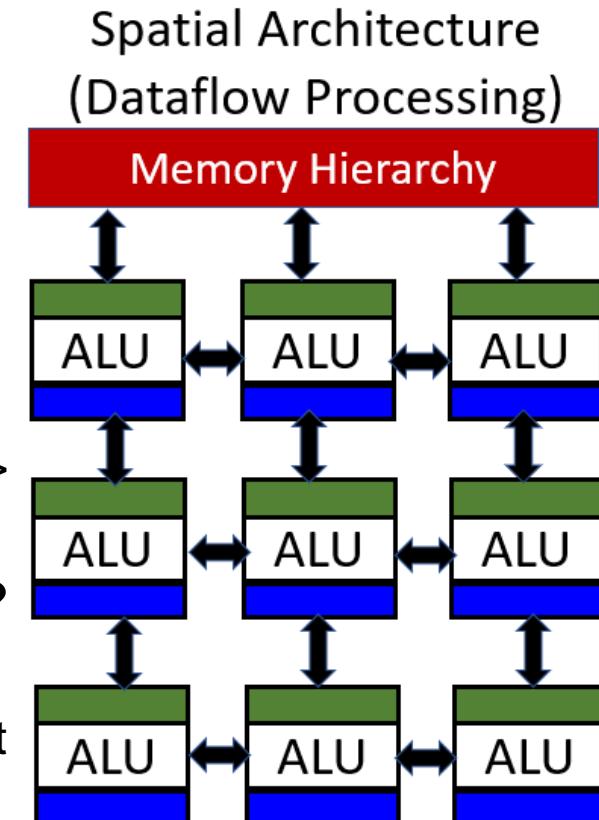
- Centralized control for ALUs
- ALUs can only fetch data from the memory hierarchy
- ALUs “cannot” communicate directly with each other
- Why TA becomes popular? Parallelism
- Design aspects for DNN workloads
 - **Reduce # of multiplication** -> increase throughput
 - **Ordered computation (tiling)** -> improve memory subsystem





Design Aspects of Spatial Accelerator (SA)

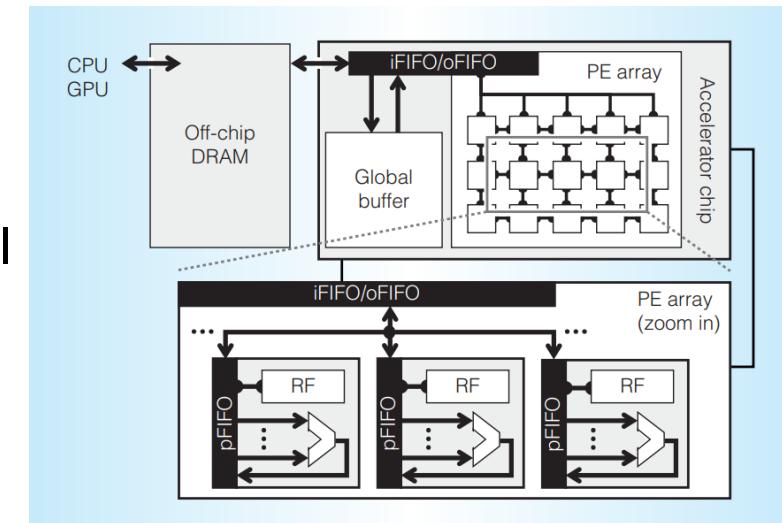
- **ALUs**
 - Can pass data from one to another directly
 - Can have its own control logics and local memory (registers)
- **Dataflow processing**
 - Programmable -> dynamic vs static graphs
 - Dynamic Mapping -> increase data reuse -> energy-efficiency
- **Why SA are popular on DNN workloads?**
 - Consume lower power & high throughput
 - Why? Data reuse -> reduce data movement





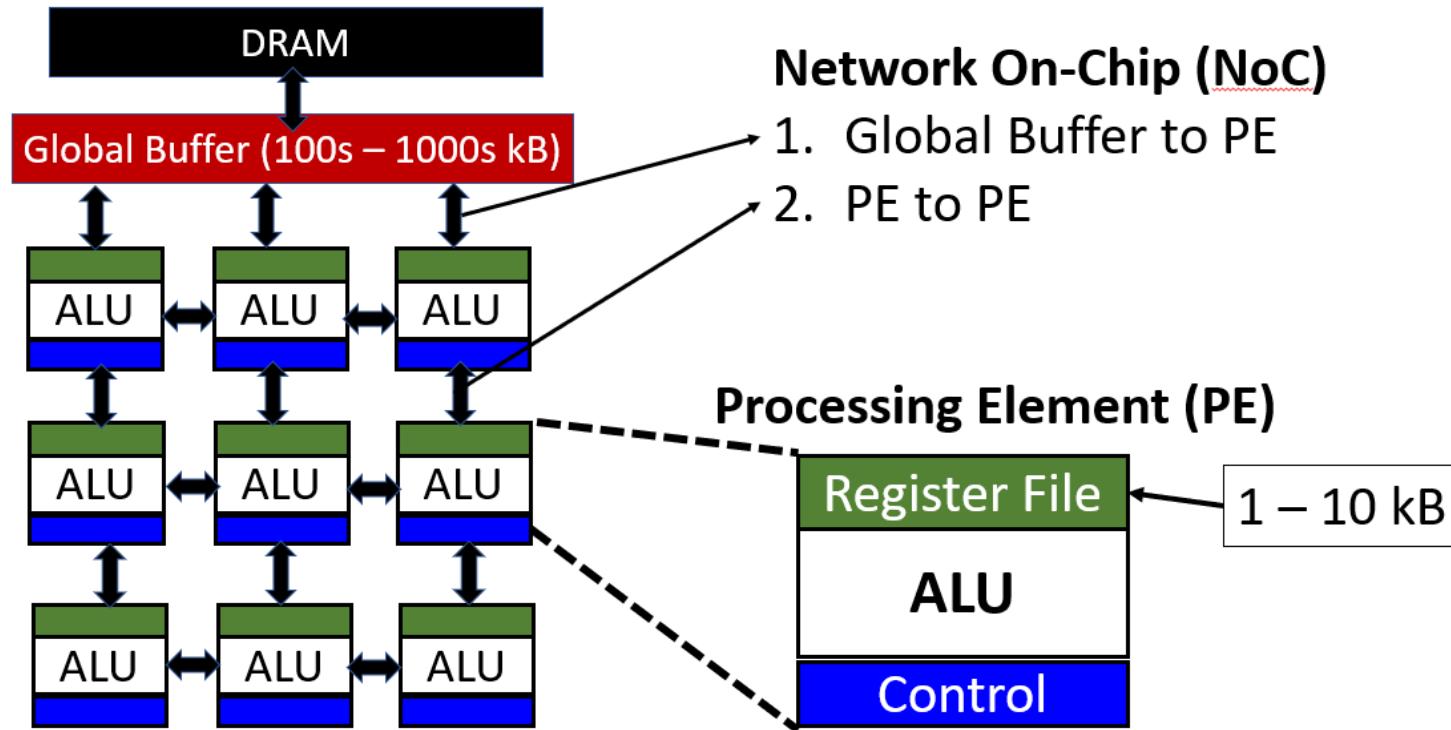
Spatial Array Architecture

- **Spatial array architecture comprises**
 - An array of processing elements (PE)
 - Off-chip DRAM
 - Global buffer
 - Network-on-chip (NOC)
 - Register file (RF) in the PE
- **Input and output FIFO (i/oFIFO)**
 - Use to communicate DRAM, global buffer, and PE
- **PE FIFO (pFIFO)**
 - Control the traffic going in and out of ALU





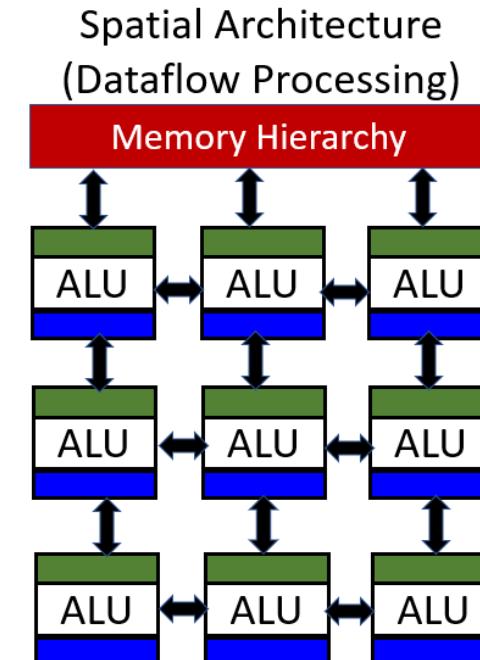
Spatial Architecture for DNN





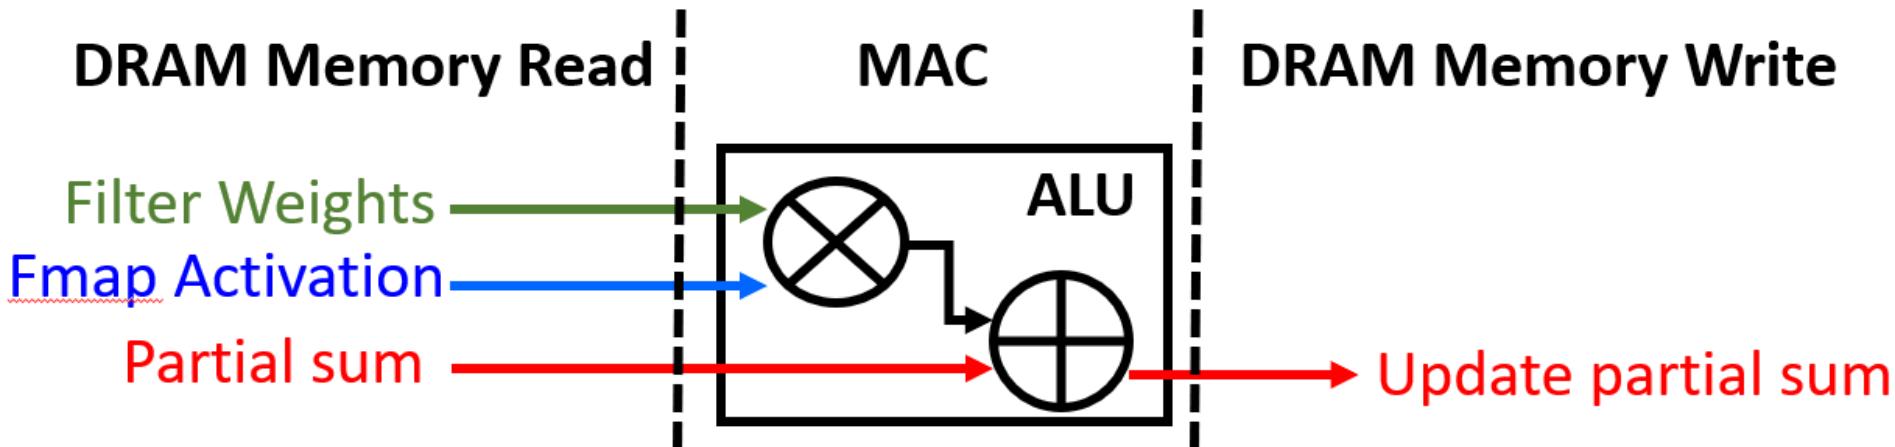
Challenges of Spatial Accelerators

- Memory access is the bottleneck
 - AlexNet has 2896M DRAM accesses required
 - How to decrease expensive DRAM accesses ?
 - Intelligent distributed data allocation
- Varying parameters in DNN models
 - Each layer has different computation volume
 - Different operations in DNN layers and models





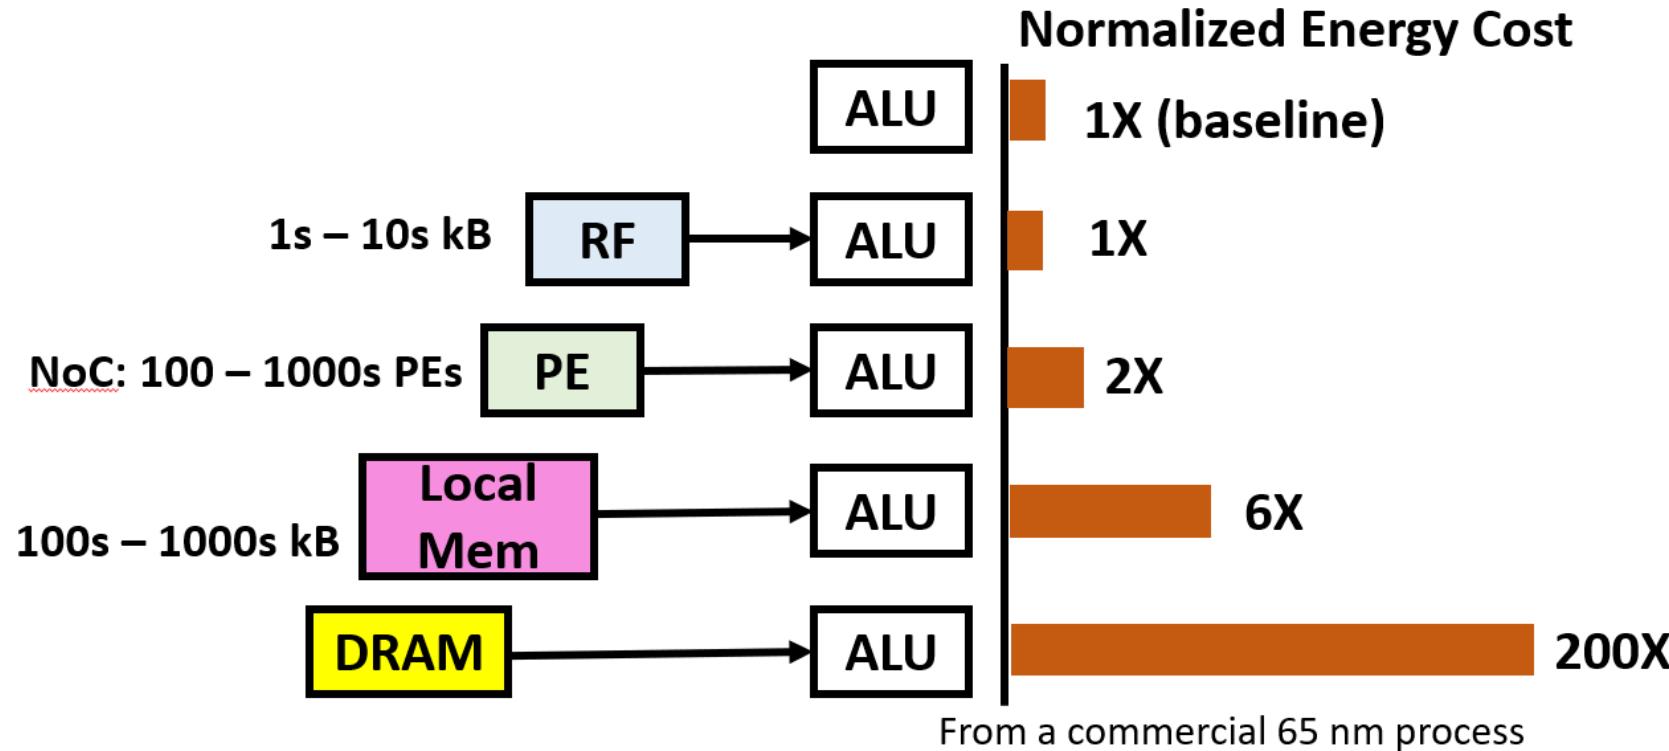
Improve Spatial Accelerator Energy-Efficiency ?



Worst Case: All memory R/W accesses from DRAM

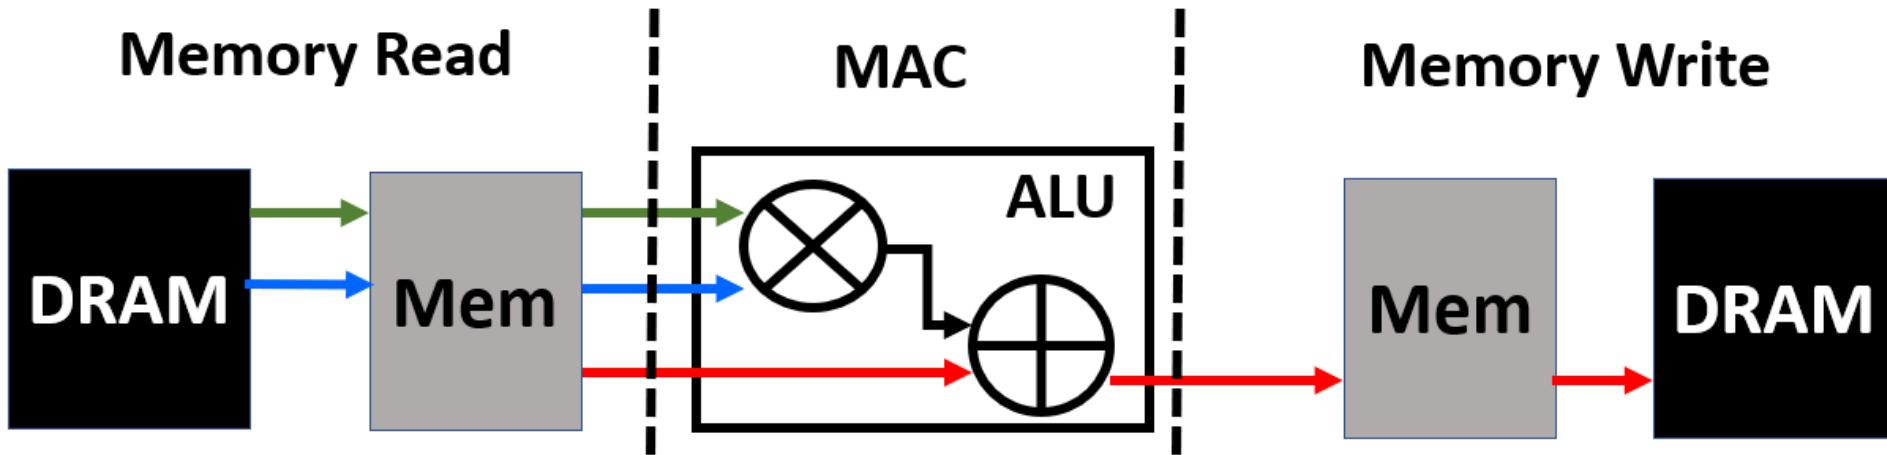


Energy Cost of Memory Access





Data Reuse on Local Memory



How to leverage local memory to reduce the times of remote DRAM access on DNN workloads ?

Optimal case: reduce **2896 M** to **61 M** DRAM accesses on AlexNet



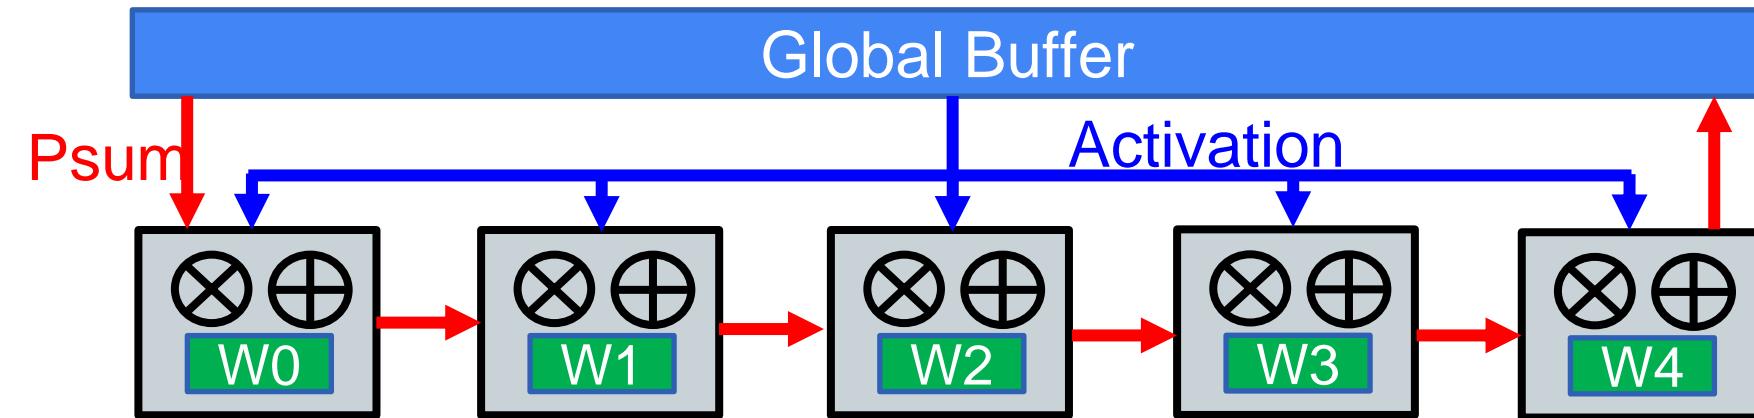
Dataflow Taxonomy

- Output Stationary (OS)
- Weight Stationary (WS)
- Input Stationary (IS)
- Dataflow
 - Indicates the matrix which is “pinned” to a given PE
 - The **ordering** of the operations
 - Data prioritization across the memory hierarchy and compute data paths



Weight Stationary (WS)

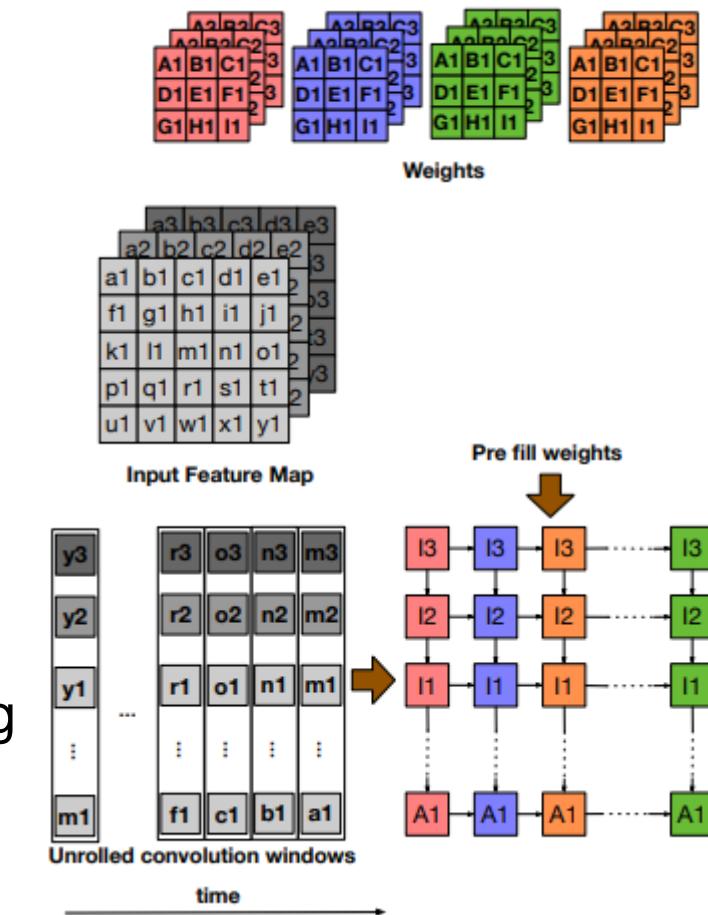
- Minimize weight read energy consumption
- Broadcast activations and accumulate psums spatially across PEs
- Each weight stays stationary in RF of each PE





Weight Stationary (WS)

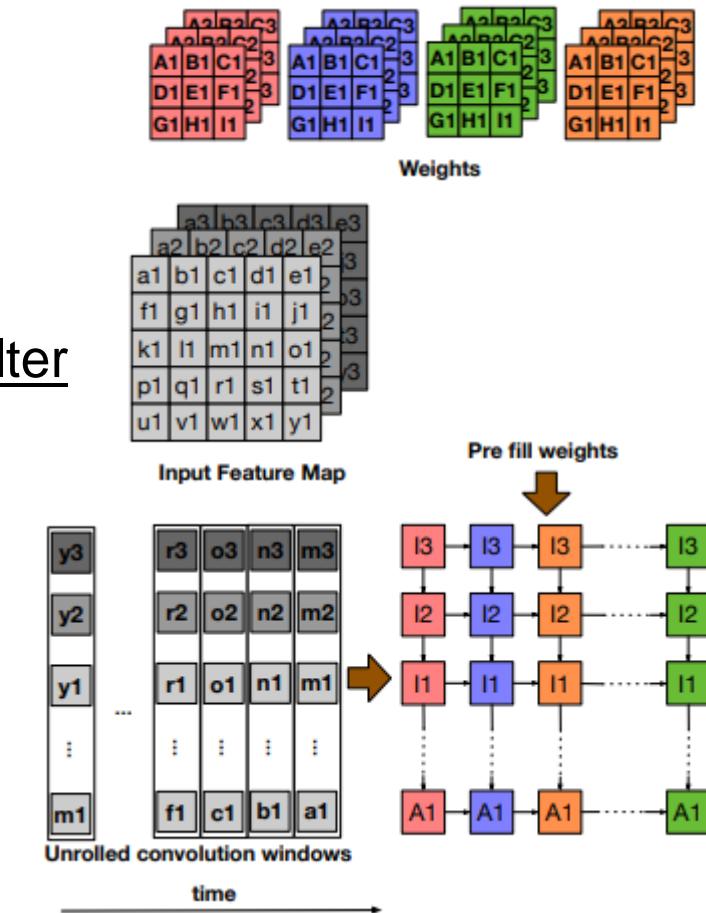
- Each element of the **weight matrix** is uniquely mapped to a given MAC unit
- Every cycle the **input elements** are multiplied with the currently mapped weights
- **Partial sums** are stored within the array
- **Reduction** takes place by communicating the partial sums across the MAC units
 - Take multiple cycles





Weight Stationary (WS)

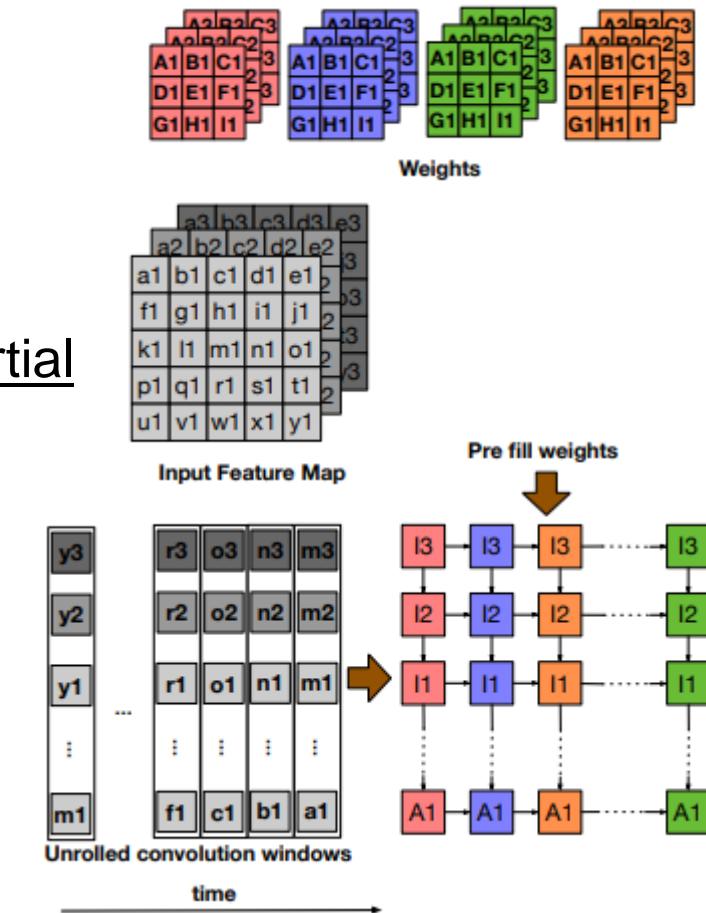
- First step in WS data mapping
 - Each column is assigned to a given filter
 - The elements of the assigned filter matrix are fed in from the top edge
 - After the filter elements are placed, the pixels of inputs are then fed in from the left edge
 - Partial sums for a given output is generated every cycle





Weight Stationary (WS)

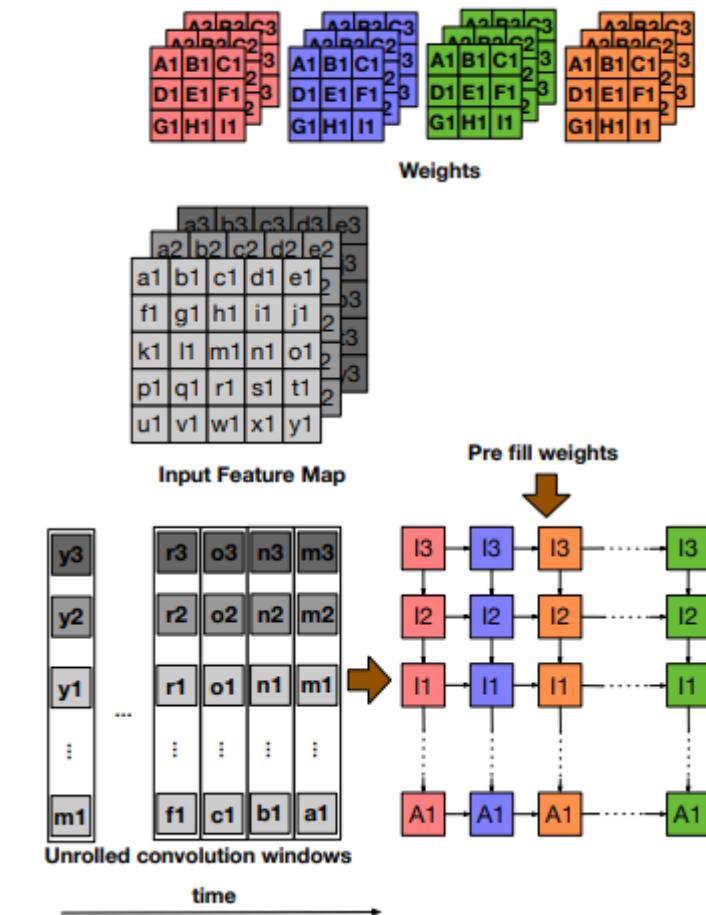
- Second step in WS data mapping
 - For a given output, corresponding partial sums are distributed over a column
 - Partial sums are reduced over the given column in next n cycles
 - n is the number of partial sums generated for a given pixel
 - Once the mapped weight are done, the mapping is replaced with new set of weights





Weight Stationary (WS)

- Shortcoming of the WS
 - Partial sums corresponding multiple outputs are required to be kept in the array until they are reduced
 - Leads to increase in implementation cost (why?)

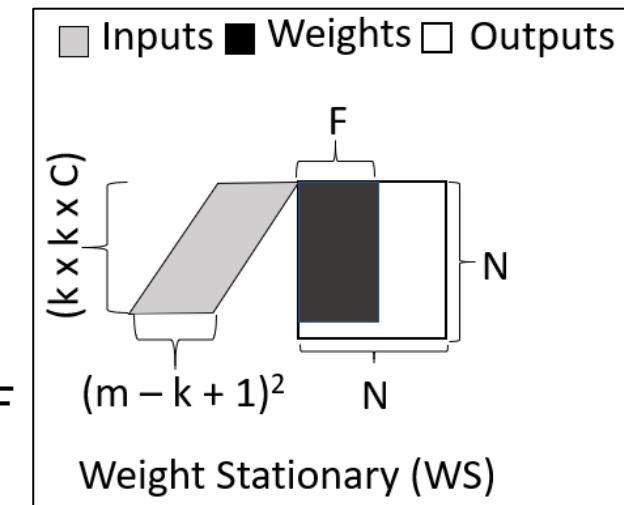




Latency Analysis of Weight Stationary

- **The weight stationary in the systolic array**

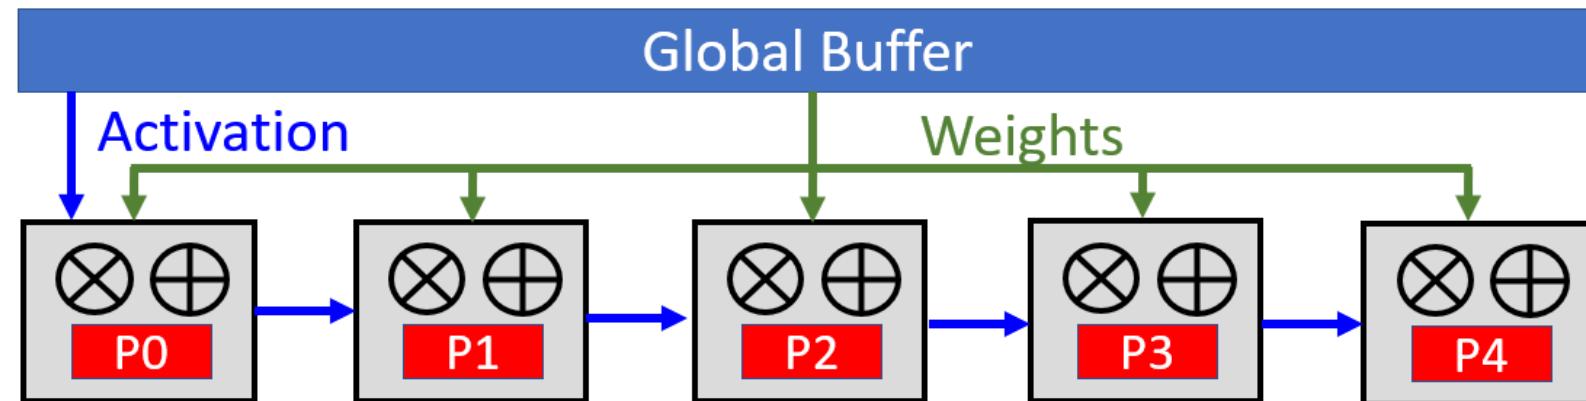
- Inputs take $(m - k + 1)^2 + (k \times k \times C - 1)$ cycles to flow in the spatial array horizontally
- Inputs also need to take F cycles to pass through each filter
- Pre-load weights take $(k \times k \times C)$ cycles
- Total cycles
 - $(m - k + 1)^2 + (k \times k \times C - 1) + (k \times k \times C) + F$





Output Stationary (OS)

- Minimize partial sum R/W energy consumption
- Keep the accumulation of psums stationary in the RF
- Stream input activations across PE array
- Broadcast the weights to all PE array from the global buffer

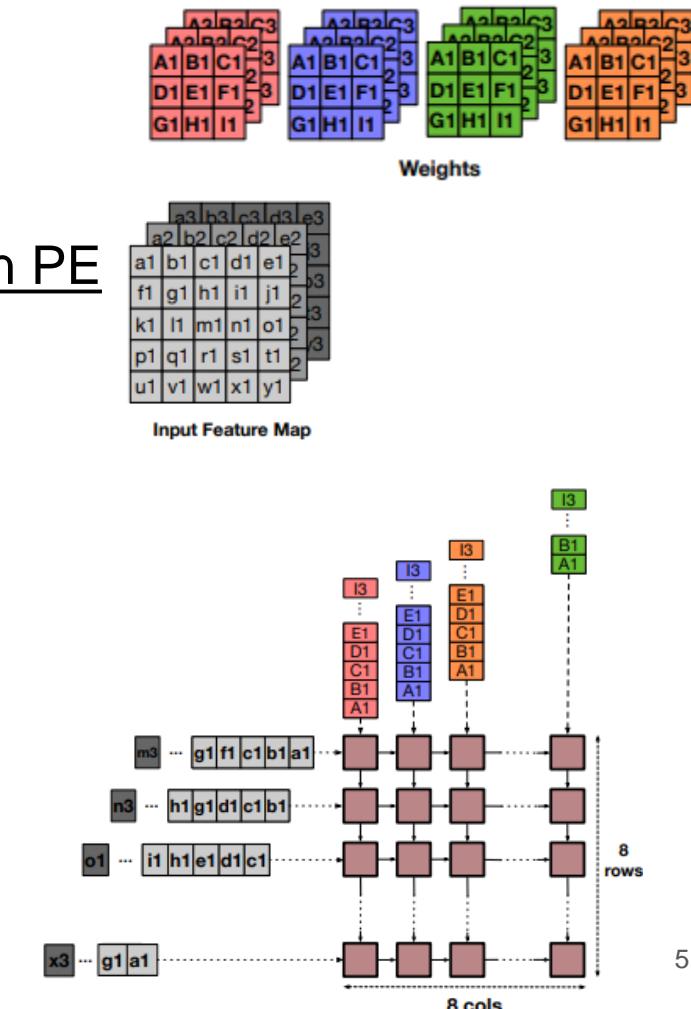




Output Stationary (OS)

- Each pixel of output is assigned to a given PE
- All compute necessary for generating the given output is done on the PE
- The input and weight are streamed in every cycle
- Reduction operation is done in place, no further communication is needed
- Once one output pixel is generated by a given PE, the result is transferred to the memory, and the PE is assigned another pixel to compute

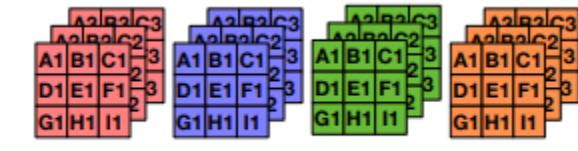
<https://arxiv.org/pdf/1811.02883>



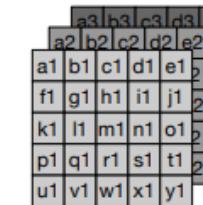


Output Stationary (OS)

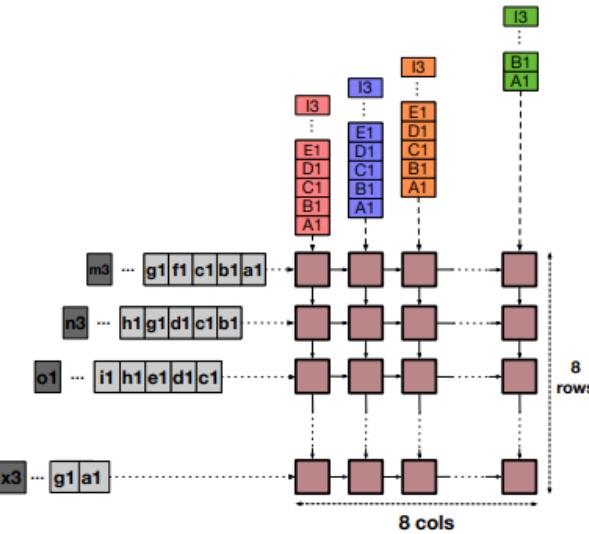
- In a given column PEs in each row
 - Generating adjacent output in a single channel
 - Each column generates pixels corresponding to different output channels
- Shortcoming of the OS
 - The data transferred overhead of generated outputs



Weights



Input Feature Map

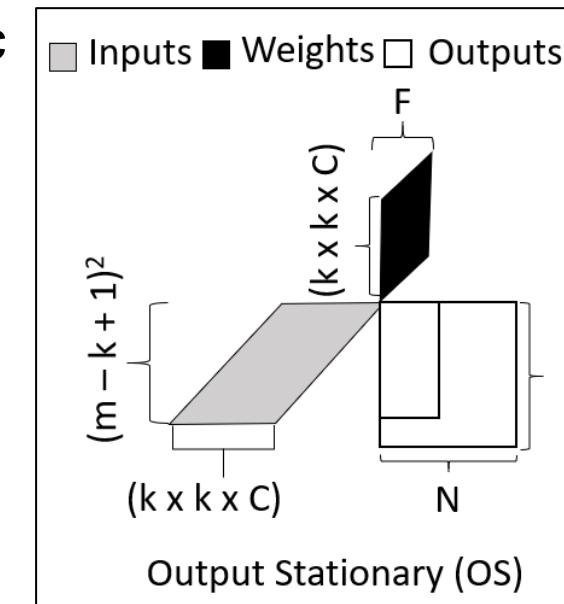




Latency Analysis of Output Stationary

- **The output stationary in the systolic array**

- Inputs and weights are pushed in the systolic array and takes $(k \times k \times C - 1) + (m - k + 1)^2$
- Taking F cycles to pass through outputs
- Outputs are accumulated in-place
- Total cycles
 - $(k \times k \times C - 1) + (m - k + 1)^2 + F$

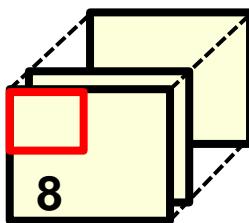
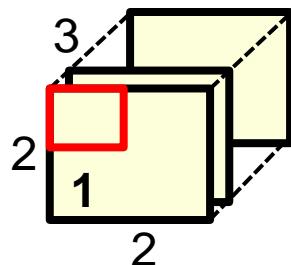




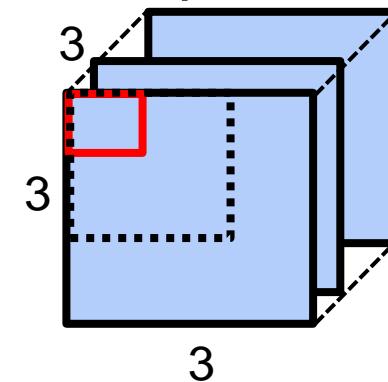
OS Dataflow Example

- Cycle through input fmap and weights (psum of output is stationary)

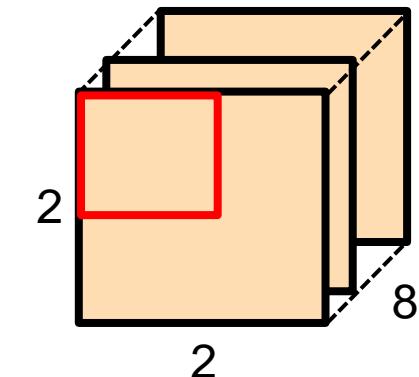
Filter



Input fmap



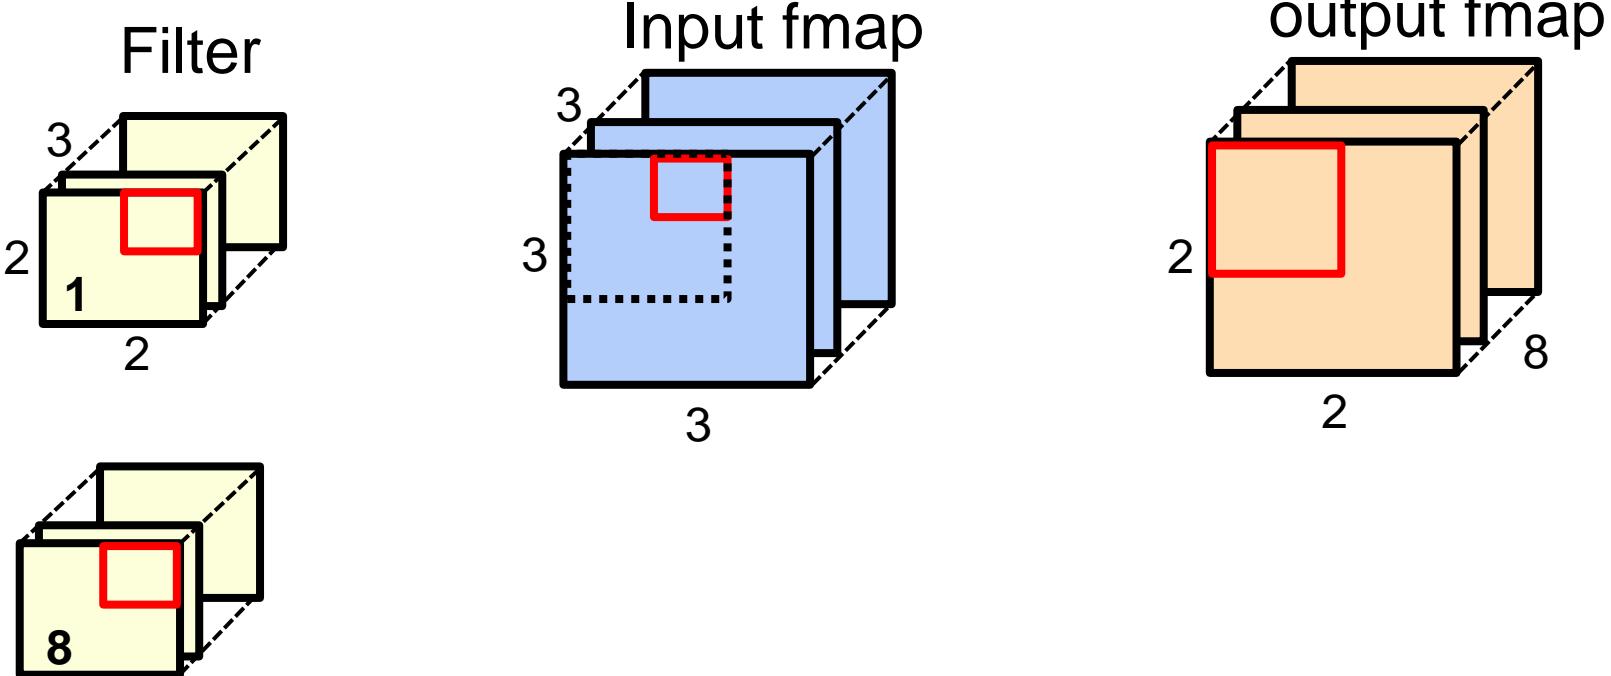
output fmap





OS Dataflow Example

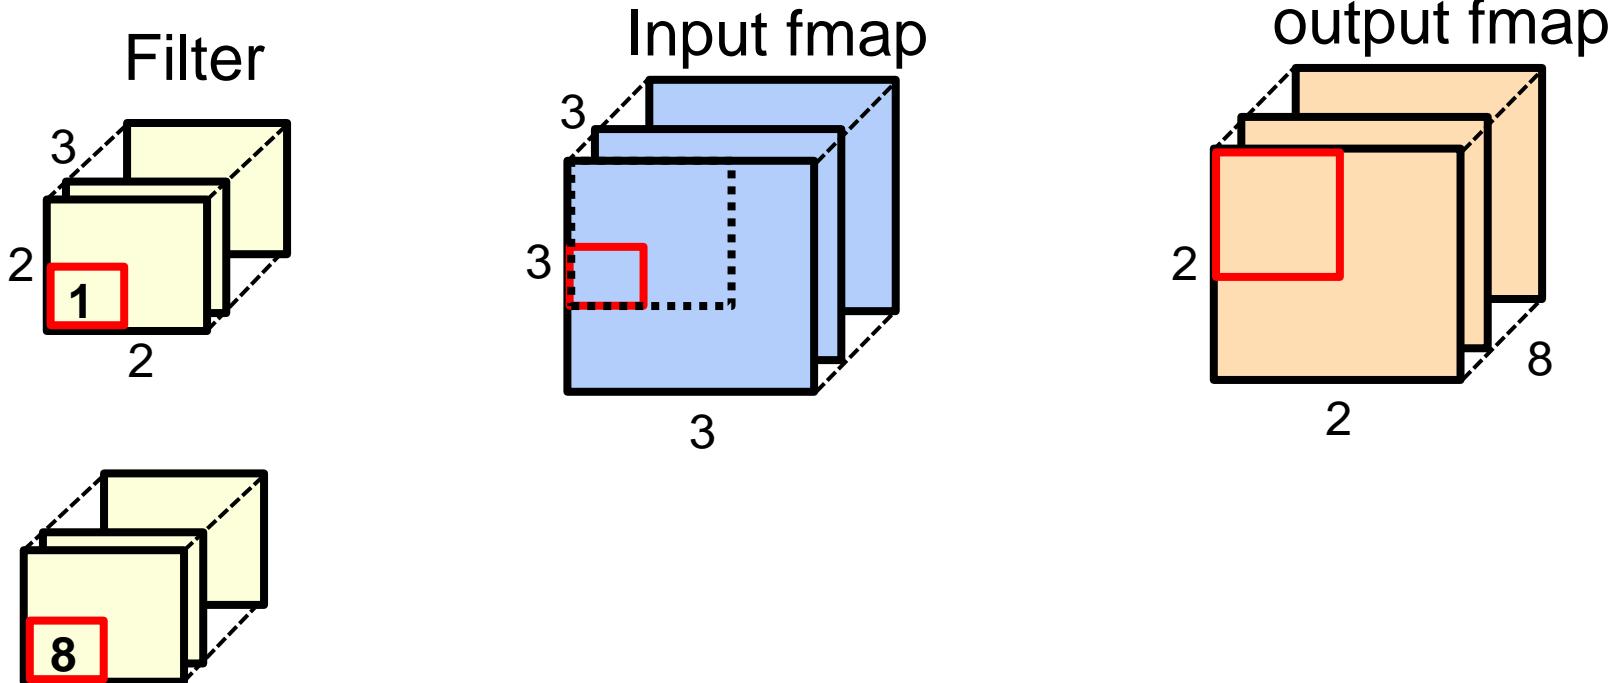
- Cycle through input fmap and weights (psum of output is stationary)





OS Dataflow Example

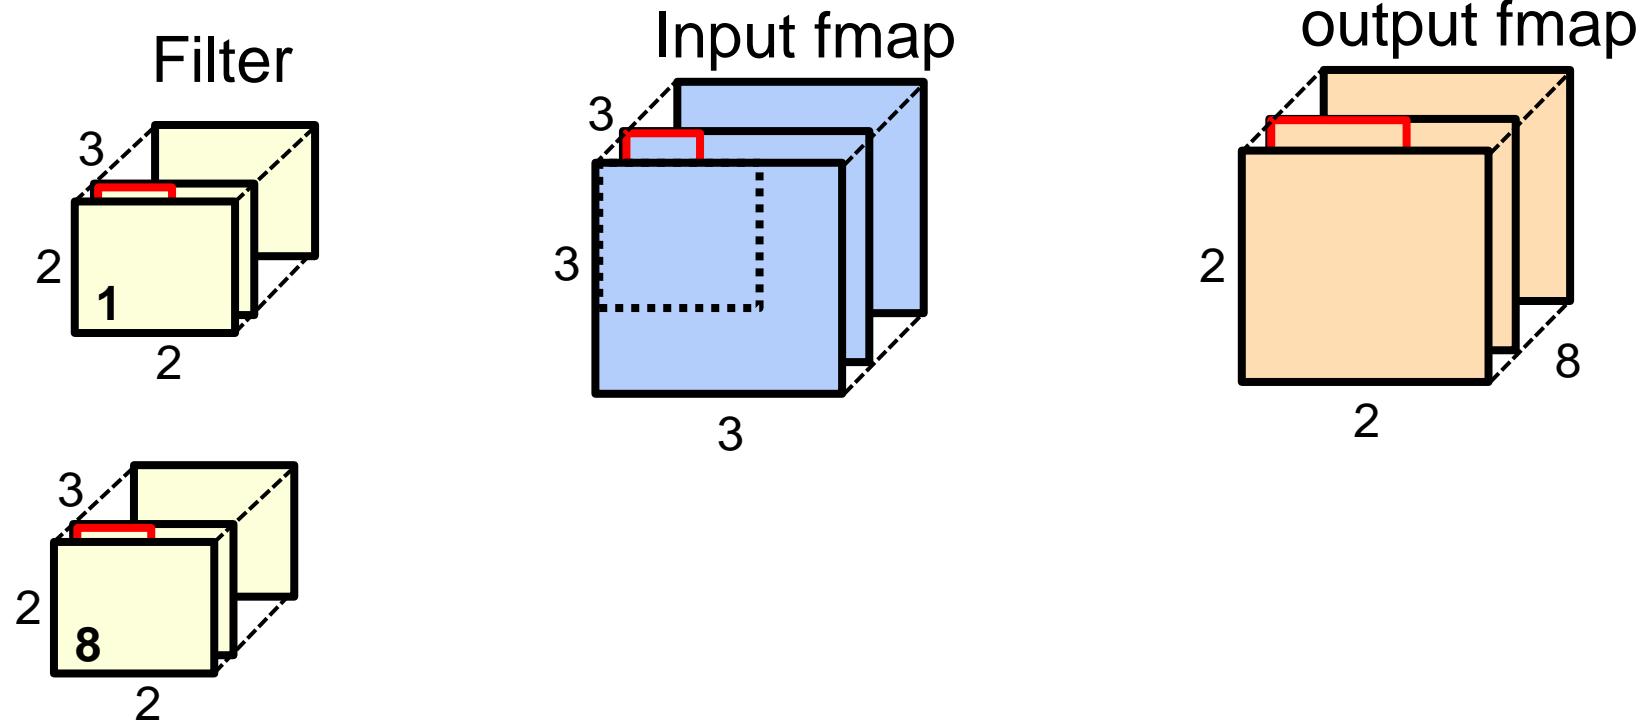
- Cycle through input fmap and weights (psum of output is stationary)





OS Dataflow Example

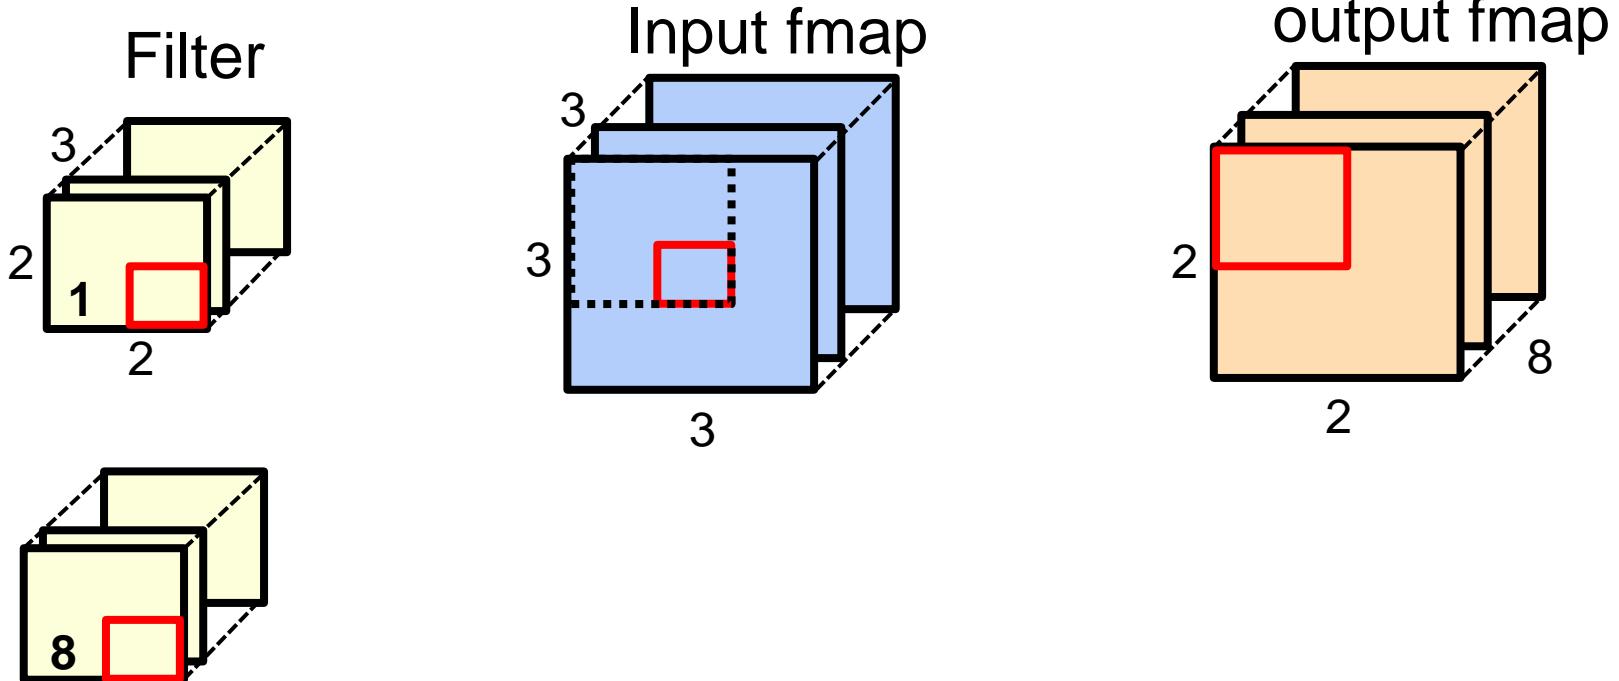
- Cycle through input fmap and weights (psum of output is stationary)





OS Dataflow Example

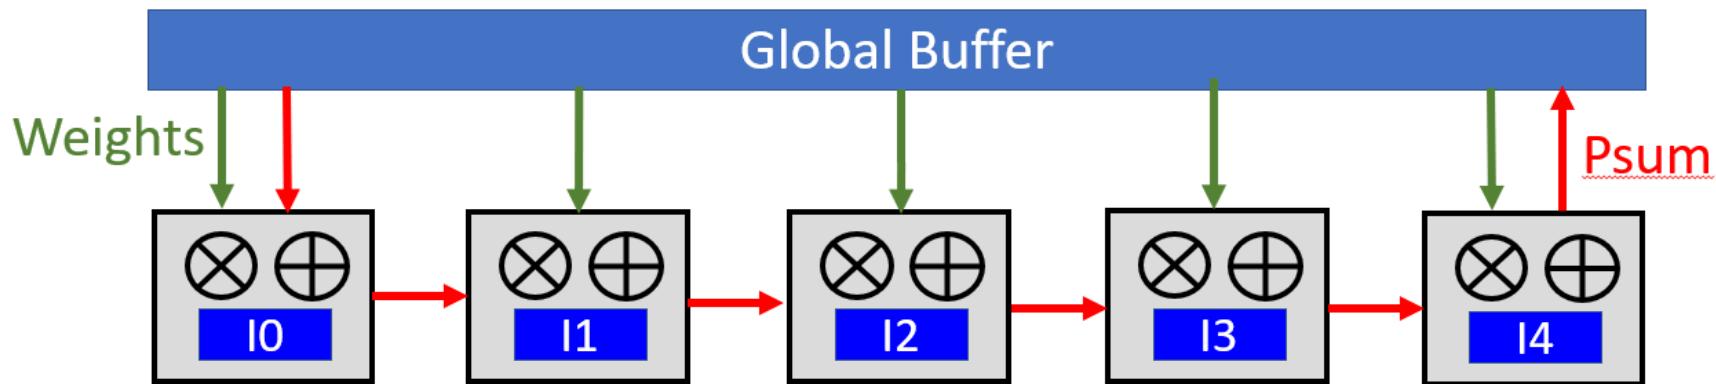
- Cycle through input fmap and weights (psum of output is stationary)





Input Stationary (IS)

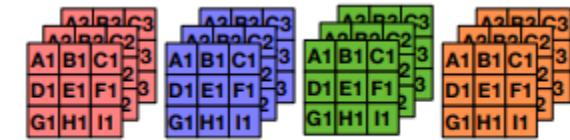
- Minimize the energy consumption of reading input activations
- Unique filter weights are uni-cast into PEs at each cycle
- Psums are spatially accumulated across PEs





Input Stationary (IS)

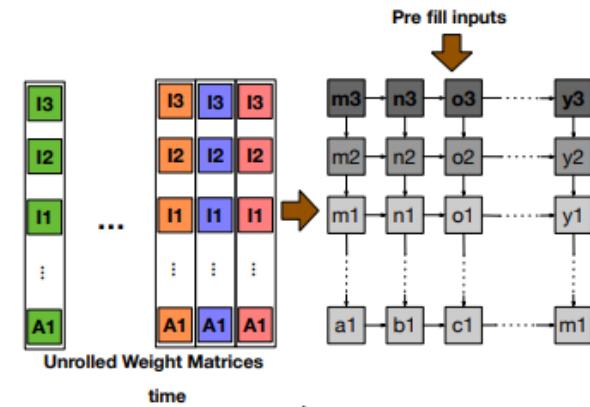
- Input feature map (IFMAP) are “pinned” with the PEs
- The elements of the weight matrices are streamed in
- Each column is assigned to a convolution window
- The convolution window is a set of all the pixels in the IFMAP which are required to generate a single OFMAP



Weights

a3	b3	c3	d3	e3	
a2	b2	c2	d2	e2	3
a1	b1	c1	d1	e1	2
f1	g1	h1	i1	j1	2
k1	l1	m1	n1	o1	2
p1	q1	r1	s1	t1	2
u1	v1	w1	x1	y1	3

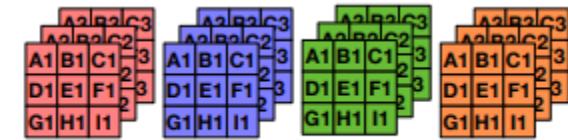
Input Feature Map



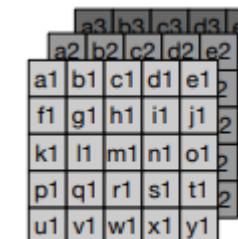


Input Stationary (IS)

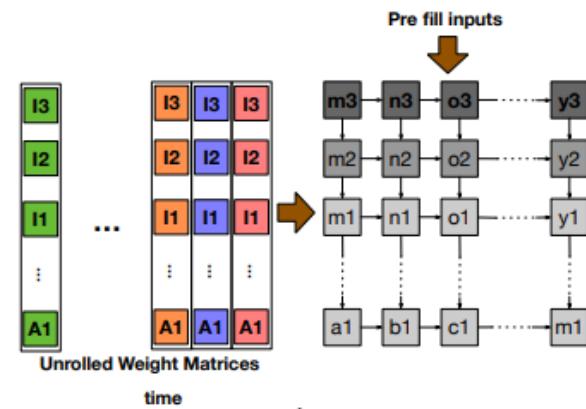
- Once the inputs are fed in, the elements of the weight matrices are streamed in from the left edge
- The reduction is performed over a given column
- The convolution windows are kept around until all the computations requiring these elements are done



Weights



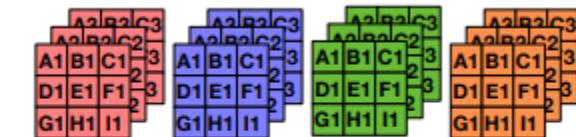
Input Feature Map





Input Stationary (IS)

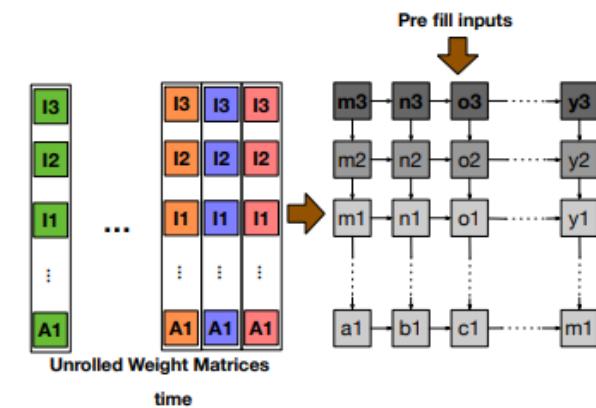
- Benefits
 - Lower SRAM bank requirements as compared to OS
- Shortcoming
 - The cost and runtime compared to WS varies by workload



Weights

a3	b3	c3	d3	e3	
a2	b2	c2	d2	e2	3
a1	b1	c1	d1	e1	2
f1	g1	h1	i1	j1	2
k1	l1	m1	n1	o1	2
p1	q1	r1	s1	t1	2
u1	v1	w1	x1	y1	

Input Feature Map

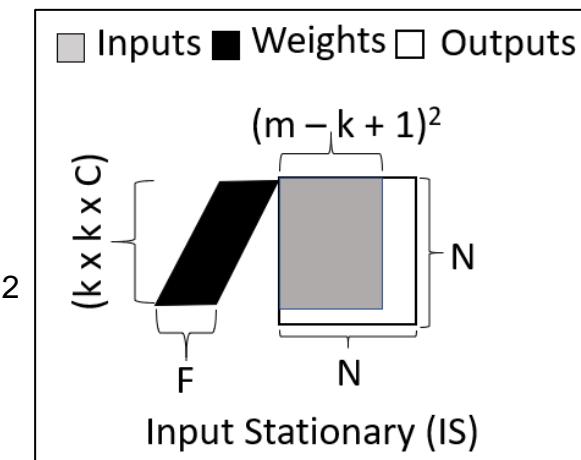




Latency Analysis of Input Stationary

- **The input stationary in the systolic array**

- Weights stream into the systolic array horizontally and takes $(k \times k \times C - 1) + F$ cycles
- Weights also take $(m - k + 1)^2$ cycles to pass through entire inputs
- Pre-load inputs takes $(k \times k \times C)$ cycles
- Total cycles
 - $(k \times k \times C) + (k \times k \times C - 1) + F + (m - k + 1)^2$





Parameters of CNN Network

Parameters	
m	The width and height of input feature map
K	The width and height of filter
F	The number of filters
C	The number of channels
N	The width and height of spatial array



Dataflow Cost Analysis

- OS minimizes output reads (0)
- WS saves # of weight reads (E)
- IS saves # of input reads (E)

R: size of filter weight
E: size of output activations

These dataflows only
reduce a specific reads.
Could we do better ?

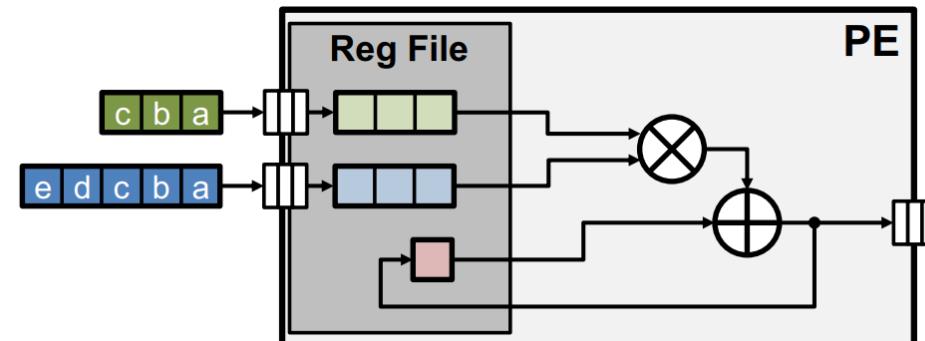
	OS	WS	IS
MACs	E^*R	E^*R	E^*R
Weight Reads	E^*R	R	E^*R
Input Reads	E^*R	E^*R	E
Output Reads	0	E^*R	E^*R
Output Writes	E	E^*R	E^*R



Row Stationary (RS)

- Minimize data reuse at
- Optimize for overall data type energy efficiency

$$\begin{array}{c} \text{Filter} \\ \begin{matrix} a & b & c \\ \hline \end{matrix} \end{array} * \begin{array}{c} \text{Input Fmap} \\ \begin{matrix} a & b & c & d & e \\ \hline \end{matrix} \end{array} = \begin{array}{c} \text{Partial Sums} \\ \begin{matrix} a & b & c \\ \hline \end{matrix} \end{array}$$

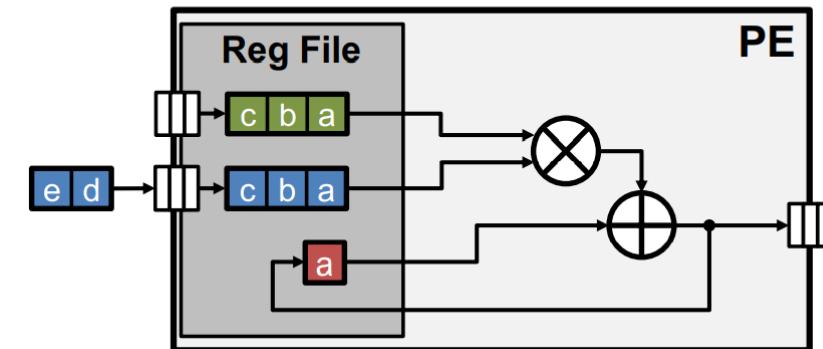
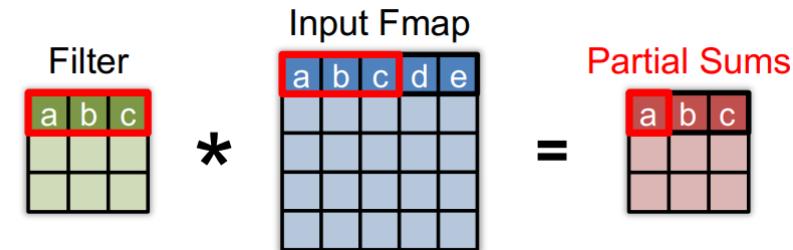


Chen et al., ISCA 2017



How does RS work ?

- Keep the row of filter weights stationary in RF of a PE
- PE does MACs for each sliding window of ifmap at a time
- Use only one memory space to accumulate Psums
- Overlap ifmap between different sliding windows -> reuse ifmap

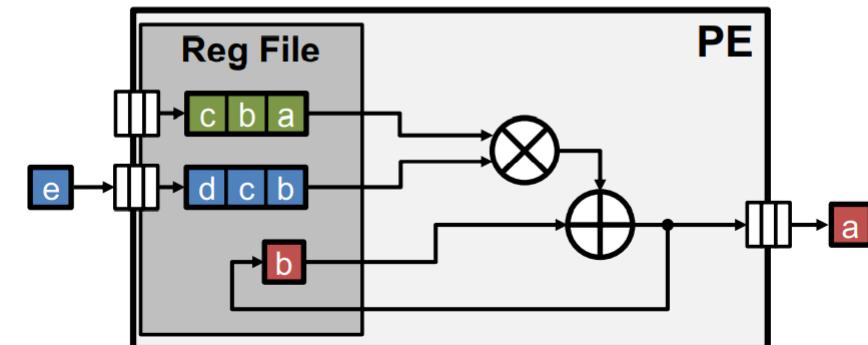
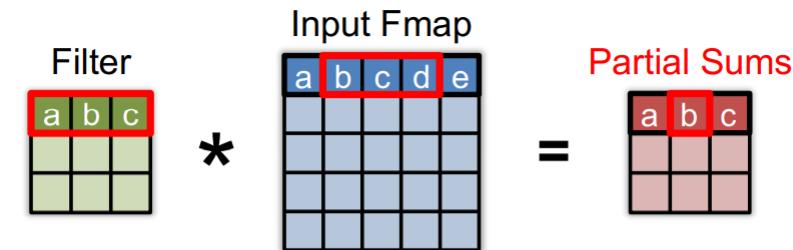


Chen et al., ISCA 2017



How does RS work ?

- Ifmap sliding window right shifts
- Pop the value “a” out of RF
- Accumulate Psum “b”

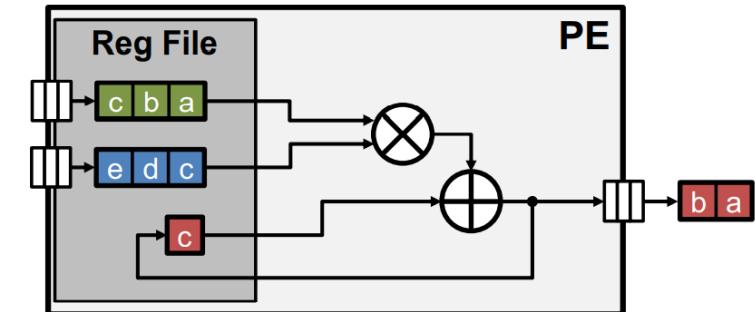
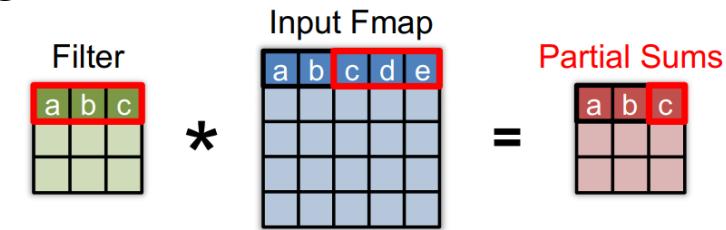


Chen et al., ISCA 2017



How does RS work ?

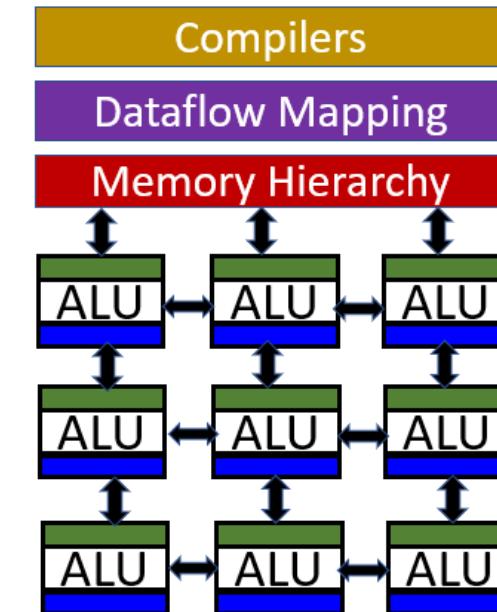
- Ifmap sliding window continues to right shift
- Pop out the value “b” in RF
- Accumulate psum “c”





What do we learn from DNN Dataflow ?

- DNN layer shape and hardware resources provided determine the energy efficiency of dataflow mapping
- How can the fixed-size PE array accommodate different layer shapes?
- Known DNN layer shapes offline, could compiler/runtime system guide the mapping ?





Takeaway Questions

- What are the purposes of dataflow used by DNN applications?
 - (A) Reduce the data movement across off-chip memory
 - (B) Improve the clock frequency of PE
 - (C) Decrease the energy consumption of spatial array accelerator
- What kind of dataflow implemented by the PE on the right-hand side?
 - (A) WS
 - (B) IS
 - (C) OS

