



Accelerator Architectures for Machine Learning (AAML)

Lecture 6: Digital DNN Accelerator

Tsung Tai Yeh
Department of Computer Science
National Yang-Ming Chiao Tung University



Acknowledgements and Disclaimer

- Slides was developed in the reference with
Joel Emer, Vivienne Sze, Yu-Hsin Chen, Tien-Ju Yang, ISCA 2019
tutorial
Efficient Processing of Deep Neural Network, Vivienne Sze, Yu-Hsin
Chen, Tien-Ju Yang, Joel Emer, Morgan and Claypool Publisher, 2020
Yakun Sophia Shao, EE290-2: Hardware for Machine Learning, UC
Berkeley, 2020
CS231n Convolutional Neural Networks for Visual Recognition,
Stanford University, 2020
CS224W: Machine Learning with Graphs, Stanford University, 2021



Outline


- Reconfigurable Deep Learning Accelerators
 - Reconfigurable FPGA
 - SambaNova Reconfigurable Dataflow Unit (RDU)
 - Coarse grained reconfigurable array (CGRA)
 - Meta MTIA AI Accelerator
 - Tenstorrent AI Processor
 - Wafer-scale AI chip -- Cerebras



Reconfigurable Deep Learning on FPGA



Spectrum of Architectures for Deep Learning






Why Reconfigurable Computing?

- AI accelerators improves 100X performance/energy compared to general-purpose processor
- But new hardware is sophisticated and expensive
 - Especially in cutting-edge manufacturing processes


Can we bridge the gap btwn general-purpose & accelerators w/out custom hardware?






Reconfigurable Computing

- Basic idea
 - A spatial array of processing elements (PEs) & memories with a configurable network
 - Map your computation spatially onto the array
 - Goal: programmable with near ASIC efficiency






Basic FPGA Design






Modern FPGAs

- FPGAs are coarse-grain today
 - Hardened logic in LUTs
 - “DSP blocks” to implement wide add/mul efficiently
 - Dense memories distributed throughout fabric





Mode Name	Mathematical Function
Multiplication Mode	$X \cdot Y$
Adder or Subtract Mode	$(X + Y)$ or $(X - Y)$
Multiply-Add/Subtract	$(X \cdot Y) + Z$ or $(X \cdot Y) - Z$
Multiply Accumulate Mode	$(X \cdot Y) + Acc$ or $(X \cdot Y) - Acc$
Vector One Mode	$(X \cdot Y) + Chain In$





High-Level Synthesis (HLS)

- Automated optimization and scheduling
- High portability against different PDK or PPA requirements
- Short design cycle






Challenges of HLS Accelerator Design

- **Time consuming**
 - Manual architecture and micro-architecture design, manual C/C++ code rewriting
- **Suboptimal**
 - Empirical parameter tuning, like parallel factors, buffer sizes, tiling sizes, etc..
- **Low flexibility**
 - Only support a small set of models



Accelerator Design Languages (ADLs)

- **Pythonic**
- **Maintainability**
 - Decoupled hardware customizations
- **Composability**
 - All the kernels, primitives, and schedules should be composable to form complex designs





SambaNova Reconfigurable Dataflow Unit (RDU)




Plasticine Architecture

- **Plasticine architecture**
 - A reconfigurable architecture for parallel patterns (Raghu, ISCA 2017)
 - **Pattern Compute Unit (PCU)**
 - Reconfigurable pipeline with multiple stages of SIMD functional units (FUs)
 - **Pattern Memory Unit (PMU)**
 - A banked scratchpad memory
 - **The compiler**
 - Maps the computation of inner loops to PCUs
 - Most operands are transferred directly between FUs without scratchpad access or inter-PCU communication



Plasticine Architecture Overview


- Data access address calculation occurs while the PCU is working
- Each DRAM channel is accessed using several address generators (AG) on two sides of the chip
- Multiple AGs connect to an address coalescing unit for memory requests





Plasticine PCU Architecture

- **Pattern Compute Unit (PCU)**
 - Each stage's SIMD lane contains a FU and associated pipeline register (PR)




1. **Scalar**: uses to communicate single words of data
2. Each **vector** communicates one word per line in the PCU
3. **Control** signals at the start or end of execution of a PCU



Plasticine PMU Architecture

- **Pattern Memory Unit (PMU)**


- Contains a scratchpad memory and address calculation
- Calculates address only needs simple scalar math
- Has simpler FUs than ones in PCUs





Reconfigurable Dataflow Unit (RDU)

- **SambaNova RDU**
 - **Pattern Compute Units**
 - BF16 with FP32 accumulation
 - Support FP32, Int32, Int16, Int8
 - **Pattern Memory Unit**
 - Memory transformation
 - **Dataflow optimization**
 - Tiling
 - Nested pipelining
 - Operator parallel streaming






Dataflow Exploits Data Locality / Parallelism

- **Software-hardware co-design architecture**
 - Dataflow captures data locality and parallelism
 - Flexible time and space scheduling to achieve higher utilization
 - Flexible memory system and interconnect to sustain high compute throughput
 - Custom dataflow pipeline





Chip and Architecture Overview

- **RDU Tile**
 - Compute and memory components
 - A programmable interconnect
- **Tile resource management**
 - Combine adjacent tiles to form a larger logical tile
 - Each tile controlled independently
 - Allow different applications on separate tiles concurrently
- **Memory access**
 - Direct access to TBs DDR4 off-chip memory
 - Memory-mapped access to host memory






RDU Tile





Pattern Compute Unit (PCU)


- **Pattern Compute Unit (PCU)**
 - Compute engine
- **Reconfigurable SIMD data path**
 - For dense and sparse tensor algebra in FP32, BF16, and integer data format
- **Programmable counters**
 - Program loop iterators
- **Tail unit**
 - Accelerates functions such as exp, sigmoid





Pattern Memory Unit (PMU)


- **Pattern Memory Unit (PMU)**
 - **On-chip memory system**
 - **Banked SRAM arrays**
 - Write and read multiple high bandwidth SIMD data stream concurrently
 - **Address ALUs**
 - Address calculation for arbitrarily complex accesses
 - **Data align**
 - Tensor layout transformation





Switch and On-chip Interconnect


- **Switch**
 - Programmable packet-switched interconnect fabric
- **Independent data and control buses**
 - Suit different traffic classes
- **Programmable routing**
 - Flexible chip bandwidth allocation to concurrent stream
- **Programmable counters**
 - Outer loop iterators
 - On-chip metric collection





Interface to I/O Subsystem


- **Address ALUs**
 - Address calculation for arbitrarily complex accesses
- **Coalescing Units**
 - Enable transparent access to memories across RDUs and host memory
- **Address space manager**
 - Programmable, variable length segments






Operator Mapping

SOFTMAX: $\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$





LAYERNORM: $y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$





Operator Mapping (Softmax)



SOTMAX: $\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$





Pipelined in Space



LAYERNORM: $y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$





Pipelined in Space + Fused


LAYERNORM: $y = \frac{x - E[x]}{\sqrt{Var[x] + \epsilon}} * \gamma + \beta$





Spatial Dataflow within an RDU

- The dataflow removes
 - Memory traffic and host communication overhead






CGRA



Coarse grained reconfigurable array (CGRA)

- **Coarse grained reconfigurable array (CGRA)**
 - Multiple processing elements (PEs)
 - Each PE has ALU-like functional unit
 - **Array configurations vary by**
 - Array size
 - Functional units
 - Interconnection network
 - Register file architectures
 - CGRAs can achieve **power-efficiency** of several 10s of GOps/sec per Watt (why?)
 - Samsung SRP processor (embedded and multimedia apps)






Key features of CGRA accelerators

- **Software-pipelining execution mapping**
 - Accelerate loops with low parallelism
 - Loops with loop-carried dependence, loops with high branch divergence
- **Avoid von-Neumann architecture bottleneck**
 - CGRAs are not subjected to dynamic fetch and decoding of instructions
 - CGRA instructions are in a **pre-decoded form** in the instruction memory
 - **PE transfers data directly** among each another
 - Without going through a centralized registers and memory



Loop execution on the CGRA

Data dependency graph




Mapping data dependency graph to CGRA



Loop:

$$t1 = (a[i]+b[i]-k)*c[i]$$
$$d[i] = \sim t1 \& 0xFFFF$$


Execution time: 1





Loop execution on the CGRA

Data dependency graph




Mapping data dependency graph to CGRA



Loop:

$$t1 = (a[i]+b[i]-k)*c[i]$$
$$d[i] = \sim t1 \& 0xFFFF$$

Execution time: 2






Loop execution on the CGRA

Loop:

$$t1 = (a[i]+b[i]-k)*c[i]$$
$$d[i] = \sim t1 \& 0xFFFF$$


Data dependency graph



Mapping data dependency graph to CGRA



Execution time: 3






Loop execution on the CGRA

Loop:

$$t1 = (a[i]+b[i]-k)*c[i]$$
$$d[i] = \sim t1 \& 0xFFFF$$


Data dependency graph



Mapping data dependency graph to CGRA



Execution time: 6





Takeaway Questions

- What are hardware components used by RDU ?
 - (A) Pattern computer unit (PCU)
 - (B) Pattern memory unit (PMU)
 - (C) Interconnect network router
- What are features of CGRAs ?
 - (A) Customized PEs
 - (B) Software-pipelining execution mapping
 - (C) Reconfigurable dataflow



Meta MTIA2 ASIC



MTIA Model Inference Engine

- **MTIA Architecture**

- 8x8 array of processing elements (PEs)
- Network-on-chip (NoC) connects to a set of on-chip SRAMs
- SRAMs are shared by the PEs and to off-chip memory
- Control core is quad-core RISC-V processor






MTIA Model Inference Engine

- **MTIA PE**

- Each PE includes 384 KB local memory
- RISC-V core issues commands to Command Processor to offload computations to fixed-function units






MTIA Model Inference Engine

- **MTIA PE**

- RISC-V Vector extension (64B wide)
- Memory Layout Unit (MLU) performs memory-layout transformation
- Dot Product Engine (DPE) performs GEMM






MTIA Model Inference Engine

- **MTIA PE**

- Dot Product Engine (DPE)
performs GEMM

- Two $32 \times 32B \times 32$ MAC tiles
- 2.76 TFLOPS/s per PE with FP16/ BF16 and output in FP32 \rightarrow 2:4 sparsity for weights






MTIA Model Inference Engine

- **MTIA PE**


- Reduction Engine (RE)
 - Stores matrix multiply results as they are accumulated
- SIMD Engine (SE)
 - Performs quantization and nonlinear funcs
 - Includes LUT for approximating nonlinear functions





MTIA Model Inference Engine

- **MTIA PE**
 - Fabric Interface (FI)
 - Acts as DMA engine to transfer data in and out of PE's local memory through NoC





MTIA Model Inference Engine

- **MTIA Unique Memory Hierarchy**
 - Uses a large SRAM (256 MB) backed by LPDDR DRAM
 - Avoids HBM to reduce cost and power consumption
 - Meet latency requirements of recommendation models
 - The recommendation models exhibit significant locality
 - Similar to Cerebras and Groq accelerator
 - SRAM provides 2.7TB/s of bandwidth
 - Performance drops sharply as models reach a complexity and size that exceeds the SRAM capacity




MTIA Model Inference Engine

- **New Feature in MTIA2i**
 - Dynamic INT8 quantization
 - Leveraging the reduction engine to identify the min and max values per batch
 - Channel-wise symmetric dynamic INT8 for FC layers
 - Compression
 - Supports lossless asymmetric numerical system (ANS) compression for weights
 - Achieves up to a 50% compression ratio

MTIA Model Inference Engine

- **MTIA2i Software Stack**
 - Support PyTorch eager mode
 - TorchDynamo enables symbolic tracing to capture models with dynamic shapes
 - TorchInductor generates Triton codes for PyTorch operator and operator fusion






Tenstorrent AI Hardware



Tenstorrent AI Processor


- **The NoC grid of Tenstorrent Wormhole**
 - Most compute elements called a Tensix core
 - D = DRAM
 - T = Tensix
 - E = Ethernet
 - A = ARC
 - P = PCI/e





Tenstorrent AI Processor

- **Each Tensix contains**
 - 5 “Baby” RISC-V CPUs
 - 5 stage pipelined, single issue CPU
 - Handle instruction dispatch
 - 2 NoC interface
 - A vector unit
 - A matrix/tensor unit





Tenstorrent AI Processor


- **The Dataflow of Tensix**
 - NoC 0 reads data from DRAM
 - Unpacker unpacks the data into a format that can be processed by the matrix/tensor unit
 - Matrix/tensor unit performs the computation
 - Packer packs the result back into a format for storage
 - NoC 1 sends the result to DRAM



Tenstorrent AI Processor

- **The Tensix NoCs**


- The NoCs are full-duplex
- Both NoCs supports sending and receiving data at the same time
- NoC 0 running in the opposite direction of NoC1
- The uni-directional design of the NoCs reduce power and area
- The 2D torus topology ensures that every points on the chip remain accessible from every other point





Tenstorrent AI Processor


- **3 kernels on Tensix**
 - 2 data movement kernels can be developed separately
 - The 3 compute cores work cooperatively to perform the computation





Tenstorrent AI Processor

- **Kernel Synchronization**
 - Circular buffers backed by hardware mutexes and SRAM
 - Kernels wait for available space in the circular buffer, write data to it, and then mark that data as ready






Tenstorrent AI Processor

- **SRAM, interleaved and shared buffers**
 - Allows data, intermediate tensor or operator buffers to live in SRAM
 - This reduces the need for expensive DRAM accesses
 - The “interleaved” mode is used for memory access
 - The “shared mode” reduces the distance and cross talk when accessing DRAM for certain operations



Tenstorrent AI Processor

- **Native tile based computing**
 - The Tensix natively performs operations on 32×32 tiles
 - 32×32 tiles are small enough for hardware to digest in a few cycles





Tenstorrent AI Processor

- **Cache Hierarchy**
 - There is no cache hierarchy on Tenstorrent chips
 - Provides direct access to SRAM across the entire chip
 - SRAM is not cache and no automatic caching occurs
 - Data must be explicitly brought into SRAM




Wafer-scale AI chip -- Cerebras



Largest AI chip

- 46,225 mm² silicon
- 1.2 trillion transistors
- 400,000 AI optimized cores
- 18 Gigabytes of on-chip memory
- 9 Pbyte/s memory bandwidth
- 100 Pbit/s fabric bandwidth
- TSMC 16 nm process



Cerebras WSE

21.1 Billion
Transistors
815 mm² silicon



GPU



Why big chips ?

- Big chips process data more quickly
 - Cluster scale performance on a single chip
 - GB of fast memory 1 clock cycle from core
 - On-chip interconnect orders of magnitude faster than off-chip
 - Model-parallel, linear performance scaling
 - Training at scale, with any batch size, at full utilization




Cerebras Architecture

- Core optimized for neural network primitives
- **Flexible, programmable core**
 - NN models are evolving
- **Designed for sparse compute**
 - Workloads contain fine-grained sparsity (where are these sparsity from ?)
- **Local memory**
 - reusing weight & activations
- **Fast interconnect**
 - Layer-to-layer with high bandwidth and low latency



Cerebras programmable core


- Flexible cores optimized for **tensor operations**
 - General ops for control processing
 - e.g. arithmetic, logical, LD/ST, branch
 - Optimized tensor ops for data processing
 - **Tensor operands**
 - e.g. fmac $[Z] = [Z], [W]$, a 3D 3D 2D





Sparse compute engine


- **Nonlinear activations** naturally create **fine-grained sparsity**
- **Dataflow scheduling in hardware**
 - Triggered by data
 - Filters out sparse zero data
 - Skips unnecessary processing
- Fine-grained execution datapaths
 - Small cores with independent instructions
 - Efficiently processes dynamic, non-uniform work





Cerebras memory architecture

- **Traditional memory designs**
 - Centralized shared memory is slow & far away
 - Requires high data reuse (caching)
 - Local weights and activations are local -> low data reuse
- **Cerebras memory architecture**
 - All memory is fully distributed along compute
 - Datapath has full performance from memory



■ Core ■ Memory



High-bandwidth low-latency interconnect

- **2D mesh topology** effective for local communication
 - High bandwidth and low latency for local communication
 - All HW-based communication avoids SW overhead
 - Small single-word message



Challenges of wafer scale

- Building a 46,225 mm², 1.2 trillion transistor chip
- **Challenges include**
 - Cross-die connectivity
 - Yield
 - Thermal expansion
 - Package assembly
 - Power and cooling



Takeaway Questions

- What are challenges to build a large chip for NN applications ?
 - (A) Power and cooling
 - (B) Fault tolerance for defected dies
 - (C) Package assembly
- How does Cerebras tackle the DNN sparsity ?
 - (A) Customized sparse core
 - (B) Data-driven dataflow scheduling
 - (C) Filters out sparse zero data