



ML Compiler on Heterogenous Computer Architecture

Tsung Tai Yeh
Department of Computer Science
National Yang-Ming Chiao Tung University



Acknowledgements and Disclaimer

- Slides was developed in the reference with CS 15-779, Advanced Topics in Machine Learning Systems (LLM Edition), CMU, 2025
- AI System: <https://github.com/Infrasys-AI/AISystem>




Outline

- Heterogeneous Computer Architecture
 - CPU+GPU
 - CPU+ASIC
- ML Compiler
 - MLIR
 - IREE
- MegaKernel + Mirage on GPU
 - Domain-Specific Language



What is heterogeneous SoC?

- **Heterogenous computer architecture**
 - A chip contains CPU and multiple specialized functional units




Chip	Tesla - FSD Chip	Qualcomm - Snapdragon 865 (Galaxy S20, March 6 2020)
Technology Node	Samsung 14 nm process	TSMC's advanced 7nm (N7P)
CPU	3x (4-core) Cortex-A72	4x Cortex-A77, 4x Cortex-A55 (4 high power, 4 low power)
GPU	Custom GPU, 0.6 TFLOPS @ 1 Ghz	Adreno 650, 1.25 TFLOPS @ 700 MHz-ish
NPU (AI accelerator)	2x Tesla NPU, each 37 TOPS (total 74 TOPS)	Hexagon 698 @ 15 TOPS
Memory (Cache)	2x 32MB SRAM for NPUs	1 MB L2, 4 MB L3, and 3 MB system wide cache
Memory (RAM)	8GB LPDDR4X, 2x 64-bit, Bandwidth 111 GB/s	16GB LPDDR5, 4x 16-bit , Bandwidth 71.30 GB/s
ISP (Image signal processor)	24-bit? 1 billion pixels per second	Spectra 480, dual 14-bit CV-ISP 2 Gpixel/s, H.265 (HEVC)
Secure Processing Unit	"Security system", verify code has been signed by Tesla.	Qualcomm SPU230, EAL4+ certified



Why heterogeneous computer architecture?

Stuttering


- Transistors per chip, '000
- Clock speed (max), MHz
- Thermal design power*, w
- Chip introduction dates, selected



General purpose processor is not getting faster and power-efficient because of **Slowdown of Moore's Law and Dennard Scaling**




Why heterogeneous computer architecture?






Evolution of Computer Architecture





Hetero Computer Architecture (CPU + GPU)



- Step 1: CPU sends data from its host memory to device memory
- Step 2: CPU asks GPU to begin the execution
- Step 3: GPU sends results back to the CPU
- What are advantages when using this hetero. architecture?





Hetero Computer Architecture (CPU + ASIC)

- Two types of heterogeneous computer architecture
 - Discretized CPU+ASIC (separated DRAM)
 - Integrated CPU+ASIC (shared DRAM)





Hetero Computer Architecture (CPU+ASIC)


- **Post-Moore era and dark silicon**
 - A suite of accelerators on chip are rising
 - Applications will only use a subset of processors/accelerators at a time
 - Such a heterogeneous architecture is compatible with dark silicon



2010 Apple A4
65 nm TSMC 53 mm²
4 accelerators



2014 Apple A8
20 nm TSMC 89 mm²
28 accelerators




2019 Apple A12
7 nm TSMC 83 mm²
42 accelerators



Hetero Computer Architecture (GPU)


- GPU includes FP, SFU (Special Functional Unit), Ray Tracing (RT) Core, and Tensor Core





Challenges of Hetero. Computer Architecture?


- Program Compilation
 - Programming model?
 - Data/Kernel mapping/partition?
 - Concurrent execution?






Challenges of Hetero. Computer Architecture?


- Program Compilation



Hardware defines
software



Software defines
hardware





Hardware-Software Co-
Design



Challenges of Hetero. Computer Architecture?


- Hardware
 - Packaging on Chiplet
 - Network-on-Chip (NoC)
 - Photonic Integrated Circuit





Challenges of Hetero. Computer Architecture?

- Trade-off the performance and flexibility






Takeaway Questions

- How to improve the performance of processor?
 - (A) Increase the size of cache
 - (B) Add specialized engines in the processor
 - (C) Utilize high bandwidth memory (HBM)
- What are benefits of heterogeneous computer architecture?
 - (A) Improve energy efficiency of the processor
 - (B) Facilitate parallel computing
 - (C) Reduce memory access latency




Computer Language Stacks






Compiler Basics






LLVM Compiler Architecture






LLVM Compiler Architecture






What is AI Compiler?

- Translate the operators of ML models to hardware






What is AI Compiler?





AI Compiler: Stage I


- ML Model graph
 - Static model graph
 - Python->Onnx
 - Graph rewrite/Optimizer
- Performance
 - Op kernel libraries (cuDNN, CMSIS-NN ...)
 - More performance improve using Op scheduling, tiling, fusion





AI Compiler: Stage II

- ML Model graph
 - Transforms PyTorch expression into IR
 - Optimizes Tensor IR
- Performance
 - Operator lowering
 - Inter-op optimization
 - Static/dynamic graphs
 - Not only rely on the customized Op Lib





AI Compiler Frontend

- **Front-end compilation**

- **Goal**
 - Parse model graphs from different AI system frameworks
 - Transforms model graphs into IR
- **Tasks**
 - Input format of ML models ((TensorFlow, PyTorch, ONNX ...))
 - Transformation: transform model into united expression
 - TVM Relay, PyTorch Aten (TorchScript)
 - High-level IR/Graph IR
 - Hardware independent
 - Operator/Tensor expression



AI Compiler Frontend


- **Front-end compilation**
 - **Tasks**
 - Computational Graph Optimizations
 - Algebraic simplification
 - Operator Fusion
 - Operator Sinking
 - Static memory planning
 - Tensor Layout transformation



AI Compiler High-Level IR

- Layer-level IR
 - Express ML model structure as a calculation graph
 - High-level abstraction
 - Optimization
 - DSE, operator fusion..
 - Cross-platform


Represent High level Deep Learning Computations





AI Compiler Graph IR


- Graph IR
 - Express ML model as a computation graph
 - Tensor
 - Operator
 - Dependency





AI Compiler Graph IR

- Tensor
 - Shape [2, 3, 4, 5]
 - [N, C, H, W] [N, H, W, C]
 - Type [int, float, string, ...]
- Operator
 - Algebra operator
 - Pre-defined operators





Add	Log	While
Sub	MatMul	Merge
Mul	Conv	BroadCast
Div	BatchNorm	Reduce
Relu	Loss	Map
Floor	Sigmoid



AI Compiler Graph IR

- Directed Acyclic Graph (DAG)
 - Operator, Tensor, control flow (For/While), dependency



Forward


Backward



AI Compiler Graph IR

- Static Computational Graph
 - AI system framework (e.g. TensorFlow) parses API used to describe ML model
 - Fixed before execution
 - Use static data structure to describe model graph topology

```
1  class Network(nn.Cell):
2      def __init__(self):
3          super().__init__()
4          self.flatten = nn.Flatten()
5          self.dense_relu_sequential = nn.SequentialCell(
6              nn.Dense(28*28, 512),
7              nn.ReLU())
8
9      def construct(self, x):
10         x = self.flatten(x)
11         logits = self.dense_relu_sequential(x)
12         return logits
13
```






AI Compiler Graph IR

- Dynamic Computational Graph
 - Built on-the-fly as operations are performed
 - Define-by-run offers greater flexibility
 - Good for handling complex and variable-structured data
 - Time-series data: audio
 - Graph data: social networks
 - Multi-modal data: combinations of different variable-structured data types




AI Compiler Graph IR

- Operator fusion



A is called twice




A is called once
Buffer A's output




AI Compiler Graph IR

- Operator fusion



Three kernel calls
(A, B, C)




Reuse the intermediate
data buffer



AI Compiler Graph IR

- Operator fusion
 - Reduce memory RW of intermediate tensors






AI Compiler Graph IR

CFG (Control Flow Graph)

- How to fuse operators ?
 - TVM dominator tree (In a DAG)
 - Dominator
 - Node X dominates node Y iff all paths from the entry to Y go through X.
 - Node A dominates node C (A dom C)







AI Compiler Graph IR

- How to fuse operators ?
 - TVM dominator tree (In a DAG)


Dominator Tree:





AI Compiler Graph IR


- The purpose of dominator tree
 - Check the path of each node to dominator node
 - Fuses the node that does not affect the rest of nodes
 - How to create a dominate tree?
 - Create DFS tree based on DAG
 - Create DOM (dominator) tree
 - Examine a group of nodes to check if multiple nodes can be fused





AI Compiler Graph IR


- Rule of operator fusion
 - **Injective (one-to-one map)** : Add, pointwise
 - **Reduction**: sum/max/min
 - **Complex-out-fusable**
: conv2D
 - **Opaque** (cannot be fused): sort





AI Compiler Graph IR


- Data layout alignment
 - Unaligned tensor data will increase the memory transactions





AI Compiler Graph IR



- Data layout (N, C, H, W)
 - N: batch; N: Height; W: width; C: Channels
 - NCHW: arrange data in the same channel in the a consecutive memory space
 - Good for the computations of GPU (data parallel)





AI Compiler Graph IR

- Data layout (N, H, W, C)
 - NHWC: arrange the data having the same location in different channels in a consecutive memory space e.g. Conv1x1





AI Compiler Graph IR

- Data layout (N, C, H, W)
 - PyTorch on NPU/GPU uses **NCHW** data layout
 - TensorFlow use **NHWC** data layout





AI Compiler Graph IR

- Memory optimization
 - Attention memory usage for a deep Transformer (64 layer and 4 heads), recomputed during the backward pass
 - BERT (768 hidden layers) and needs **73GB** memory when the batch size is 64

Data type	Stored	Recomputed
1024 text tokens (several paragraphs)	1.0 GB	16 MB
32×32×3 pixels (CIFAR-10 image)	9.6 GB	151 MB
64×64×3 pixels (Imagenet 64 image)	154 GB	2.4 GB
24,000 samples (~2 seconds of 12 kHz audio)	590 GB	9.2GB




AI Compiler Graph IR

- Memory optimization
 - **Static memory allocation**
 - Parameters, constant, output
 - Allocate memory in the model initialization stage
 - **Dynamic memory allocation**
 - Output tensor, workspace tensor (intermediate tensor)
 - Allocate memory (dynamic: varying batch size, static: fixed batch size)



AI Compiler Graph IR


- Memory optimization
 - **Inplace operation:** overwrite when the next operator is element-wise operator
 - **Memory sharing:** the size of both operators is the same and no data dependency in these two operators





AI Compiler Low-Level IR

- Low-level IR
 - Describes the computation of a ML model in a more fine-grained representation than that in high-level IR
 - Enable the target-dependent optimization
 - Halide-based IR
 - Separation of comp. and schedule
 - Choose the best schedule to specific target platform






AI Compiler Backend

- Back-end compilation
 - Goal
 - Transform ML graph to specific hardware
 - Code generation: LLVM/CUDA/OpenCL ...
 - Tasks
 - Hardware Specific Optimization
 - Memory allocation
 - Parallelization
 - Scheduling
 - Auto Scheduling: polyhedral, Halide



AI Compiler Backend


- Hardware-specific optimizations





AI Compiler Backend


- Hardware intrinsic mapping
 - Transform a certain set of low-level IR to kernels
 - TVM extensible tensorization
 - Declare the behavior of hardware intrinsic and lowering the rule for intrinsic mapping
 - Enable compiler backend apply optimized micro-kernels to a specific pattern of operations





AI Compiler Backend


- Memory allocation and fetching
 - E.g. GPU memory hierarchy requires efficient memory allocation and fetching techniques for improving data locality
 - TVM memory scope
 - Tag a compute stage as shared or thread-local
 - Shared: generates code with shared memory allocation
 - Properly insert memory barrier





AI Compiler Backend


- Memory latency hiding
 - Reordering the execution pipeline
 - In TPU-Accel with decoupled access-execute (DAE)
 - Backend needs to perform scheduling and fine-grained sync to produce the correct and efficient code
 - TVM virtual threading schedule primitive
 - Virtually parallelized threads
 - Barriers + operations = a single instruction stream





AI Compiler Backend


- Loop oriented optimization
 - Loop fusion
 - fuse loops with the same boundaries for better data reuse
 - Sliding window
 - Compute values when needed
 - Store them for data reuse until they no longer required





AI Compiler Backend

- Parallelization
 - Halide uses a schedule primitive called parallel
 - Specify the parallelized dimension of the loops
 - Nested polyhedral model – detect hierarchy parallelization among levels of tiling and striding





AI Compiler Backend

- Back-end compilation
 - Tasks
 - Auto-tuning
 - Parameterization cost model
 - Using kernel libraries
 - NVIDIA cuDNN/TensorRT, AMD MIOpen
 - Low-level IR/ Operator IR
 - Halide IR
 - Compilation scheme
 - Just-In-Time (JIT), Ahead-Of-Time (AOT)




Takeaway Questions

- What are jobs of AI compiler?
 - (A) Handle tensor memory allocation
 - (B) Reorder the execution of the DL operators
 - (C) Generate assembly codes
- How does AI compiler improve the data reuse on the local memory?
 - (A) Use the NCHW data layout
 - (B) Operator fusion
 - (C) Operator lowering



MLIR – Multi-Level Intermediate Representation


- Most high-level languages have their own AST
- ML graphs compilation process is fragmented
- MLIR allows developers to use a unified codebase/framework to do their optimizations and develop some optimizations for multiple inputs





MLIR – Multi-Level Intermediate Representation


- MLIR's input
 - applications, compilers, C program, etc.
- Within MLIR
 - Implement multiple Dialects for distinct inputs
 - Use Dialect to deal with tensors





MLIR – Multi-Level Intermediate Representation


- Once we have an optimal IR
 - MLIR can lower it onto the backends such as LLVM for CPU ...
 - If the targeting hardware is FPGA, TPU, need vendor-tools for final compilation





MLIR – Multi-Level Intermediate Representation

- MLIR Compiler Infrastructure
 - A set of optimization/code conversion/code generation pipeline





MLIR – Multi-Level Intermediate Representation

- MLIR Dialect
 - One way to express IR from other specific IRs
 - Every IR can be transformed in the corresponding MLIR dialect
 - Each programming language's dialect (Tensor dialect, HLO dialect, LLM IR dialect) is inherent from **mlir::Dialect**
 - AST (Abstract Syntax Tree)






MLIR – Multi-Level Intermediate Representation

- MLIR Dialect

- DRR(Dynamic Reconstructed Radiography) - transform different dialect
- ODS(Operation Definition System) – define operation





MLIR – Multi-Level Intermediate Representation

- MLIR Operation
 - Output: %tensor
 - Operation: toy.transpose
 - Input: %tensor
 - Transform tensor $<2x3xf64>$ to tensor $<3x2xf64>$
 - The location of transpose is in “example/file/path”, line 12, 1st word
- Operations
 - [gpu.all_reduce_\(gpu::AllReduceOp\)](#)
 - [gpu.alloc_\(gpu::AllocOp\)](#)
 - [gpu.barrier_\(gpu::BarrierOp\)](#)
 - [gpu.binary_\(gpu::BinaryOp\)](#)
 - [gpu.block_dim_\(gpu::BlockDimOp\)](#)
 - [gpu.block_id_\(gpu::BlockIdOp\)](#)
 - [gpu.cluster_block_id_\(gpu::ClusterBlockIdOp\)](#)

```
%t_tensor = "toy.transpose"(%tensor) {inplace = true} : (tensor<2x3xf64>) ->
tensor<3x2xf64> loc("example/file/path":12:1)
```



MLIR – Multi-Level Intermediate Representation

- Simple Matmul Kernel

```
M = 2          # Rows in arg0
K = 2816       # Columns in arg0, Rows in arg1
N = 1280       # Columns in arg1

# Matrix multiplication with f16 -> f32 promotion
for i in range(M):
    for j in range(N):
        acc = 0.0  # float32 accumulator
        for k in range(K):
            a = float(arg0[i][k])  # f16 -> f32
            b = float(arg1[k][j])  # f16 -> f32
            acc += a * b
        result[i][j] = acc  # store result as float32
```



MLIR – Multi-Level Intermediate Representation

- matmul.mlir


```
#map = affine_map<(d0, d1, d2) -> (d0, d2)>
#map1 = affine_map<(d0, d1, d2) -> (d2, d1)>
#map2 = affine_map<(d0, d1, d2) -> (d0, d1)>
func.func @matmul(%arg0: tensor<2x2816xf16>, %arg1: tensor<2816x1280xf16>) -> tensor<2x1280xf32> {
  %cst = arith.constant 0.000000e+00 : f32
  %0 = tensor.empty() : tensor<2x1280xf32>
  %1 = linalg.fill ins(%cst : f32) outs(%0 : tensor<2x1280xf32>) -> tensor<2x1280xf32>
  %2 = linalg.generic {indexing_maps = [#map, #map1, #map2], iterator_types = ["parallel", "parallel",
"reduction"]} ins(%arg0, %arg1 : tensor<2x2816xf16>, tensor<2816x1280xf16>) outs(%1 : tensor<2x1280xf32>) {
  ^bb0(%in: f16, %in_0: f16, %out: f32):
    %3 = arith.extf %in : f16 to f32
    %4 = arith.extf %in_0 : f16 to f32
    %5 = arith.mulf %3, %4 : f32
    %6 = arith.addf %out, %5 : f32
    linalg.yield %6 : f32
  } -> tensor<2x1280xf32>
  return %2 : tensor<2x1280xf32>
}
```



IREE – Intermediate Representation Execution Environment

- IREE


- A MLIR-based compiler for ML programs
- Takes ML workloads from various frontends (PyTorch ..) and execute on different backends (x86, Arm, NVIDIA GPUs, AMD GPUs ..)





IREE – Intermediate Representation Execution Environment


- IREE Compiler Design






Mega-Kernel on the GPU

- GPU includes multiple specialized engines (CUDA core, Tensor core ..)






Existing Kernel-Per-Operator Approach





Limitations





No Inter-Layer Pipelining
Kernel barriers prevent inter-layer pipelining






Limitations

No Overlapping
Coarse-grained dependency
prevents comp. & comm. overlap





Kernel-Per-Operator v.s. Mega-Kernel





Key Challenges of Mega-Kernel

1. How to manage dependency?

No kernel barriers in mega-kernel

Task Graph

2. How to handle dynamism?

Continuous batching, prefill/decode,
paged/radix attention, speculative decoding

In-Kernel
Parallel Runtime

3. How to optimize performance?

Existing compilers target individual kernels



Mirage Superoptimizer*



Mirage: A SuperOptimizer for ML

- Can we represent FlashAttention as a graph optimization?

Is it possible to implement FlashAttention as a combination of matmul, exp, add, mul, div, etc?





Mirage: A SuperOptimizer for ML


- Key idea: automatically generate highly-optimized GPU kernels for DNNs



- **Less engineering effort:** thousands lines of CUDA code → a few lines of Python code in Mirage
- **Better performance:** outperform existing systems by 1.1-3.5x
- **Faster adaptation:** day-0 support for new models; no manual effort




Hierarchical Graph Representation





Example: RMSNorm & MatMul in LLMs

- Existing systems launch two kernels since Y does not fit in shared memory



$$y_i = \frac{x_i g_i}{\sqrt{\frac{1}{N} \sum_j x_j^2}}$$
$$z_i = \sum_k w_{ik} y_k$$


Performance issues:

1. No shared memory reuse
2. Kernel launch overhead



μ Graph for RMSNorm & MatMul


- Existing systems launch two kernels since Y does not fit in shared memory





High-Performance μ Graph



- Key Challenges to discover High-performance μ Graph
 - How to generate potential μ Graph?
 - How to verify their correctness?
 - Mirage system





Hardware-Customized μ Graphs


- Find μ Graphs similar to expert-written implementations for attention on NVIDIA A100 GPU





Neural Processing Unit (NPU)



- Intel NPU
 - Hardware Acceleration Blocks
 - Handle GEMM, CONV ...
 - Streaming Hybrid Architecture Vector Engines (SHAVE)
 - Perform parallel computing for general needs
 - DMA Engines
 - Moving data between DRAM and software-managed cache





Heterogeneity on ASIC

- NPU
 - MAC engine + Specialized engines
- CPU on AI PC
 - AMD Ryzen AI Pro 300 (CPU+NPU+GPU)





Takeaway Questions

- What are jobs of MLIR?
 - (A) Operator definition
 - (B) Operator lowering
 - (C) Instruction selection
- What are benefits of MegaKernel?
 - (A) Overlapping GPU specialized engine execution
 - (B) GPU register reuse
 - (C) Decrease the kernel launch overhead



Future of AI Compiler

- Future of AI compiler
 - In model inference
 - Ahead-of-Time (AoT) compilation
 - In model training
 - Just-in-Time (JIT) compilation
- The form of IR
 - Need one IR that can support diverse programming language and ML frameworks
 - Good for cross-platform



Future of AI Compiler

- Auto-parallelization
 - Automatic execute ML models through different parallelization approaches
 - Distributed computing (Model training)
 - Parallel computing in one chip
- Auto Code/kernel generation
 - Not only Domain-Specific Language (DSL)
 - Match diverse hardware platforms